
Jupyter Notebook Documentation

Release 6.4.0.dev0

<https://jupyter.org>

Apr 25, 2021

USER DOCUMENTATION

1	The Jupyter Notebook	3
2	User interface components	11
3	Notebook Examples	15
4	What to do when things go wrong	57
5	Changelog	63
6	Comms	99
7	Configuration Overview	101
8	Config file and command line options	105
9	Running a notebook server	125
10	Security in the Jupyter notebook server	133
11	Security in notebook documents	135
12	Distributing Jupyter Extensions as Python Packages	139
13	Extending the Notebook	147
14	Contributing to the Jupyter Notebook	165
15	Developer FAQ	169
16	My Notebook	171
17	Other notebook	173

- [Installation](#)
- [Starting the Notebook](#)

THE JUPYTER NOTEBOOK

1.1 Introduction

The notebook extends the console-based approach to interactive computing in a qualitatively new direction, providing a web-based application suitable for capturing the whole computation process: developing, documenting, and executing code, as well as communicating the results. The Jupyter notebook combines two components:

A web application: a browser-based tool for interactive authoring of documents which combine explanatory text, mathematics, computations and their rich media output.

Notebook documents: a representation of all content visible in the web application, including inputs and outputs of the computations, explanatory text, mathematics, images, and rich media representations of objects.

See also:

See the [installation guide](#) on how to install the notebook and its dependencies.

1.1.1 Main features of the web application

- In-browser editing for code, with automatic syntax highlighting, indentation, and tab completion/introspection.
- The ability to execute code from the browser, with the results of computations attached to the code which generated them.
- Displaying the result of computation using rich media representations, such as HTML, LaTeX, PNG, SVG, etc. For example, publication-quality figures rendered by the [matplotlib](#) library, can be included inline.
- In-browser editing for rich text using the [Markdown](#) markup language, which can provide commentary for the code, is not limited to plain text.
- The ability to easily include mathematical notation within markdown cells using LaTeX, and rendered natively by [MathJax](#).

1.1.2 Notebook documents

Notebook documents contains the inputs and outputs of a interactive session as well as additional text that accompanies the code but is not meant for execution. In this way, notebook files can serve as a complete computational record of a session, interleaving executable code with explanatory text, mathematics, and rich representations of resulting objects. These documents are internally [JSON](#) files and are saved with the `.ipynb` extension. Since JSON is a plain text format, they can be version-controlled and shared with colleagues.

Notebooks may be exported to a range of static formats, including HTML (for example, for blog posts), reStructuredText, LaTeX, PDF, and slide shows, via the [nbconvert](#) command.

Furthermore, any `.ipynb` notebook document available from a public URL can be shared via the [Jupyter Notebook Viewer \(nbviewer\)](#). This service loads the notebook document from the URL and renders it as a static web page. The results may thus be shared with a colleague, or as a public blog post, without other users needing to install the Jupyter notebook themselves. In effect, [nbviewer](#) is simply [nbconvert](#) as a web service, so you can do your own static conversions with `nbconvert`, without relying on `nbviewer`.

See also:

[Details on the notebook JSON file format](#)

1.1.3 Notebooks and privacy

Because you use Jupyter in a web browser, some people are understandably concerned about using it with sensitive data. However, if you followed the standard [install instructions](#), Jupyter is actually running on your own computer. If the URL in the address bar starts with `http://localhost:` or `http://127.0.0.1:`, it's your computer acting as the server. Jupyter doesn't send your data anywhere else—and as it's open source, other people can check that we're being honest about this.

You can also use Jupyter remotely: your company or university might run the server for you, for instance. If you want to work with sensitive data in those cases, talk to your IT or data protection staff about it.

We aim to ensure that other pages in your browser or other users on the same computer can't access your notebook server. See [Security in the Jupyter notebook server](#) for more about this.

1.2 Starting the notebook server

You can start running a notebook server from the command line using the following command:

```
jupyter notebook
```

This will print some information about the notebook server in your console, and open a web browser to the URL of the web application (by default, `http://127.0.0.1:8888`).

The landing page of the Jupyter notebook web application, the **dashboard**, shows the notebooks currently available in the notebook directory (by default, the directory from which the notebook server was started).

You can create new notebooks from the dashboard with the **New Notebook** button, or open existing ones by clicking on their name. You can also drag and drop `.ipynb` notebooks and standard `.py` Python source code files into the notebook list area.

When starting a notebook server from the command line, you can also open a particular notebook directly, bypassing the dashboard, with `jupyter notebook my_notebook.ipynb`. The `.ipynb` extension is assumed if no extension is given.

When you are inside an open notebook, the *File | Open...* menu option will open the dashboard in a new browser tab, to allow you to open another notebook from the notebook directory or to create a new notebook.

Note: You can start more than one notebook server at the same time, if you want to work on notebooks in different directories. By default the first notebook server starts on port 8888, and later notebook servers search for ports near that one. You can also manually specify the port with the `--port` option.

1.2.1 Creating a new notebook document

A new notebook may be created at any time, either from the dashboard, or using the *File* → *New* menu option from within an active notebook. The new notebook is created within the same directory and will open in a new browser tab. It will also be reflected as a new entry in the notebook list on the dashboard.

1.2.2 Opening notebooks

An open notebook has **exactly one** interactive session connected to a kernel, which will execute code sent by the user and communicate back results. This kernel remains active if the web browser window is closed, and reopening the same notebook from the dashboard will reconnect the web application to the same kernel. In the dashboard, notebooks with an active kernel have a `Shutdown` button next to them, whereas notebooks without an active kernel have a `Delete` button in its place.

Other clients may connect to the same kernel. When each kernel is started, the notebook server prints to the terminal a message like this:

```
[NotebookApp] Kernel started: 87f7d2c0-13e3-43df-8bb8-1bd37aaf3373
```

This long string is the kernel's ID which is sufficient for getting the information necessary to connect to the kernel. If the notebook uses the IPython kernel, you can also see this connection data by running the `%connect_info` magic, which will print the same ID information along with other details.

You can then, for example, manually start a Qt console connected to the *same* kernel from the command line, by passing a portion of the ID:

```
$ jupyter qtconsole --existing 87f7d2c0
```

Without an ID, `--existing` will connect to the most recently started kernel.

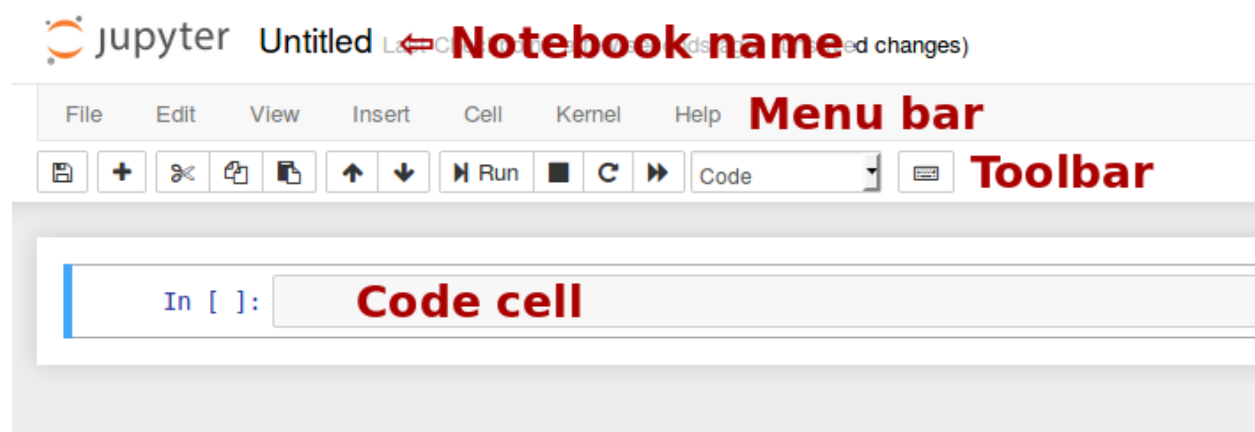
With the IPython kernel, you can also run the `%qtconsole` magic in the notebook to open a Qt console connected to the same kernel.

See also:

[Decoupled two-process model](#)

1.3 Notebook user interface

When you create a new notebook document, you will be presented with the **notebook name**, a **menu bar**, a **toolbar** and an empty **code cell**.



Notebook name: The name displayed at the top of the page, next to the Jupyter logo, reflects the name of the `.ipynb` file. Clicking on the notebook name brings up a dialog which allows you to rename it. Thus, renaming a notebook from “Untitled0” to “My first notebook” in the browser, renames the `Untitled0.ipynb` file to `My first notebook.ipynb`.

Menu bar: The menu bar presents different options that may be used to manipulate the way the notebook functions.

Toolbar: The tool bar gives a quick way of performing the most-used operations within the notebook, by clicking on an icon.

Code cell: the default type of cell; read on for an explanation of cells.

1.4 Structure of a notebook document

The notebook consists of a sequence of cells. A cell is a multiline text input field, and its contents can be executed by using `Shift-Enter`, or by clicking either the “Play” button the toolbar, or *Cell, Run* in the menu bar. The execution behavior of a cell is determined by the cell’s type. There are three types of cells: **code cells**, **markdown cells**, and **raw cells**. Every cell starts off being a **code cell**, but its type can be changed by using a drop-down on the toolbar (which will be “Code”, initially), or via *keyboard shortcuts*.

For more information on the different things you can do in a notebook, see the [collection of examples](#).

1.4.1 Code cells

A *code cell* allows you to edit and write new code, with full syntax highlighting and tab completion. The programming language you use depends on the *kernel*, and the default kernel (IPython) runs Python code.

When a code cell is executed, code that it contains is sent to the kernel associated with the notebook. The results that are returned from this computation are then displayed in the notebook as the cell’s *output*. The output is not limited to text, with many other possible forms of output are also possible, including `matplotlib` figures and HTML tables (as used, for example, in the `pandas` data analysis package). This is known as IPython’s *rich display* capability.

See also:

[Rich Output example notebook](#)

1.4.2 Markdown cells

You can document the computational process in a literate way, alternating descriptive text with code, using *rich text*. In IPython this is accomplished by marking up text with the Markdown language. The corresponding cells are called *Markdown cells*. The Markdown language provides a simple way to perform this text markup, that is, to specify which parts of the text should be emphasized (italics), bold, form lists, etc.

If you want to provide structure for your document, you can use markdown headings. Markdown headings consist of 1 to 6 hash # signs # followed by a space and the title of your section. The markdown heading will be converted to a clickable link for a section of the notebook. It is also used as a hint when exporting to other document formats, like PDF.

When a Markdown cell is executed, the Markdown code is converted into the corresponding formatted rich text. Markdown allows arbitrary HTML code for formatting.

Within Markdown cells, you can also include *mathematics* in a straightforward way, using standard LaTeX notation: $...$ for inline mathematics and
$$...$$
 for displayed mathematics. When the Markdown cell is executed, the LaTeX portions are automatically rendered in the HTML output as equations with high quality typography. This is made possible by [MathJax](#), which supports a [large subset](#) of LaTeX functionality

Standard mathematics environments defined by LaTeX and AMS-LaTeX (the `amsmath` package) also work, such as `\begin{equation}...\end{equation}`, and `\begin{align}...\end{align}`. New LaTeX macros may be defined using standard methods, such as `\newcommand`, by placing them anywhere *between math delimiters* in a Markdown cell. These definitions are then available throughout the rest of the IPython session.

See also:

[Working with Markdown Cells](#) example notebook

1.4.3 Raw cells

Raw cells provide a place in which you can write *output* directly. Raw cells are not evaluated by the notebook. When passed through [nbconvert](#), raw cells arrive in the destination format unmodified. For example, you can type full LaTeX into a raw cell, which will only be rendered by LaTeX after conversion by nbconvert.

1.5 Basic workflow

The normal workflow in a notebook is, then, quite similar to a standard IPython session, with the difference that you can edit cells in-place multiple times until you obtain the desired results, rather than having to rerun separate scripts with the `%run` magic command.

Typically, you will work on a computational problem in pieces, organizing related ideas into cells and moving forward once previous parts work correctly. This is much more convenient for interactive exploration than breaking up a computation into scripts that must be executed together, as was previously necessary, especially if parts of them take a long time to run.

To interrupt a calculation which is taking too long, use the *Kernel, Interrupt* menu option, or the `i, i` keyboard shortcut. Similarly, to restart the whole computational process, use the *Kernel, Restart* menu option or `0, 0` shortcut.

A notebook may be downloaded as a `.ipynb` file or converted to a number of other formats using the menu option *File, Download as*.

See also:

[Running Code in the Jupyter Notebook](#) example notebook

[Notebook Basics](#) example notebook

1.5.1 Keyboard shortcuts

All actions in the notebook can be performed with the mouse, but keyboard shortcuts are also available for the most common ones. The essential shortcuts to remember are the following:

- **Shift-Enter: run cell** Execute the current cell, show any output, and jump to the next cell below. If `Shift-Enter` is invoked on the last cell, it makes a new cell below. This is equivalent to clicking the *Cell, Run* menu item, or the Play button in the toolbar.
- **Esc: Command mode** In command mode, you can navigate around the notebook using keyboard shortcuts.
- **Enter: Edit mode** In edit mode, you can edit text in cells.

For the full list of available shortcuts, click *Help, Keyboard Shortcuts* in the notebook menus.

1.6 Plotting

One major feature of the Jupyter notebook is the ability to display plots that are the output of running code cells. The IPython kernel is designed to work seamlessly with the `matplotlib` plotting library to provide this functionality. Specific plotting library integration is a feature of the kernel.

1.7 Installing kernels

For information on how to install a Python kernel, refer to the [IPython install page](#).

The Jupyter wiki has a long list of [Kernels for other languages](#). They usually come with instructions on how to make the kernel available in the notebook.

1.8 Trusting Notebooks

To prevent untrusted code from executing on users' behalf when notebooks open, we store a signature of each trusted notebook. The notebook server verifies this signature when a notebook is opened. If no matching signature is found, Javascript and HTML output will not be displayed until they are regenerated by re-executing the cells.

Any notebook that you have fully executed yourself will be considered trusted, and its HTML and Javascript output will be displayed on load.

If you need to see HTML or Javascript output without re-executing, and you are sure the notebook is not malicious, you can tell Jupyter to trust it at the command-line with:

```
$ jupyter trust mynotebook.ipynb
```

See [Security in notebook documents](#) for more details about the trust mechanism.

1.9 Browser Compatibility

The Jupyter Notebook aims to support the latest versions of these browsers:

- Chrome
- Safari
- Firefox

Up to date versions of Opera and Edge may also work, but if they don't, please use one of the supported browsers.

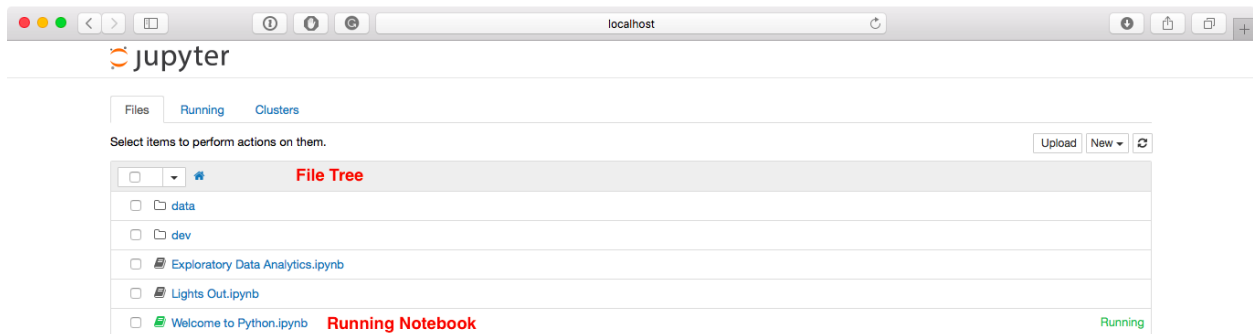
Using Safari with HTTPS and an untrusted certificate is known to not work (websockets will fail).

USER INTERFACE COMPONENTS

When opening bug reports or sending emails to the Jupyter mailing list, it is useful to know the names of different UI components so that other developers and users have an easier time helping you diagnose your problems. This section will familiarize you with the names of UI elements within the Notebook and the different Notebook modes.


2.1 Notebook Dashboard

When you launch `jupyter notebook` the first page that you encounter is the Notebook Dashboard.



2.2 Notebook Editor

Once you've selected a Notebook to edit, the Notebook will open in the Notebook Editor.

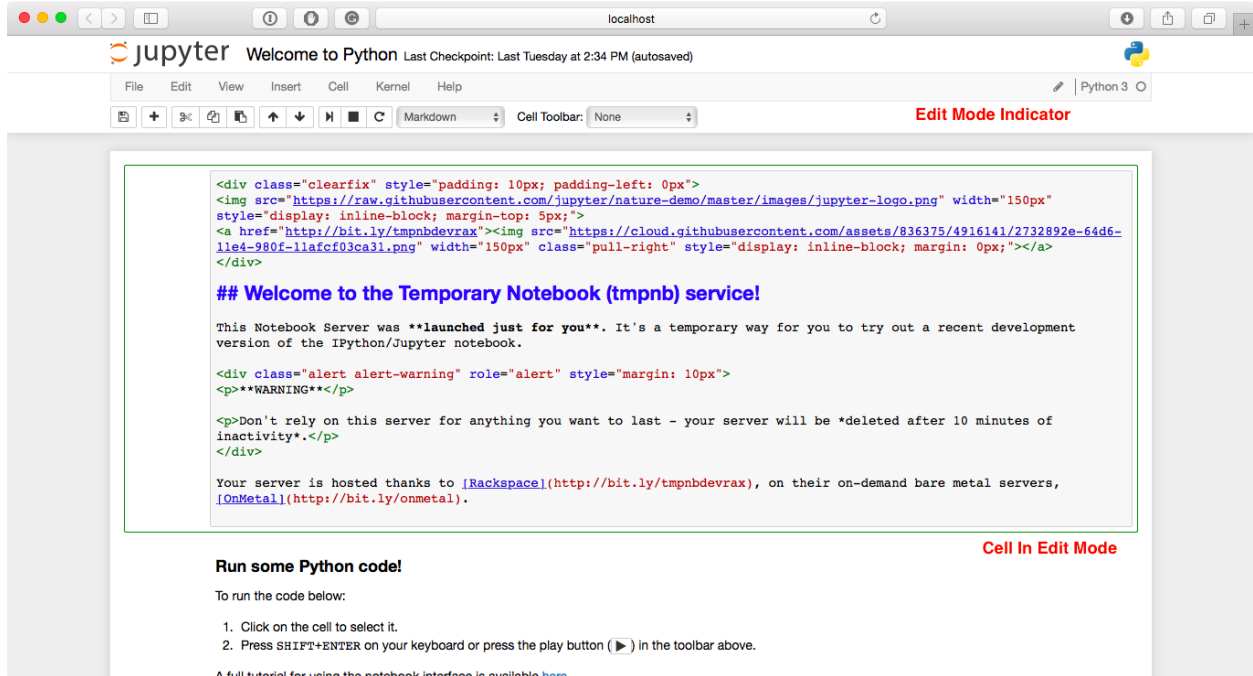


2.3 Interactive User Interface Tour of the Notebook

If you would like to learn more about the specific elements within the Notebook Editor, you can go through the user interface tour by selecting *Help* in the menubar then selecting *User Interface Tour*.


2.3.1 Edit Mode and Notebook Editor

When a cell is in edit mode, the Cell Mode Indicator will change to reflect the cell's state. This state is indicated by a small pencil icon on the top right of the interface. When the cell is in command mode, there is no icon in that location.



2.4 File Editor

Now let's say that you've chosen to open a Markdown file instead of a Notebook file whilst in the Notebook Dashboard. If so, the file will be opened in the File Editor.



NOTEBOOK EXAMPLES

The pages in this section are all converted notebook files. You can also [view these notebooks on nbviewer](#).

3.1 What is the Jupyter Notebook?

3.1.1 Introduction

The Jupyter Notebook is an **interactive computing environment** that enables users to author notebook documents that include: - Live code - Interactive widgets - Plots - Narrative text - Equations - Images - Video

These documents provide a **complete and self-contained record of a computation** that can be converted to various formats and shared with others using email, [Dropbox](#), version control systems (like [git/GitHub](#)) or [nbviewer.jupyter.org](#).

Components

The Jupyter Notebook combines three components:

- **The notebook web application:** An interactive web application for writing and running code interactively and authoring notebook documents.
- **Kernels:** Separate processes started by the notebook web application that runs users' code in a given language and returns output back to the notebook web application. The kernel also handles things like computations for interactive widgets, tab completion and introspection.
- **Notebook documents:** Self-contained documents that contain a representation of all content visible in the notebook web application, including inputs and outputs of the computations, narrative text, equations, images, and rich media representations of objects. Each notebook document has its own kernel.

3.1.2 Notebook web application

The notebook web application enables users to:

- **Edit code in the browser**, with automatic syntax highlighting, indentation, and tab completion/introspection.
- **Run code from the browser**, with the results of computations attached to the code which generated them.
- See the results of computations with **rich media representations**, such as HTML, LaTeX, PNG, SVG, PDF, etc.
- Create and use **interactive JavaScript widgets**, which bind interactive user interface controls and visualizations to reactive kernel side computations.

- Author **narrative text** using the [Markdown](#) markup language.
- Include mathematical equations using **LaTeX syntax in Markdown**, which are rendered in-browser by [MathJax](#).

3.1.3 Kernels

Through Jupyter’s kernel and messaging architecture, the Notebook allows code to be run in a range of different programming languages. For each notebook document that a user opens, the web application starts a kernel that runs the code for that notebook. Each kernel is capable of running code in a single programming language and there are kernels available in the following languages:

- Python(<https://github.com/ipython/ipython>)
- Julia (<https://github.com/JuliaLang/Julia.jl>)
- R (<https://github.com/IRkernel/IRkernel>)
- Ruby (<https://github.com/minrk/iruby>)
- Haskell (<https://github.com/gibiansky/IHaskell>)
- Scala (<https://github.com/Bridgewater/scala-notebook>)
- node.js (<https://gist.github.com/Carreau/4279371>)
- Go (<https://github.com/takluyver/igo>)

The default kernel runs Python code. The notebook provides a simple way for users to pick which of these kernels is used for a given notebook.

Each of these kernels communicate with the notebook web application and web browser using a JSON over ZeroMQ/WebSockets message protocol that is described [here](#). Most users don’t need to know about these details, but it helps to understand that “kernels run code.”

3.1.4 Notebook documents

Notebook documents contain the **inputs and outputs** of an interactive session as well as **narrative text** that accompanies the code but is not meant for execution. **Rich output** generated by running code, including HTML, images, video, and plots, is embedded in the notebook, which makes it a complete and self-contained record of a computation.

When you run the notebook web application on your computer, notebook documents are just **files on your local filesystem with a .ipynb extension**. This allows you to use familiar workflows for organizing your notebooks into folders and sharing them with others.

Notebooks consist of a **linear sequence of cells**. There are three basic cell types:

- **Code cells:** Input and output of live code that is run in the kernel
- **Markdown cells:** Narrative text with embedded LaTeX equations
- **Raw cells:** Unformatted text that is included, without modification, when notebooks are converted to different formats using nbconvert

Internally, notebook documents are **JSON data** with **binary values base64** encoded. This allows them to be **read and manipulated programmatically** by any programming language. Because JSON is a text format, notebook documents are version control friendly.

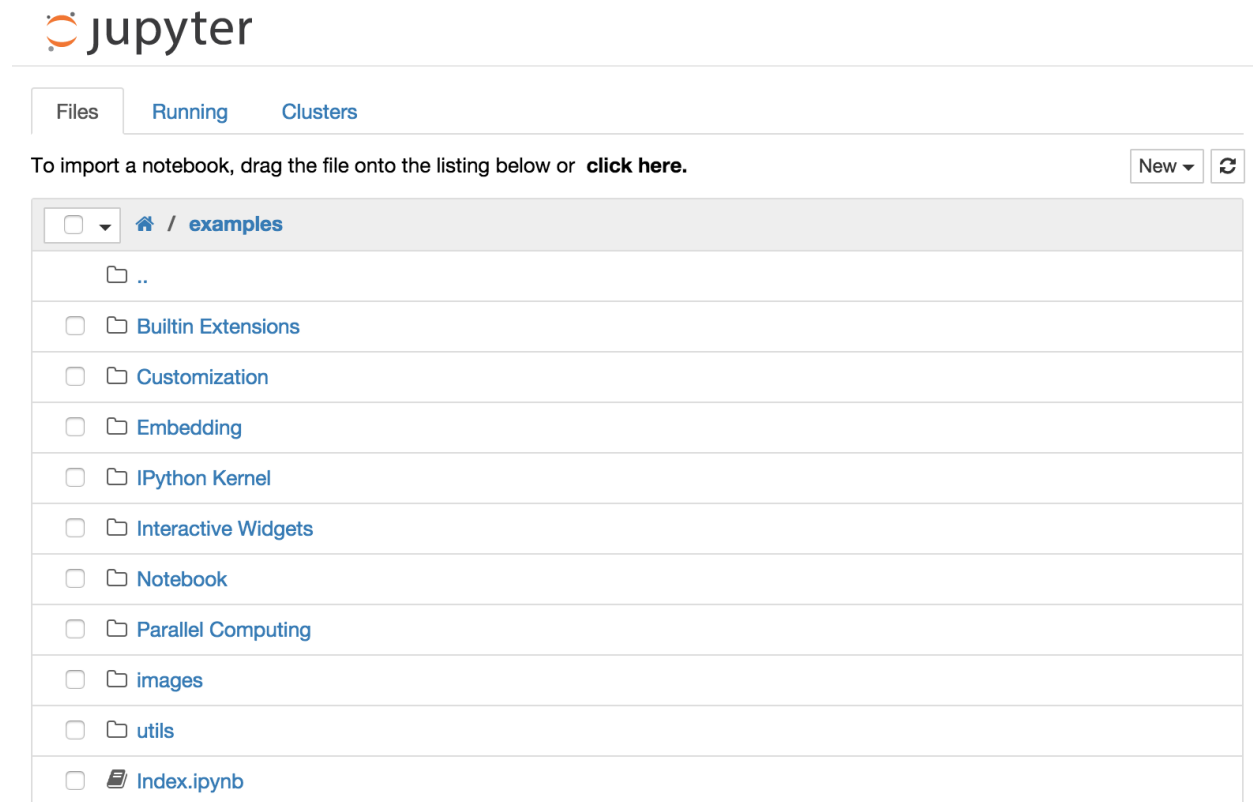
Notebooks can be exported to different static formats including HTML, reStructuredText, LaTeX, PDF, and slide shows ([reveal.js](#)) using Jupyter’s nbconvert utility.

Furthermore, any notebook document available from a **public URL or on GitHub can be shared** via [nbviewer](#). This service loads the notebook document from the URL and renders it as a static web page. The resulting web page may thus be shared with others **without their needing to install the Jupyter Notebook**.

3.2 Notebook Basics


3.2.1 The Notebook dashboard

When you first start the notebook server, your browser will open to the notebook dashboard. The dashboard serves as a home page for the notebook. Its main purpose is to display the notebooks and files in the current directory. For example, here is a screenshot of the dashboard page for the `examples` directory in the Jupyter repository:



The top of the notebook list displays clickable breadcrumbs of the current directory. By clicking on these breadcrumbs or on sub-directories in the notebook list, you can navigate your file system.

To create a new notebook, click on the “New” button at the top of the list and select a kernel from the dropdown (as seen below). Which kernels are listed depend on what’s installed on the server. Some of the kernels in the screenshot below may not exist as an option to you.



Notebooks and files can be uploaded to the current directory by dragging a notebook file onto the notebook list or by the “click here” text above the list.

The notebook list shows green “Running” text and a green notebook icon next to running notebooks (as seen below). Notebooks remain running until you explicitly shut them down; closing the notebook’s page is not sufficient.

 ▾	 / examples
	..
<input type="checkbox"/> 	Builtin Extensions
<input type="checkbox"/> 	Customization
<input type="checkbox"/> 	Embedding
<input type="checkbox"/> 	IPython Kernel
<input type="checkbox"/> 	Interactive Widgets
<input type="checkbox"/> 	Notebook
<input type="checkbox"/> 	Parallel Computing
<input type="checkbox"/> 	images
<input type="checkbox"/> 	utils
<input checked="" type="checkbox"/> 	Index.ipynb Running

To shutdown, delete, duplicate, or rename a notebook check the checkbox next to it and an array of controls will appear at the top of the notebook list (as seen below). You can also use the same operations on directories and files when applicable.



To see all of your running notebooks along with their directories, click on the “Running” tab:



Files Running Clusters

Currently running Jupyter processes



Terminals ▾

There are no terminals running.

Notebooks ▾

examples/Notebook/Index.ipynb	Shutdown
examples/Notebook/What is the IPython Notebook.ipynb	Shutdown
examples/Notebook/Running the Notebook Server.ipynb	Shutdown
examples/Notebook/Notebook Basics.ipynb	Shutdown
examples/Notebook/Running Code.ipynb	Shutdown
examples/Notebook/Working With Markdown Cells.ipynb	Shutdown
examples/Notebook/Custom Keyboard Shortcuts.ipynb	Shutdown
examples/Notebook/JavaScript Notebook Extensions.ipynb	Shutdown
examples/Notebook/Notebook Security.ipynb	Shutdown

This view provides a convenient way to track notebooks that you start as you navigate the file system in a long running notebook server.

3.2.2 Overview of the Notebook UI

If you create a new notebook or open an existing one, you will be taken to the notebook user interface (UI). This UI allows you to run code and author notebook documents interactively. The notebook UI has the following main areas:

- Menu
- Toolbar
- Notebook area and cells

The notebook has an interactive tour of these elements that can be started in the “Help:User Interface Tour” menu item.

3.2.3 Modal editor

Starting with IPython 2.0, the Jupyter Notebook has a modal user interface. This means that the keyboard does different things depending on which mode the Notebook is in. There are two modes: edit mode and command mode.

Edit mode

Edit mode is indicated by a green cell border and a prompt showing in the editor area:

```
In [1]: a = 10
```

When a cell is in edit mode, you can type into the cell, like a normal text editor.

Enter edit mode by pressing `Enter` or using the mouse to click on a cell's editor area.

Command mode

Command mode is indicated by a grey cell border with a blue left margin:

```
In [1]: a = 10
```


When you are in command mode, you are able to edit the notebook as a whole, but not type into individual cells. Most importantly, in command mode, the keyboard is mapped to a set of shortcuts that let you perform notebook and cell actions efficiently. For example, if you are in command mode and you press `C`, you will copy the current cell - no modifier is needed.

Don't try to type into a cell in command mode; unexpected things will happen!

Enter command mode by pressing `Esc` or using the mouse to click *outside* a cell's editor area.

3.2.4 Mouse navigation

All navigation and actions in the Notebook are available using the mouse through the menubar and toolbar, which are both above the main Notebook area:



The first idea of mouse based navigation is that **cells can be selected by clicking on them**. The currently selected cell gets a grey or green border depending on whether the notebook is in edit or command mode. If you click inside a cell's editor area, you will enter edit mode. If you click on the prompt or output area of a cell you will enter command mode.

If you are running this notebook in a live session (not on <http://nbviewer.jupyter.org>) try selecting different cells and going between edit and command mode. Try typing into a cell.

The second idea of mouse based navigation is that **cell actions usually apply to the currently selected cell**. Thus if you want to run the code in a cell, you would select it and click the

button in the toolbar or the “Cell:Run” menu item. Similarly, to copy a cell you would select it and click the

button in the toolbar or the “Edit:Copy” menu item. With this simple pattern, you should be able to do most everything you need with the mouse.

Markdown cells have one other state that can be modified with the mouse. These cells can either be rendered or unrendered. When they are rendered, you will see a nice formatted representation of the cell's contents. When they are unrendered, you will see the raw text source of the cell. To render the selected cell with the mouse, click the

button in the toolbar or the “Cell:Run” menu item. To unrender the selected cell, double click on the cell.

3.2.5 Keyboard Navigation

The modal user interface of the Jupyter Notebook has been optimized for efficient keyboard usage. This is made possible by having two different sets of keyboard shortcuts: one set that is active in edit mode and another in command mode.

The most important keyboard shortcuts are `Enter`, which enters edit mode, and `Esc`, which enters command mode.

In edit mode, most of the keyboard is dedicated to typing into the cell's editor. Thus, in edit mode there are relatively few shortcuts. In command mode, the entire keyboard is available for shortcuts, so there are many more. The `Help->`Keyboard Shortcuts`` dialog lists the available shortcuts.

We recommend learning the command mode shortcuts in the following rough order:

1. Basic navigation: `enter`, `shift-enter`, `up/k`, `down/j`
2. Saving the notebook: `s`
3. Change Cell types: `y`, `m`, `1-6`, `t`
4. Cell creation: `a`, `b`
5. Cell editing: `x`, `c`, `v`, `d`, `z`
6. Kernel operations: `i`, `0` (press twice)

3.3 Running Code

First and foremost, the Jupyter Notebook is an interactive environment for writing and running code. The notebook is capable of running code in a wide range of languages. However, each notebook is associated with a single kernel. This notebook is associated with the IPython kernel, therefore runs Python code.

3.3.1 Code cells allow you to enter and run code

Run a code cell using `Shift-Enter` or pressing the button in the toolbar above:

```
[2]: a = 10
```

```
[3]: print(a)
```

```
10
```

There are two other keyboard shortcuts for running code:

- `Alt-Enter` runs the current cell and inserts a new one below.
- `Ctrl-Enter` run the current cell and enters command mode.

3.3.2 Managing the Kernel

Code is run in a separate process called the Kernel. The Kernel can be interrupted or restarted. Try running the following cell and then hit the

button in the toolbar above.

```
[4]: import time
time.sleep(10)
```

If the Kernel dies you will be prompted to restart it. Here we call the low-level system `libc.time` routine with the wrong argument via `ctypes` to segfault the Python interpreter:

```
[5]: import sys
from ctypes import CDLL
# This will crash a Linux or Mac system
# equivalent calls can be made on Windows

# Uncomment these lines if you would like to see the segfault

# dll = 'dylib' if sys.platform == 'darwin' else 'so.6'
# libc = CDLL("libc.%s" % dll)
# libc.time(-1) # BOOM!!
```

3.3.3 Cell menu

The “Cell” menu has a number of menu items for running code in different ways. These includes:

- Run and Select Below
- Run and Insert Below
- Run All
- Run All Above
- Run All Below

3.3.4 Restarting the kernels

The kernel maintains the state of a notebook’s computations. You can reset this state by restarting the kernel. This is done by clicking on the

in the toolbar above.

3.3.5 `sys.stdout` and `sys.stderr`

The `stdout` and `stderr` streams are displayed as text in the output area.

```
[6]: print("hi, stdout")

hi, stdout
```

```
[7]: from __future__ import print_function
print('hi, stderr', file=sys.stderr)
```

```
hi, stderr
```

3.3.6 Output is asynchronous

All output is displayed asynchronously as it is generated in the Kernel. If you execute the next cell, you will see the output one piece at a time, not all at the end.

```
[8]: import time, sys
     for i in range(8):
         print(i)
         time.sleep(0.5)
```

```
0
1
2
3
4
5
6
7
```

3.3.7 Large outputs

To better handle large outputs, the output area can be collapsed. Run the following cell and then single- or double-click on the active area to the left of the output:

```
[9]: for i in range(50):
     print(i)
```

```
0
1
2
3
4
5
6
7
8
9
10
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
```

(continues on next page)

(continued from previous page)

```
26
27
28
29
30
31
32
33
34
35
36
37
38
39
40
41
42
43
44
45
46
47
48
49
```

Beyond a certain point, output will scroll automatically:

```
[10]: for i in range(500):
      print(2**i - 1)
```

```
0
1
3
7
15
31
63
127
255
511
1023
2047
4095
8191
16383
32767
65535
131071
262143
524287
1048575
2097151
4194303
8388607
16777215
33554431
67108863
```

(continues on next page)

(continued from previous page)

```
134217727
268435455
536870911
1073741823
2147483647
4294967295
8589934591
17179869183
34359738367
68719476735
137438953471
274877906943
549755813887
1099511627775
2199023255551
4398046511103
8796093022207
17592186044415
35184372088831
70368744177663
140737488355327
281474976710655
562949953421311
1125899906842623
2251799813685247
4503599627370495
9007199254740991
18014398509481983
36028797018963967
72057594037927935
144115188075855871
288230376151711743
576460752303423487
1152921504606846975
2305843009213693951
4611686018427387903
9223372036854775807
18446744073709551615
36893488147419103231
73786976294838206463
147573952589676412927
295147905179352825855
590295810358705651711
1180591620717411303423
2361183241434822606847
4722366482869645213695
9444732965739290427391
18889465931478580854783
37778931862957161709567
75557863725914323419135
151115727451828646838271
302231454903657293676543
604462909807314587353087
1208925819614629174706175
2417851639229258349412351
4835703278458516698824703
9671406556917033397649407
```

(continues on next page)

(continued from previous page)

```
19342813113834066795298815
38685626227668133590597631
77371252455336267181195263
154742504910672534362390527
309485009821345068724781055
618970019642690137449562111
1237940039285380274899124223
2475880078570760549798248447
4951760157141521099596496895
9903520314283042199192993791
19807040628566084398385987583
39614081257132168796771975167
79228162514264337593543950335
158456325028528675187087900671
316912650057057350374175801343
633825300114114700748351602687
1267650600228229401496703205375
2535301200456458802993406410751
5070602400912917605986812821503
10141204801825835211973625643007
20282409603651670423947251286015
40564819207303340847894502572031
81129638414606681695789005144063
162259276829213363391578010288127
324518553658426726783156020576255
649037107316853453566312041152511
1298074214633706907132624082305023
2596148429267413814265248164610047
5192296858534827628530496329220095
10384593717069655257060992658440191
20769187434139310514121985316880383
41538374868278621028243970633760767
83076749736557242056487941267521535
166153499473114484112975882535043071
332306998946228968225951765070086143
664613997892457936451903530140172287
1329227995784915872903807060280344575
2658455991569831745807614120560689151
5316911983139663491615228241121378303
10633823966279326983230456482242756607
21267647932558653966460912964485513215
42535295865117307932921825928971026431
85070591730234615865843651857942052863
170141183460469231731687303715884105727
340282366920938463463374607431768211455
680564733841876926926749214863536422911
1361129467683753853853498429727072845823
2722258935367507707706996859454145691647
5444517870735015415413993718908291383295
10889035741470030830827987437816582766591
21778071482940061661655974875633165533183
43556142965880123323311949751266331066367
87112285931760246646623899502532662132735
174224571863520493293247799005065324265471
348449143727040986586495598010130648530943
696898287454081973172991196020261297061887
1393796574908163946345982392040522594123775
```

(continues on next page)

(continued from previous page)

```

2787593149816327892691964784081045188247551
5575186299632655785383929568162090376495103
11150372599265311570767859136324180752990207
22300745198530623141535718272648361505980415
44601490397061246283071436545296723011960831
89202980794122492566142873090593446023921663
178405961588244985132285746181186892047843327
356811923176489970264571492362373784095686655
713623846352979940529142984724747568191373311
1427247692705959881058285969449495136382746623
2854495385411919762116571938898990272765493247
5708990770823839524233143877797980545530986495
11417981541647679048466287755595961091061972991
22835963083295358096932575511191922182123945983
45671926166590716193865151022383844364247891967
91343852333181432387730302044767688728495783935
182687704666362864775460604089535377456991567871
365375409332725729550921208179070754913983135743
730750818665451459101842416358141509827966271487
1461501637330902918203684832716283019655932542975
2923003274661805836407369665432566039311865085951
5846006549323611672814739330865132078623730171903
11692013098647223345629478661730264157247460343807
23384026197294446691258957323460528314494920687615
46768052394588893382517914646921056628989841375231
93536104789177786765035829293842113257979682750463
187072209578355573530071658587684226515959365500927
374144419156711147060143317175368453031918731001855
748288838313422294120286634350736906063837462003711
1496577676626844588240573268701473812127674924007423
2993155353253689176481146537402947624255349848014847
5986310706507378352962293074805895248510699696029695
11972621413014756705924586149611790497021399392059391
23945242826029513411849172299223580994042798784118783
47890485652059026823698344598447161988085597568237567
95780971304118053647396689196894323976171195136475135
191561942608236107294793378393788647952342390272950271
383123885216472214589586756787577295904684780545900543
766247770432944429179173513575154591809369561091801087
1532495540865888858358347027150309183618739122183602175
3064991081731777716716694054300618367237478244367204351
6129982163463555433433388108601236734474956488734408703
12259964326927110866866776217202473468949912977468817407
24519928653854221733733552434404946937899825954937634815
49039857307708443467467104868809893875799651909875269631
98079714615416886934934209737619787751599303819750539263
196159429230833773869868419475239575503198607639501078527
392318858461667547739736838950479151006397215279002157055
784637716923335095479473677900958302012794430558004314111
1569275433846670190958947355801916604025588861116008628223
3138550867693340381917894711603833208051177722232017256447
6277101735386680763835789423207666416102355444464034512895
12554203470773361527671578846415332832204710888928069025791
25108406941546723055343157692830665664409421777856138051583
50216813883093446110686315385661331328818843555712276103167
100433627766186892221372630771322662657637687111424552206335
200867255532373784442745261542645325315275374222849104412671

```

(continues on next page)

(continued from previous page)

```
401734511064747568885490523085290650630550748445698208825343
803469022129495137770981046170581301261101496891396417650687
160693804425899027554196209234116260252202993782792835301375
3213876088517980551083924184682325205044405987565585670602751
6427752177035961102167848369364650410088811975131171341205503
12855504354071922204335696738729300820177623950262342682411007
25711008708143844408671393477458601640355247900524685364822015
51422017416287688817342786954917203280710495801049370729644031
102844034832575377634685573909834406561420991602098741459288063
205688069665150755269371147819668813122841983204197482918576127
411376139330301510538742295639337626245683966408394965837152255
822752278660603021077484591278675252491367932816789931674304511
1645504557321206042154969182557350504982735865633579863348609023
3291009114642412084309938365114701009965471731267159726697218047
6582018229284824168619876730229402019930943462534319453394436095
13164036458569648337239753460458804039861886925068638906788872191
26328072917139296674479506920917608079723773850137277813577744383
52656145834278593348959013841835216159447547700274555627155488767
105312291668557186697918027683670432318895095400549111254310977535
210624583337114373395836055367340864637790190801098222508621955071
421249166674228746791672110734681729275580381602196445017243910143
842498333348457493583344221469363458551160763204392890034487820287
1684996666696914987166688442938726917102321526408785780068975640575
3369993333393829974333376885877453834204643052817571560137951281151
6739986666787659948666753771754907668409286105635143120275902562303
13479973333575319897333507543509815336818572211270286240551805124607
26959946667150639794667015087019630673637144422540572481103610249215
53919893334301279589334030174039261347274288845081144962207220498431
107839786668602559178668060348078522694548577690162289924414440996863
215679573337205118357336120696157045389097155380324579848828881993727
431359146674410236714672241392314090778194310760649159697657763987455
862718293348820473429344482784628181556388621521298319395315527974911
1725436586697640946858688965569256363112777243042596638790631055949823
345087317339528189371737793113851272622554486085193277581262111899647
6901746346790563787434755862277025452451108972170386555162524223799295
13803492693581127574869511724554050904902217944340773110325048447598591
27606985387162255149739023449108101809804435888681546220650096895197183
55213970774324510299478046898216203619608871777363092441300193790394367
110427941548649020598956093796432407239217743554726184882600387580788735
220855883097298041197912187592864814478435487109452369765200775161577471
441711766194596082395824375185729628956870974218904739530401550323154943
883423532389192164791648750371459257913741948437809479060803100646309887
1766847064778384329583297500742918515827483896875618958121606201292619775
3533694129556768659166595001485837031654967793751237916243212402585239551
7067388259113537318333190002971674063309935587502475832486424805170479103
14134776518227074636666380005943348126619871175004951664972849610340958207
28269553036454149273332760011886696253239742350009903329945699220681916415
56539106072908298546665520023773392506479484700019806659891398441363832831
113078212145816597093331040047546785012958969400039613319782796882727665663
226156424291633194186662080095093570025917938800079226639565593765455331327
452312848583266388373324160190187140051835877600158453279131187530910662655
904625697166532776746648320380374280103671755200316906558262375061821325311
1809251394333065553493296640760748560207343510400633813116524750123642650623
3618502788666131106986593281521497120414687020801267626233049500247285301247
7237005577332262213973186563042994240829374041602535252466099000494570602495
14474011154664524427946373126085988481658748083205070504932198000989141204991
28948022309329048855892746252171976963317496166410141009864396001978282409983
```

(continues on next page)

(continued from previous page)

```

57896044618658097711785492504343953926634992332820282019728792003956564819967
115792089237316195423570985008687907853269984665640564039457584007913129639935
231584178474632390847141970017375815706539969331281128078915168015826259279871
463168356949264781694283940034751631413079938662562256157830336031652518559743
926336713898529563388567880069503262826159877325124512315660672063305037119487
1852673427797059126777135760139006525652319754650249024631321344126610074238975
3705346855594118253554271520278013051304639509300498049262642688253220148477951
7410693711188236507108543040556026102609279018600996098525285376506440296955903
14821387422376473014217086081112052205218558037201992197050570753012880593911807
29642774844752946028434172162224104410437116074403984394101141506025761187823615
59285549689505892056868344324448208820874232148807968788202283012051522375647231
118571099379011784113736688648896417641748464297615937576404566024103044751294463
237142198758023568227473377297792835283496928595231875152809132048206089502588927
474284397516047136454946754595585670566993857190463750305618264096412179005177855
948568795032094272909893509191171341133987714380927500611236528192824358010355711
1897137590064188545819787018382342682267975428761855001222473056385648716020711423
3794275180128377091639574036764685364535950857523710002444946112771297432041422847
7588550360256754183279148073529370729071901715047420004889892225542594864082845695
15177100720513508366558296147058741458143803430094840009779784451085189728165691391
30354201441027016733116592294117482916287606860189680019559568902170379456331382783
60708402882054033466233184588234965832575213720379360039119137804340758912662765567
121416805764108066932466369176469931665150427440758720078238275608681517825325531135
242833611528216133864932738352939863330300854881517440156476551217363035650651062271
485667223056432267729865476705879726660601709763034880312953102434726071301302124543
971334446112864535459730953411759453321203419526069760625906204869452142602604249087
1942668892225729070919461906823518906642406839052139521251812409738904285205208498175
3885337784451458141838923813647037813284813678104279042503624819477808570410416996351
7770675568902916283677847627294075626569627356208558085007249638955617140820833992703
15541351137805832567355695254588151253139254712417116170014499277911234281641667985407
31082702275611665134711390509176302506278509424834232340028998555822468563283335970815
62165404551223330269422781018352605012557018849668464680057997111644937126566671941631
124330809102446660538845562036705210025114037699336929360115994223289874253133343883263
248661618204893321077691124073410420050228075398673858720231988446579748506266687766527
497323236409786642155382248146820840100456150797347717440463976893159497012533375533055
994646472819573284310764496293641680200912301594695434880927953786318994025066751066111
1989292945639146568621528992587283360401824603189390869761855907572637988050133502132223
3978585891278293137243057985174566720803649206378781739523711815145275976100267004264447
7957171782556586274486115970349133441607298412757563479047423630290551952200534008528895
15914343565113172548972231940698266883214596825515126958094847260581103904401068017057791
31828687130226345097944463881396533766429193651030253916189694521162207808802136034115583
63657374260452690195888927762793067532858387302060507832379389042324415617604272068231167
127314748520905380391777855525586135065716774604121015664758778084648831235208544136462335
254629497041810760783555711051172270131433549208242031329517556169297662470417088272924671
509258994083621521567111422102344540262867098416484062659035112338595324940834176545849343
1018517988167243043134222844204689080525734196832968125318070224677190649881668353091698687
203703597633448608626844568840937816105146839366593625063614044935438129976336706183397375
4074071952668972172536891376818756322102936787331872501272280898708762599526673412366794751
8148143905337944345073782753637512644205873574663745002544561797417525199053346824733589503
16296287810675888690147565507275025288411747149327490005089123594835050398106693649467179007
32592575621351777380295131014550050576823494298654980010178247189670100796213387298934358015
65185151242703554760590262029100101153646988597309960020356494379340201592426774597868716031
130370302485407109521180524058200202307293977194619920040712988758680403184853549195737432063
260740604970814219042361048116400404614587954389239840081425977517360806369707098391474864127
521481209941628438084722096232800809229175908778479680162851955034721612739414196782949728255
1042962419883256876169444192465601618458351817556959360325703910069443225478828393565899456511
2085924839766513752338888384931203236916703635113918720651407820138886450957656787131798913023
41718496795330275046777676986240647383340727022783744130281564027772901915313574263597826047

```

(continues on next page)

(continued from previous page)

```
8343699359066055009355553539724812947666814540455674882605631280555545803830627148527195652095
16687398718132110018711107079449625895333629080911349765211262561111091607661254297054391304191
33374797436264220037422214158899251790667258161822699530422525122222183215322508594108782608383
6674959487252844007484442831779850358133451632364539906084505024444366430645017188217565216767
13349918974505688014968885663559700716266903264729079812169010048888732861290034376435130433535
266998379490113760299377713271194014325338065294581596243380200977777465722580068752870260867071
533996758980227520598755426542388028650676130589163192486760401955554931445160137505740521734143
1067993517960455041197510853084776057301352261178326384973520803911109862890320275011481043468287
2135987035920910082395021706169552114602704522356652769947041607822219725780640550022962086936575
4271974071841820164790043412339104229205409044713305539894083215644439451561281100045924173873151
8543948143683640329580086824678208458410818089426611079788166431288878903122562200091848347746303
1708789628736728065916017364935641691682163617885322215957633286257775780624512440018369695492607
34175792574734561318320347298712833833643272357706444319152665725155515612490248800367393390985215
68351585149469122636640694597425667667286544715412888638305331450311031224980497600734786781970431
136703170298938245273281389194851335334573089430825777276610662900622062449960995201469573563940863
273406340597876490546562778389702670669146178861651554553221325801244124899921990402939147127881727
546812681195752981093125556779405341338292357723303109106442651602488249799843980805878294255763455
1093625362391505962186251113558810682676584715446606218212885303204976499599687961611756588511526911
2187250724783011924372502227117621365353169430893212436425770606409952999199375923223513177023053823
4374501449566023848745004454235242730706338861786424872851541212819905998398751846447026354046107647
8749002899132047697490008908470485461412677723572849745703082425639811996797503692894052708092215295
17498005798264095394980017816940970922825355447145699491406164851279623993595007385788105416184430592
34996011596528190789960035633881941845650710894291398982812329702559247987190014771576210832368861183
69992023193056381579920071267763883691301421788582797965624659405118495974380029543152421664737722367
13998404638611276315984014253552776738260284357716559593124931881023699194876005908630484332947544473
27996809277222552631968028507105553476520568715433119186249863762047398389752011817260968665895088947
55993618554445105263936057014211106953041137430866238372499727524094796779504023634521937331790177894
11198723710889021052787211402842221390608227486173247674499945504818959355900804726904387466358035578
22397447421778042105574422805684442781216454972346495348999891009637918711801609453808774932716071157
44794894843556084211148845611368885562432909944692990697999782019275837423603218907617549865432142319
89589789687112168422297691222737771124865819889385981395999564038551674847206437815235099730864284631
17917957937422433684459538244547554224973163977877196279199912807710334969441287563047019946172856920
3583591587484867368919076489095108449946327955754392558399825615420669938882575126094039892345713852
71671831749689734737838152978190216899892655911508785116799651230841339877765150252188079784691427704
14334366349937946947567630595638043379978531182301757023359930246168267975553030050437615956938285540
28668732699875893895135261191276086759957062364603514046719860492336535951106060100875231913876571083
57337465399751787790270522382552173519914124729207028093439720984673071902212120201750463827753142163
11467493079950357558054104476510434703982824945841405618687944196934614380442424040350092765550628432
22934986159900715116108208953020869407965649891682811237375888393869228760884848080700185531101256863
45869972319801430232216417906041738815931299783365622474751776787738457521769696161400371062202513733
91739944639602860464432835812083477631862599566731244949503553575476915043539392322800742124405027462
18347988927920572092886567162416695526372519913346248989900710715095383008707878464560148424881005492
36695977855841144185773134324833391052745039826692497979801421430190766017415756929120296849762010984
73391955711682288371546268649666782105490079653384995959602842860381532034831513858240593699524021963
14678391142336457674309253729933356421098015930676999191920568572076306406966302771648118739904804392
29356782284672915348618507459866712842196031861353998383841137144152612813932605543296237479809608787
58713564569345830697237014919733425684392063722707996767682274288305225627865211086592474959619217579
11742712913869166139447402983946685136878412744541599353536454857661045125573042217318494991923843519
23485425827738332278894805967893370273756825489083198707072909715322090251146084434636989983847687030
46970851655476664557789611935786740547513650978166397414145819430644180502292168869273979967695374060
93941703310953329115579223871573481095027301956332794828291638861288361004584337738547959935390748123
1878834066219066582311584477431469621900546039126655896565832772257672200916867547709591987078149624
3757668132438133164623168954862939243801092078253311793131665554451534440183373509541918397415629924
75153362648762663292463379097258784876021841565066235862633311089030688803667470190838367948312598497
15030672529752532658492675819451756975204368313013247172526662217806137760733494038167673589662519699
30061345059505065316985351638903513950408736626026494345053324435612275521466988076335347179325039398
60122690119010130633970703277807027900817473252052988690106648871224551042933976152670694358650078797
```

(continues on next page)

(continued from previous page)

```

1202453802380202612679414065556140558016349465041059773802132977424491020858679523053413887173001575
2404907604760405225358828131112281116032698930082119547604265954848982041717359046106827774346003151
480981520952081045071765626224562232065397860164239095208531909697964083434718092213655548692006303
9619630419041620901435312524449124464130795720328478190417063819395928166869436184427311097384012607
1923926083808324180287062504889824892826159144065695638083412763879185633373887236885462219476802521
3847852167616648360574125009779649785652318288131391276166825527758371266747774473770924438953605043
769570433523329672114825001955929957130463657626278255233651055516742533495548947541848877907210086
1539140867046659344229650003911859914260927315252556510466730211103348506699109789508369775581442017
3078281734093318688459300007823719828521854630505113020933460422206697013398219579016739551162884034
6156563468186637376918600015647439657043709261010226041866920844413394026796439158033479102325768068
1231312693637327475383720003129487931408741852202045208373384168882678805359287831606695820465153613
2462625387274654950767440006258975862817483704404090416746768337765357610718575663213391640930307227
4925250774549309901534880012517951725634967408808180833493536675530715221437151326426783281860614455
9850501549098619803069760025035903451269934817616361666987073351061430442874302652853566563721228910
197010030981972396061395200500718069025398696352327233397414670212286088574860530570713312744245782
3940200619639447921227904010014361380507973927046544666794829340424572177149721061141426625488491564
7880401239278895842455808020028722761015947854093089333589658680849144354299442122282853250976983128
1576080247855779168491161604005744552203189570818617866717931736169828870859888424456570650195396625
3152160495711558336982323208011489104406379141637235733435863472339657741719776848913141300390793251
6304320991423116673964646416022978208812758283274471466871726944679315483439553697826282600781586502
1260864198284623334792929283204595641762551656654894293374345388935863096687910739565256520156317300
2521728396569246669585858566409191283525103313309788586748690777871726193375821479130513040312634601
504345679313849333917171713281838256705020626619577173497381555743452386751642958261026080625269202
1008691358627698667834343426563676513410041325323915434699476311148690477350328591652205216125053840
2017382717255397335668686853127353026820082650647830869398952622297380954700657183304410432250107680
4034765434510794671337373706254706053640165301295661738797905244594761909401314366608820864500215361
8069530869021589342674747412509412107280330602591323477595810489189523818802628733217641729000430723
1613906173804317868534949482501882421456066120518264695519162097837904763760525746643528345800086144
3227812347608635737069898965003764842912132241036529391038324195675809527521051493287056691600172289
6455624695217271474139797930007529685824264482073058782076648391351619055042102986574113383200344578
1291124939043454294827959586001505937164852896414611756415329678270323811008420597314822676640068915
2582249878086908589655919172003011874329705792829223512830659356540647622016841194629645353280137831
5164499756173817179311838344006023748659411585658447025661318713081295244033682389259290706560275662
103289951234763435862367668801204749731882317131689405132263742616259048806736477851858141312055132
2065799902469526871724735337602409499463764634263378810264527485232518097613472955703716282624110265
4131599804939053743449470675204818998927529268526757620529054970465036195226945911407432565248220530
8263199609878107486898941350409637997855058537053515241058109940930072390453891822814865130496441060
1652639921975621497379788270081927599571011707410703048211621988186014478090778364562973026099288212
3305279843951242994759576540163855199142023414821406096423243976372028956181556729125946052198576424
6610559687902485989519153080327710398284046829642812192846487952744057912363113458251892104397152848
1322111937580497197903830616065542079656809365928562438569297590548811582472622691650378420879430569
2644223875160994395807661232131084159313618731857124877138595181097623164945245383300756841758861139
5288447750321988791615322464262168318627237463714249754277190362195246329890490766601513683517722278
1057689550064397758323064492852433663725447492742849950855438072439049265978098153320302736703544455
2115379100128795516646128985704867327450894985485699901710876144878098531956196306640605473407088911
4230758200257591033292257971409734654901789970971399803421752289756197063912392613281210946814177823
8461516400515182066584515942819469309803579941942799606843504579512394127824785226562421893628355646
1692303280103036413316903188563893861960715988388559921368700915902478825564957045312484378725671129
3384606560206072826633806377127787723921431976777119842737401831804957651129914090624968757451342258
6769213120412145653267612754255575447842863953554239685474803663609915302259828181249937514902684516
1353842624082429130653522550851115089568572790710847937094960732721983060451965636249987502980536903
2707685248164858261307045101702230179137145581421695874189921465443966120903931272499975005961073806
5415370496329716522614090203404460358274291162843391748379842930887932241807862544999950011922147613
1083074099265943304522818040680892071654858232568678349675968586177586448361572508999990002384429522
2166148198531886609045636081361784143309716465137356699351937172355172896723145017999980004768859045
433229639706377321809127216272356828661943293027471339870387434471034579344629003599960009537718090
8664592794127546436182544325447136573238865860549426797407748689420691586892580071999920019075436181

```

(continues on next page)

(continued from previous page)

```
17329185588255092872365088650894273146477731721098853594815497378841383173785160143999840038150872363
3465837117651018574473017730178854629295546344219770718963099475768276634757032028799968007630174472
6931674235302037148946035460357709258591092688439541437926198951536553269514064057599936015260348945
1386334847060407429789207092071541851718218537687908287585239790307310653902812811519987203052069789
2772669694120814859578414184143083703436437075375816575170479580614621307805625623039974406104139578
5545339388241629719156828368286167406872874150751633150340959161229242615611251246079948812208279156
1109067877648325943831365673657233481374574830150326630068191832245848523122250249215989762441655831
2218135755296651887662731347314466962749149660300653260136383664491697046244500498431979524883311662
4436271510593303775325462694628933925498299320601306520272767328983394092489000996863959049766623324
8872543021186607550650925389257867850996598641202613040545534657966788184978001993727918099533246649
1774508604237321510130185077851573570199319728240522608109106931593357636995600398745583619906649329
3549017208474643020260370155703147140398639456481045216218213863186715273991200797491167239813298659
7098034416949286040520740311406294280797278912962090432436427726373430547982401594982334479626597319
1419606883389857208104148062281258856159455782592418086487285545274686109596480318996466895925319463
2839213766779714416208296124562517712318911565184836172974571090549372219192960637992933791850638927
5678427533559428832416592249125035424637823130369672345949142181098744438385921275985867583701277855
1135685506711885766483318449825007084927564626073934469189828436219748887677184255197173516740255571
2271371013423771532966636899650014169855129252147868938379656872439497775354368510394347033480511142
4542742026847543065933273799300028339710258504295737876759313744878995550708737020788694066961022284
9085484053695086131866547598600056679420517008591475753518627489757991101417474041577388133922044569
1817096810739017226373309519720011335884103401718295150703725497951598220283494808315477626784408913
3634193621478034452746619039440022671768206803436590301407450995903196440566989616630955253568817827
7268387242956068905493238078880045343536413606873180602814901991806392881133979233261910507137635655
1453677448591213781098647615776009068707282721374636120562980398361278576226795846652382101427527131
2907354897182427562197295231552018137414565442749272241125960796722557152453591693304764202855054262
5814709794364855124394590463104036274829130885498544482251921593445114304907183386609528405710108524
1162941958872971024878918092620807254965826177099708896450384318689022860981436677321905681142021704
2325883917745942049757836185241614509931652354199417792900768637378045721962873354643811362284043409
4651767835491884099515672370483229019863304708398835585801537274756091443925746709287622724568086819
9303535670983768199031344740966458039726609416797671171603074549512182887851493418575245449136173639
1860707134196753639806268948193291607945321883359534234320614909902436577570298683715049089827234727
3721414268393507279612537896386583215890643766719068468641229819804873155140597367430098179654469455
7442828536787014559225075792773166431781287533438136937282459639609746310281194734860196359308938911
1488565707357402911845015158554633286356257506687627387456491927921949262056238946972039271861787782
2977131414714805823690030317109266572712515013375254774912983855843898524112477893944078543723575564
5954262829429611647380060634218533145425030026750509549825967711687797048224955787888157087447151129
1190852565885922329476012126843706629085006005350101909965193542337559409644991157577631417489430225
2381705131771844658952024253687413258170012010700203819930387084675118819289982315155262834978860451
4763410263543689317904048507374826516340024021400407639860774169350237638579964630310525669957720903
9526820527087378635808097014749653032680048042800815279721548338700475277159929260621051339915441806
1905364105417475727161619402949930606536009608560163055944309667740095055431985852124210267983088361
3810728210834951454323238805899861213072019217120326111888619335480190110863971704248420535966176722
762145642166990290864647761179972242614403843424065222377238670960380221727943408496841071932353445
1524291284333980581729295522359944485228807686848130444755447734192076044345588681699368214386470689
3048582568667961163458591044719888970457615373696260889510895468384152088691177363398736428772941378
6097165137335922326917182089439777940915230747392521779021790936768304177382354726797472857545882756
1219433027467184465383436417887955588183046149478504355804358187353660835476470945359494571509176551
2438866054934368930766872835775911176366092298957008711608716374707321670952941890718989143018353102
4877732109868737861533745671551822352732184597914017423217432749414643341905883781437978286036706204
9755464219737475723067491343103644705464369195828034846434865498829286683811767562875956572073412409
1951092843947495144613498268620728941092873839165606969286973099765857336762353512575191314414682481
3902185687894990289226996537241457882185747678331213938573946199531714673524707025150382628829364963
7804371375789980578453993074482915764371495356662427877147892399063429347049414050300765257658729927
1560874275157996115690798614896583152874299071332485575429578479812685869409882810060153051531745985
3121748550315992231381597229793166305748598142664971150859156959625371738819765620120306103063491971
6243497100631984462763194459586332611497196285329942301718313919250743477639531240240612206126983942
1248699420126396892552638891917266522299439257065988460343662783850148695527906248048122441225396788
```

(continues on next page)

(continued from previous page)

```

2497398840252793785105277783834533044598878514131976920687325567700297391055812496096244882450793576
4994797680505587570210555567669066089197757028263953841374651135400594782111624992192489764901587153
9989595361011175140421111135338132178395514056527907682749302270801189564223249984384979529803174307
1997919072202235028084222227067626435679102811305581536549860454160237912844649996876995905960634861
3995838144404470056168444454135252871358205622611163073099720908320475825689299993753991811921269723
7991676288808940112336888908270505742716411245222326146199441816640951651378599987507983623842539446
1598335257761788022467377781654101148543282249044465229239888363328190330275719997501596724768507889
3196670515523576044934755563308202297086564498088930458479776726656380660551439995003193449537015778
6393341031047152089869511126616404594173128996177860916959553453312761321102879990006386899074031556
1278668206209430417973902225323280918834625799235572183391910690662552264220575998001277379814806311
2557336412418860835947804450646561837669251598471144366783821381325104528441151996002554759629612622
5114672824837721671895608901293123675338503196942288733567642762650209056882303992005109519259225245
1022934564967544334379121780258624735067700639388457746713528552530041811376460798401021903851845049
2045869129935088668758243560517249470135401278776915493427057105060083622752921596802043807703690098
4091738259870177337516487121034498940270802557553830986854114210120167245505843193604087615407380196
8183476519740354675032974242068997880541605115107661973708228420240334491011686387208175230814760392
1636695303948070935006594848413799576108321023021532394741645684048066898202337277441635046162952078

```

3.4 Markdown Cells

Text can be added to Jupyter Notebooks using Markdown cells. You can change the cell type to Markdown by using the `Cell` menu, the toolbar, or the key shortcut `m`. Markdown is a popular markup language that is a superset of HTML. Its specification can be found here:

<https://daringfireball.net/projects/markdown/>

3.4.1 Markdown basics

You can make text *italic* or **bold** by surrounding a block of text with a single or double `*` respectively

You can build nested itemized or enumerated lists:

- One
 - Sublist
 - * This
- Sublist - That - The other thing
- Two
- Sublist
- Three
- Sublist

Now another list:

1. Here we go
 1. Sublist
 2. Sublist
2. There we go
3. Now this

You can add horizontal rules:

Here is a blockquote:

Beautiful is better than ugly. Explicit is better than implicit. Simple is better than complex. Complex is better than complicated. Flat is better than nested. Sparse is better than dense. Readability counts. Special cases aren't special enough to break the rules. Although practicality beats purity. Errors should never pass silently. Unless explicitly silenced. In the face of ambiguity, refuse the temptation to guess. There should be one— and preferably only one —obvious way to do it. Although that way may not be obvious at first unless you're Dutch. Now is better than never. Although never is often better than *right* now. If the implementation is hard to explain, it's a bad idea. If the implementation is easy to explain, it may be a good idea. Namespaces are one honking great idea – let's do more of those!

And shorthand for links:

[Jupyter's website](#)

You can use backslash to generate literal characters which would otherwise have special meaning in the Markdown syntax.

```
\*literal asterisks\*
*literal asterisks*
```

Use double backslash to generate the literal \$ symbol.

3.4.2 Headings

You can add headings by starting a line with one (or multiple) # followed by a space, as in the following example:

```
# Heading 1
# Heading 2
## Heading 2.1
## Heading 2.2
```

3.4.3 Embedded code

You can embed code meant for illustration instead of execution in Python:

```
def f(x):
    """a docstring"""
    return x**2
```

or other languages:

```
for (i=0; i<n; i++) {
    printf("hello %d\n", i);
    x += 4;
}
```

3.4.4 LaTeX equations

Courtesy of MathJax, you can include mathematical expressions both inline: $e^{i\pi} + 1 = 0$ and displayed:

$$e^x = \sum_{i=0}^{\infty} \frac{1}{i!} x^i \quad (3.1)$$

Inline expressions can be added by surrounding the latex code with `$`:

```
$e^{i\pi} + 1 = 0$
```

Expressions on their own line are surrounded by `\begin{equation}` and `\end{equation}`:

```
\begin{equation}
e^x = \sum_{i=0}^{\infty} \frac{1}{i!} x^i
\end{equation}
```

3.4.5 GitHub flavored markdown

The Notebook webapp supports Github flavored markdown meaning that you can use triple backticks for code blocks:

```
```python
print "Hello World"
```

```javascript
console.log("Hello World")
```
```

Gives:

```
print "Hello World"
```

```
console.log("Hello World")
```

And a table like this:

```
This	is
a	table
```

A nice HTML Table:

| | |
|------|-------|
| This | is |
| a | table |

3.4.6 General HTML

Because Markdown is a superset of HTML you can even add things like HTML tables:

Header 1

Header 2

row 1, cell 1

row 1, cell 2

row 2, cell 1

row 2, cell 2

3.4.7 Local files

If you have local files in your Notebook directory, you can refer to these files in Markdown cells directly:

```
[subdirectory/]<filename>
```

For example, in the images folder, we have the Python logo:

```

```

and a video with the HTML5 video tag:

```
<video controls src="../../../images/animation.m4v">animation</video>
```

animation

These do not embed the data into the notebook file, and require that the files exist when you are viewing the notebook.

Security of local files

Note that this means that the Jupyter notebook server also acts as a generic file server for files inside the same tree as your notebooks. Access is not granted outside the notebook folder so you have strict control over what files are visible, but for this reason it is highly recommended that you do not run the notebook server with a notebook directory at a high level in your filesystem (e.g. your home directory).

When you run the notebook in a password-protected manner, local file access is restricted to authenticated users unless read-only views are active.

Markdown attachments

Since Jupyter notebook version 5.0, in addition to referencing external file you can attach a file to a markdown cell. To do so drag the file from in a markdown cell while editing it:



Files are stored in cell metadata and will be automatically scrubbed at save-time if not referenced. You can recognize attached images from other files by their url that starts with `attachment:`. For the image above:

```
! [pycon-logo.jpg] (attachment:pycon-logo.jpg)
```

Keep in mind that attached files will increase the size of your notebook.

You can manually edit the attachment by using the `View > Cell Toolbar > Attachment` menu, but you should not need to.

3.5 Keyboard Shortcut Customization

Starting with Jupyter Notebook 5.0, you can customize the `command` mode shortcuts from within the Notebook Application itself.

Head to the `Help` menu and select the `Edit keyboard Shortcuts` item. A dialog will guide you through the process of adding custom keyboard shortcuts.

Keyboard shortcut set from within the Notebook Application will be persisted to your configuration file. A single action may have several shortcuts attached to it.

3.6 Keyboard Shortcut Customization (Pre Notebook 5.0)

Starting with IPython 2.0 keyboard shortcuts in command and edit mode are fully customizable. These customizations are made using the Jupyter JavaScript API. Here is an example that makes the `r` key available for running a cell:

```
[ ]: %%javascript

Jupyter.keyboard_manager.command_shortcuts.add_shortcut('r', {
  help : 'run cell',
  help_index : 'zz',
  handler : function (event) {
    IPython.notebook.execute_cell();
    return false;
  }}
);
```

“By default the keypress `r`, while in command mode, changes the type of the selected cell to `raw`. This shortcut is overridden by the code in the previous cell, and thus the action no longer be available via the keypress `r`.”

There are a couple of points to mention about this API:

- The `help_index` field is used to sort the shortcuts in the Keyboard Shortcuts help dialog. It defaults to `zz`.
- When a handler returns `false` it indicates that the event should stop propagating and the default action should not be performed. For further details about the `event` object or event handling, see the jQuery docs.
- If you don't need a `help` or `help_index` field, you can simply pass a function as the second argument to `add_shortcut`.

```
[ ]: %%javascript

Jupyter.keyboard_manager.command_shortcuts.add_shortcut('r', function (event) {
  IPython.notebook.execute_cell();
  return false;
});
```

Likewise, to remove a shortcut, use `remove_shortcut`:

```
[ ]: %%javascript

Jupyter.keyboard_manager.command_shortcuts.remove_shortcut('r');
```

If you want your keyboard shortcuts to be active for all of your notebooks, put the above API calls into your `custom.js` file.

Of course we provide name for majority of existing action so that you do not have to re-write everything, here is for example how to bind `r` back to it's initial behavior:

```
[ ]: %%javascript

Jupyter.keyboard_manager.command_shortcuts.add_shortcut('r', 'jupyter-notebook:change-
↩cell-to-raw');
```

3.7 Embracing web standards

One of the main reasons why we developed the current notebook web application was to embrace the web technology.

By being a pure web application using HTML, JavaScript, and CSS, the Notebook can get all the web technology improvement for free. Thus, as browser support for different media extend, the notebook web app should be able to be compatible without modification.

This is also true with performance of the User Interface as the speed of JavaScript VM increases.

The other advantage of using only web technology is that the code of the interface is fully accessible to the end user and is modifiable live. Even if this task is not always easy, we strive to keep our code as accessible and reusable as possible. This should allow us - with minimum effort - development of small extensions that customize the behavior of the web interface.

3.7.1 Tampering with the Notebook application

The first tool that is available to you and that you should be aware of are browser “developers tool”. The exact naming can change across browser and might require the installation of extensions. But basically they can allow you to inspect/modify the DOM, and interact with the JavaScript code that runs the frontend.

- In Chrome and Safari, Developer tools are in the menu View > Developer > JavaScript Console
- In Firefox you might need to install [Firebug](#)

Those will be your best friends to debug and try different approaches for your extensions.

Injecting JS

Using magics

The above tools can be tedious for editing edit long JavaScript files. Therefore we provide the `%%javascript` magic. This allows you to quickly inject JavaScript into the notebook. Still the JavaScript injected this way will not survive reloading. Hence, it is a good tool for testing and refining a script.

You might see here and there people modifying css and injecting js into the notebook by reading file(s) and publishing them into the notebook. Not only does this often break the flow of the notebook and make the re-execution of the notebook broken, but it also means that you need to execute those cells in the entire notebook every time you need to update the code.

This can still be useful in some cases, like the `%autosave` magic that allows you to control the time between each save. But this can be replaced by a JavaScript dropdown menu to select the save interval.

```
[ ]: ## you can inspect the autosave code to see what it does.  
%autosave??
```

custom.js

To inject JavaScript we provide an entry point: `custom.js` that allows the user to execute and load other resources into the notebook. JavaScript code in `custom.js` will be executed when the notebook app starts and can then be used to customize almost anything in the UI and in the behavior of the notebook.

`custom.js` can be found in the `~/.jupyter/custom/custom.js`. You can share your `custom.js` with others.

Back to theory

```
[ ]: from jupyter_core.paths import jupyter_config_dir
jupyter_dir = jupyter_config_dir()
jupyter_dir
```

and `custom.js` is in

```
[ ]: import os.path
custom_js_path = os.path.join(jupyter_dir, 'custom', 'custom.js')
```

```
[ ]: # my custom js
if os.path.isfile(custom_js_path):
    with open(custom_js_path) as f:
        print(f.read())
else:
    print("You don't have a custom.js file")
```

Note that `custom.js` is meant to be modified by user. When writing a script, you can define it in a separate file and add a line of configuration into `custom.js` that will fetch and execute the file.

Warning : even if modification of `custom.js` takes effect immediately after browser refresh (except if browser cache is aggressive), *creating* a file in `static/` directory needs a **server restart**.

3.7.2 Exercise :

- Create a `custom.js` in the right location with the following content:

```
alert("hello world from custom.js")
```

- Restart your server and open any notebook.
- Be greeted by `custom.js`

Have a look at [default custom.js](#), to see it's content and for more explanation.

For the quick ones :

We've seen above that you can change the autosave rate by using a magic. This is typically something I don't want to type every time, and that I don't like to embed into my workflow and documents. (readers don't care what my autosave time is). Let's build an extension that allows us to do it.

Create a dropdown element in the toolbar (DOM `Jupyter.toolbar.element`), you will need

- `Jupyter.notebook.set_autosave_interval(milliseconds)`
- know that 1 min = 60 sec, and 1 sec = 1000 ms

```

var label = jQuery('<label/>').text('AutoScroll Limit:');
var select = jQuery('<select/>')
    // .append(jQuery('<option/>').attr('value', '2').text('2min (default)'))
    .append(jQuery('<option/>').attr('value', undefined).text('disabled'))

    // TODO:
    // the_toolbar_element.append(label)
    // the_toolbar_element.append(select);

select.change(function() {
    var val = jQuery(this).val() // val will be the value in [2]
    // TODO
    // this will be called when dropdown changes
});

var time_m = [1,5,10,15,30];
for (var i=0; i < time_m.length; i++) {
    var ts = time_m[i];

                                //[2]  ____ this will be `val` on [1]
                                //      |
                                //      v

    select.append($('

```

A non-interactive example first

I like my cython to be nicely highlighted

```

Jupyter.config.cell_magic_highlight['magic_text/x-cython'] = {}
Jupyter.config.cell_magic_highlight['magic_text/x-cython'].reg = [/^%cython/]

```

text/x-cython is the name of CodeMirror mode name, magic_ prefix will just patch the mode so that the first line that contains a magic does not screw up the highlighting. reg is a list or regular expression that will trigger the change of mode.

Get more documentation

Sadly, you will have to read the js source file (but there are lots of comments) and/or build the JavaScript documentation using yuidoc. If you have node and yui-doc installed:

```

$ cd ~/jupyter/notebook/notebook/static/notebook/js/
$ yuidoc . --server
warn: (yuidoc): Failed to extract port, setting to the default :3000
info: (yuidoc): Starting YUIDoc@0.3.45 using YUI@3.9.1 with NodeJS@0.10.15
info: (yuidoc): Scanning for yuidoc.json file.
info: (yuidoc): Starting YUIDoc with the following options:
info: (yuidoc):

```

(continues on next page)

(continued from previous page)

```
{ port: 3000,
  nocode: false,
  paths: [ '.' ],
  server: true,
  outdir: './out' }
info: (yuidoc): Scanning for yuidoc.json file.
info: (server): Starting server: http://127.0.0.1:3000
```

and browse <http://127.0.0.1:3000> to get documentation

Some convenience methods

By browsing the documentation you will see that we have some convenience methods that allows us to avoid re-inventing the UI every time :

```
Jupyter.toolbar.add_buttons_group([
    {
        'label'    : 'run qtconsole',
        'icon'     : 'fa-terminal', // select your icon from
                                // http://fontawesome.io/icons/
        'callback': function() {Jupyter.notebook.kernel.execute('%qtconsole')}
    }
    // add more button here if needed.
]);
```

with a lot of icons you can select from.

3.7.3 Cell Metadata

The most requested feature is generally to be able to distinguish an individual cell in the notebook, or run a specific action with them. To do so, you can either use `Jupyter.notebook.get_selected_cell()`, or rely on `CellToolbar`. This allows you to register a set of actions and graphical elements that will be attached to individual cells.

Cell Toolbar

You can see some example of what can be done by toggling the `Cell Toolbar` selector in the toolbar on top of the notebook. It provides two default presets that are `Default` and `slideshow`. `Default` allows the user to edit the metadata attached to each cell manually.

First we define a function that takes at first parameter an element on the DOM in which to inject UI element. The second element is the cell this element was registered with. Then we will need to register that function and give it a name.

Register a callback

```
[ ]: %%javascript
var CellToolbar = Jupyter.CellToolbar
var toggle = function(div, cell) {
    var button_container = $(div)

    // let's create a button that shows the current value of the metadata
    var button = $('<button/>').addClass('btn btn-mini').text(String(cell.metadata.
↵foo));

    // On click, change the metadata value and update the button label
    button.click(function() {
        var v = cell.metadata.foo;
        cell.metadata.foo = !v;
        button.text(String(!v));
    })

    // add the button to the DOM div.
    button_container.append(button);
}

// now we register the callback under the name foo to give the
// user the ability to use it later
CellToolbar.register_callback('tuto.foo', toggle);
```

Registering a preset

This function can now be part of many preset of the CellToolBar.

```
[ ]: %%javascript
Jupyter.CellToolbar.register_preset('Tutorial 1', ['tuto.foo', 'default.rawedit'])
Jupyter.CellToolbar.register_preset('Tutorial 2', ['slideshow.select', 'tuto.foo'])
```

You should now have access to two presets :

- Tutorial 1
- Tutorial 2

And check that the buttons you defined share state when you toggle preset. Also check that the metadata of the cell is modified when you click the button, and that when saved or reloaded the metadata is still available.

Exercise:

Try to wrap the all code in a file, put this file in {jupyter_dir}/custom/<a-name>.js, and add

```
require(['custom/<a-name>']);
```

in custom.js to have this script automatically loaded in all your notebooks.

require is provided by a [JavaScript library](#) that allow you to express dependency. For simple extension like the previous one we directly mute the global namespace, but for more complex extension you could pass a callback to `require([...], <callback>)` call, to allow the user to pass configuration information to your plugin.

In Python language,


```
require(['a/b', 'c/d'], function( e, f){
    e.something()
    f.something()
})
```

could be read as

```
import a.b as e
import c.d as f
e.something()
f.something()
```

See for example @damianavila “ZenMode” plugin :

```
// read that as
// import custom.zenmode.main as zenmode
require(['custom/zenmode/main'], function(zenmode) {
    zenmode.background('images/back12.jpg');
})
```

For the quickest

Try to use the following to bind a dropdown list to `cell.metadata.difficulty.select`.

It should be able to take the 4 following values :

- <None>
- Easy
- Medium
- Hard

We will use it to customize the output of the converted notebook depending on the tag on each cell

```
[1]: # %load soln/celldiff.js
```

```
[ ]:
```

3.8 Importing Jupyter Notebooks as Modules

It is a common problem that people want to import code from Jupyter Notebooks. This is made difficult by the fact that Notebooks are not plain Python files, and thus cannot be imported by the regular Python machinery.

Fortunately, Python provides some fairly sophisticated [hooks](#) into the import machinery, so we can actually make Jupyter notebooks importable without much difficulty, and only using public APIs.

```
[ ]: import io, os, sys, types
```

```
[ ]: from IPython import get_ipython
     from nbformat import read
     from IPython.core.interactiveshell import InteractiveShell
```

Import hooks typically take the form of two objects:

1. a Module **Loader**, which takes a module name (e.g. 'IPython.display'), and returns a Module
2. a Module **Finder**, which figures out whether a module might exist, and tells Python what **Loader** to use

```
[ ]: def find_notebook(fullname, path=None):
    """find a notebook, given its fully qualified name and an optional path

    This turns "foo.bar" into "foo/bar.ipynb"
    and tries turning "Foo_Bar" into "Foo Bar" if Foo_Bar
    does not exist.
    """
    name = fullname.rsplit('.', 1)[-1]
    if not path:
        path = ['']
    for d in path:
        nb_path = os.path.join(d, name + ".ipynb")
        if os.path.isfile(nb_path):
            return nb_path
        # let import Notebook_Name find "Notebook Name.ipynb"
        nb_path = nb_path.replace("_", " ")
        if os.path.isfile(nb_path):
            return nb_path
```

3.8.1 Notebook Loader

Here we have our Notebook Loader. It's actually quite simple - once we figure out the filename of the module, all it does is:

1. load the notebook document into memory
2. create an empty Module
3. execute every cell in the Module namespace

Since IPython cells can have extended syntax, the IPython transform is applied to turn each of these cells into their pure-Python counterparts before executing them. If all of your notebook cells are pure-Python, this step is unnecessary.

```
[ ]: class NotebookLoader(object):
    """Module Loader for Jupyter Notebooks"""
    def __init__(self, path=None):
        self.shell = InteractiveShell.instance()
        self.path = path

    def load_module(self, fullname):
        """import a notebook as a module"""
        path = find_notebook(fullname, self.path)

        print ("importing Jupyter notebook from %s" % path)

        # load the notebook object
        with io.open(path, 'r', encoding='utf-8') as f:
            nb = read(f, 4)

        # create the module and add it to sys.modules
        # if name in sys.modules:
        #     return sys.modules[name]
        mod = types.ModuleType(fullname)
```

(continues on next page)

(continued from previous page)

```

mod.__file__ = path
mod.__loader__ = self
mod.__dict__['get_ipython'] = get_ipython
sys.modules[fullname] = mod

# extra work to ensure that magics that would affect the user_ns
# actually affect the notebook module's ns
save_user_ns = self.shell.user_ns
self.shell.user_ns = mod.__dict__

try:
    for cell in nb.cells:
        if cell.cell_type == 'code':
            # transform the input to executable Python
            code = self.shell.input_transformer_manager.transform_cell(
↪source)

            # run the code in the module
            exec(code, mod.__dict__)
finally:
    self.shell.user_ns = save_user_ns
return mod

```

3.8.2 The Module Finder

The finder is a simple object that tells you whether a name can be imported, and returns the appropriate loader. All this one does is check, when you do:

```
import mynotebook
```

it checks whether `mynotebook.ipynb` exists. If a notebook is found, then it returns a `NotebookLoader`.

Any extra logic is just for resolving paths within packages.

```

[ ]: class NotebookFinder(object):
    """Module finder that locates Jupyter Notebooks"""
    def __init__(self):
        self.loaders = {}

    def find_module(self, fullname, path=None):
        nb_path = find_notebook(fullname, path)
        if not nb_path:
            return

        key = path
        if path:
            # lists aren't hashable
            key = os.path.sep.join(path)

        if key not in self.loaders:
            self.loaders[key] = NotebookLoader(path)
        return self.loaders[key]

```

3.8.3 Register the hook

Now we register the NotebookFinder with `sys.meta_path`

```
[ ]: sys.meta_path.append(NotebookFinder())
```

After this point, my notebooks should be importable.

Let's look at what we have in the CWD:

```
[ ]: ls nbpackage
```

So I should be able to import `nbpackage.mynotebook`.

```
[ ]: import nbpackage.mynotebook
```

Aside: displaying notebooks

Here is some simple code to display the contents of a notebook with syntax highlighting, etc.

```
[ ]: from pygments import highlight
      from pygments.lexers import PythonLexer
      from pygments.formatters import HtmlFormatter

      from IPython.display import display, HTML

      formatter = HtmlFormatter()
      lexer = PythonLexer()

      # publish the CSS for pygments highlighting
      display(HTML("""
<style type='text/css'>
%s
</style>
""" % formatter.get_style_defs()
))
```

```
[ ]: def show_notebook(fname):
      """display a short summary of the cells of a notebook"""
      with io.open(fname, 'r', encoding='utf-8') as f:
          nb = read(f, 4)
          html = []
          for cell in nb.cells:
              html.append("<h4>%s cell</h4>" % cell.cell_type)
              if cell.cell_type == 'code':
                  html.append(highlight(cell.source, lexer, formatter))
              else:
                  html.append("<pre>%s</pre>" % cell.source)
          display(HTML('\n'.join(html)))

      show_notebook(os.path.join("nbpackage", "mynotebook.ipynb"))
```

So my notebook has some code cells, one of which contains some IPython syntax.

Let's see what happens when we import it

```
[ ]: from nbpackage import mynotebook
```

Hooray, it imported! Does it work?

```
[ ]: mynotebook.foo()
```

Hooray again!

Even the function that contains IPython syntax works:

```
[ ]: mynotebook.has_ip_syntax()
```

3.8.4 Notebooks in packages

We also have a notebook inside the nb package, so let's make sure that works as well.

```
[ ]: ls nbpackage/nbs
```

Note that the `__init__.py` is necessary for nb to be considered a package, just like usual.

```
[ ]: show_notebook(os.path.join("nbpackage", "nbs", "other.ipynb"))
```

```
[ ]: from nbpackage.nbs import other
other.bar(5)
```

So now we have importable notebooks, from both the local directory and inside packages.

I can even put a notebook inside IPython, to further demonstrate that this is working properly:

```
[ ]: import shutil
from IPython.paths import get_ipython_package_dir

utils = os.path.join(get_ipython_package_dir(), 'utils')
shutil.copy(os.path.join("nbpackage", "mynotebook.ipynb"),
            os.path.join(utils, "inside_ipython.ipynb"))
```

and import the notebook from `IPython.utils`

```
[ ]: from IPython.utils import inside_ipython
inside_ipython.whatsmyname()
```

This approach can even import functions and classes that are defined in a notebook using the `%%cython` magic.

3.9 Connecting to an existing IPython kernel using the Qt Console

3.9.1 The Frontend/Kernel Model

The traditional IPython (`ipython`) consists of a single process that combines a terminal based UI with the process that runs the users code.

While this traditional application still exists, the modern Jupyter consists of two processes:

- Kernel: this is the process that runs the users code.

- Frontend: this is the process that provides the user interface where the user types code and sees results.

Jupyter currently has 3 frontends:

- Terminal Console (`jupyter console`)
- Qt Console (`jupyter qtconsole`)
- Notebook (`jupyter notebook`)

The Kernel and Frontend communicate over a ZeroMQ/JSON based messaging protocol, which allows multiple Frontends (even of different types) to communicate with a single Kernel. This opens the door for all sorts of interesting things, such as connecting a Console or Qt Console to a Notebook's Kernel. For example, you may want to connect a Qt console to your Notebook's Kernel and use it as a help browser, calling `??` on objects in the Qt console (whose pager is more flexible than the one in the notebook).

This Notebook describes how you would connect another Frontend to an IPython Kernel that is associated with a Notebook. The commands currently given here are specific to the IPython kernel.

3.9.2 Manual connection

To connect another Frontend to a Kernel manually, you first need to find out the connection information for the Kernel using the `%connect_info` magic:

```
[ ]: %connect_info
```

You can see that this magic displays everything you need to connect to this Notebook's Kernel.

3.9.3 Automatic connection using a new Qt Console

You can also start a new Qt Console connected to your current Kernel by using the `%qtconsole` magic. This will detect the necessary connection information and start the Qt Console for you automatically.

```
[ ]: a = 10
```

```
[ ]: %qtconsole
```

The Markdown parser included in the Jupyter Notebook is MathJax-aware. This means that you can freely mix in mathematical expressions using the [MathJax subset of Tex and LaTeX](#). Some examples from the [MathJax demos site](#) are reproduced below, as well as the Markdown+TeX source.

3.10 Motivating Examples

3.10.1 The Lorenz Equations

Source

```
\begin{align}
\dot{x} &= \sigma(y-x) \\
\dot{y} &= \rho x - y - xz \\
\dot{z} &= -\beta z + xy
\end{align}
```

Display**3.10.2 The Cauchy-Schwarz Inequality****Source**

```
\begin{equation*}
\left( \sum_{k=1}^n a_k b_k \right)^2 \leq \left( \sum_{k=1}^n a_k^2 \right) \left( \sum_{k=1}^n b_k^2 \right)
\end{equation*}
```

Display

$$\left(\sum_{k=1}^n a_k b_k \right)^2 \leq \left(\sum_{k=1}^n a_k^2 \right) \left(\sum_{k=1}^n b_k^2 \right)$$

3.10.3 A Cross Product Formula**Source**

```
\begin{equation*}
\mathbf{V}_1 \times \mathbf{V}_2 = \begin{vmatrix} \mathbf{i} & \mathbf{j} & \mathbf{k} \\ \frac{\partial X}{\partial u} & \frac{\partial Y}{\partial u} & 0 \\ \frac{\partial X}{\partial v} & \frac{\partial Y}{\partial v} & 0 \end{vmatrix}
\end{equation*}
```

Display

$$\mathbf{V}_1 \times \mathbf{V}_2 = \begin{vmatrix} \mathbf{i} & \mathbf{j} & \mathbf{k} \\ \frac{\partial X}{\partial u} & \frac{\partial Y}{\partial u} & 0 \\ \frac{\partial X}{\partial v} & \frac{\partial Y}{\partial v} & 0 \end{vmatrix}$$

3.10.4 The probability of getting (k) heads when flipping (n) coins is**Source**

```
\begin{equation*}
P(E) = \{n \text{ choose } k\} p^k (1-p)^{n-k}
\end{equation*}
```

Display

$$P(E) = \binom{n}{k} p^k (1-p)^{n-k}$$

3.10.5 An Identity of Ramanujan

Source

```
\begin{equation*}
\frac{1}{\sqrt{\phi\sqrt{5}-\phi}} e^{\frac{25}{\pi}} =
1+\frac{e^{-2\pi}}{1+\frac{e^{-4\pi}}{1+\frac{e^{-6\pi}}{1+\frac{e^{-8\pi}}{1+\ldots}}}}
\end{equation*}
```

Display

$$\frac{1}{\left(\sqrt{\phi\sqrt{5}-\phi}\right)e^{\frac{25}{\pi}}} = 1 + \frac{e^{-2\pi}}{1 + \frac{e^{-4\pi}}{1 + \frac{e^{-6\pi}}{1 + \frac{e^{-8\pi}}{1 + \dots}}}}$$

3.10.6 A Rogers-Ramanujan Identity

Source

```
\begin{equation*}
1 + \frac{q^2}{(1-q)} + \frac{q^6}{(1-q)(1-q^2)} + \cdots =
\prod_{j=0}^{\infty} \frac{1}{(1-q^{5j+2})(1-q^{5j+3})},
\quad \text{for } |q| < 1.
\end{equation*}
```

Display

$$1 + \frac{q^2}{(1-q)} + \frac{q^6}{(1-q)(1-q^2)} + \cdots = \prod_{j=0}^{\infty} \frac{1}{(1-q^{5j+2})(1-q^{5j+3})}, \quad \text{for } |q| < 1.$$

3.10.7 Maxwell's Equations

Source


```
\begin{align}
\nabla \times \vec{\mathbf{B}} &= \frac{1}{c} \nabla \times \vec{\mathbf{E}} \\
\frac{1}{c} \nabla \times \vec{\mathbf{E}} &= -\frac{1}{c} \frac{\partial \vec{\mathbf{B}}}{\partial t} \\
\nabla \cdot \vec{\mathbf{E}} &= 4\pi \rho \\
\nabla \cdot \vec{\mathbf{B}} &= 0
\end{align}
```

Display

3.10.8 Equation Numbering and References

Equation numbering and referencing will be available in a future version of the Jupyter notebook.

3.10.9 Inline Typesetting (Mixing Markdown and TeX)

While display equations look good for a page of samples, the ability to mix math and *formatted text* in a paragraph is also important.

Source

```
This expression  $\sqrt{3x-1} + (1+x)^2$  is an example of a TeX inline equation in a
[Markdown-formatted] (https://daringfireball.net/projects/markdown/) sentence.
```

Display

This expression $\sqrt{3x-1} + (1+x)^2$ is an example of a TeX inline equation in a *Markdown-formatted* sentence.

3.10.10 Other Syntax

You will notice in other places on the web that $\$$ are needed explicitly to begin and end MathJax typesetting. This is **not** required if you will be using TeX environments, but the Jupyter notebook will accept this syntax on legacy notebooks.

3.10.11 Source

```
$$
\begin{array}{c}
y_1 \\
y_2 \\
z_{3,4}
\end{array}
$$
```

```


$$\begin{array}{c} y_1 \\ y_2 \\ y_3 \end{array}$$


```

```


$$\begin{array}{l} x' = x \sin \phi + z \cos \phi \\ z' = -x \cos \phi + z \sin \phi \end{array}$$


```

```


$$x=4$$


```

3.10.12 Display

$$\begin{array}{l} y_1 \\ y_2 t_i \\ z_{3,4} \\ y_1 \\ y_2 t_i \\ y_3 \end{array}$$

to

$$\begin{aligned} x' &= x \sin \phi + z \cos \phi \\ z' &= -x \cos \phi + z \sin \phi \end{aligned} \quad (3.2)$$

=

$$\begin{aligned} &+ \\ - &= \\ &+ x \cos \phi \\ &z \sin \phi \end{aligned}$$

$$x = 4$$

WHAT TO DO WHEN THINGS GO WRONG

First, have a look at the common problems listed below. If you can figure it out from these notes, it will be quicker than asking for help.

Check that you have the latest version of any packages that look relevant. Unfortunately it's not always easy to figure out what packages are relevant, but if there was a bug that's already been fixed, it's easy to upgrade and get on with what you wanted to do.

4.1 Jupyter fails to start

- Have you `installed it`? ;-)
- If you're using a menu shortcut or Anaconda launcher to start it, try opening a terminal or command prompt and running the command `jupyter notebook`.
- If it can't find `jupyter`, you may need to configure your `PATH` environment variable. If you don't know what that means, and don't want to find out, just (re)install Anaconda with the default settings, and it should set up `PATH` correctly.
- If Jupyter gives an error that it can't find `notebook`, check with `pip` or `conda` that the `notebook` package is installed.
- Try running `jupyter-notebook` (with a hyphen). This should normally be the same as `jupyter notebook` (with a space), but if there's any difference, the version with the hyphen is the 'real' launcher, and the other one wraps that.

4.2 Jupyter doesn't load or doesn't work in the browser

- Try in another browser (e.g. if you normally use Firefox, try with Chrome). This helps pin down where the problem is.
- Try disabling any browser extensions and/or any Jupyter extensions you have installed.
- Some internet security software can interfere with Jupyter. If you have security software, try turning it off temporarily, and look in the settings for a more long-term solution.
- In the address bar, try changing between `localhost` and `127.0.0.1`. They should be the same, but in some cases it makes a difference.

4.3 Jupyter can't start a kernel

Files called *kernel specs* tell Jupyter how to start different kinds of kernels. To see where these are on your system, run `jupyter kernelspec list`:

```
$ jupyter kernelspec list
Available kernels:
  python3      /home/takluyver/.local/lib/python3.6/site-packages/ipykernel/resources
  bash         /home/takluyver/.local/share/jupyter/kernels/bash
  ir           /home/takluyver/.local/share/jupyter/kernels/ir
```

There's a special fallback for the Python kernel: if it doesn't find a real kernelspec, but it can import the `ipykernel` package, it provides a kernel which will run in the same Python environment as the notebook server. A path ending in `ipykernel/resources`, like in the example above, is this default kernel. The default often does what you want, so if the `python3` kernelspec points somewhere else and you can't start a Python kernel, try deleting or renaming that kernelspec folder to expose the default.

If your problem is with another kernel, not the Python one we maintain, you may need to look for support about that kernel.

4.4 Python Environments

Multiple python environments, whether based on Anaconda or Python Virtual environments, are often the source of reported issues. In many cases, these issues stem from the Notebook server running in one environment, while the kernel and/or its resources, derive from another environment. Indicators of this scenario include:

- `import` statements within code cells producing `ImportError` or `ModuleNotFound` exceptions.
- General kernel startup failures exhibited by nothing happening when attempting to execute a cell.

In these situations, take a close look at your environment structure and ensure all packages required by your notebook's code are installed in the correct environment. If you need to run the kernel from different environments than your Notebook server, check out [IPython's documentation](#) for using kernels from different environments as this is the recommended approach. Anaconda's `nb_conda_kernels` package might also be an option for you in these scenarios.

Another thing to check is the `kernel.json` file that will be located in the aforementioned *kernel specs* directory identified by running `jupyter kernelspec list`. This file will contain an `argv` stanza that includes the actual command to run when launching the kernel. Oftentimes, when reinstalling python environments, a previous `kernel.json` will reference an python executable from an old or non-existent location. As a result, it's always a good idea when encountering kernel startup issues to validate the `argv` stanza to ensure all file references exist and are appropriate.

4.5 Windows Systems

Although Jupyter Notebook is primarily developed on the various flavors of the Unix operating system it also supports Microsoft Windows - which introduces its own set of commonly encountered issues, particularly in the areas of security, process management and lower-level libraries.

4.5.1 pywin32 Issues

The primary package for interacting with Windows' primitives is `pywin32`.

- Issues surrounding the creation of the kernel's communication file utilize `jupyter_core`'s `secure_write()` function. This function ensures a file is created in which only the owner of the file has access. If libraries like `pywin32` are not properly installed, issues can arise when it's necessary to use the native Windows libraries.

Here's a portion of such a traceback:

```
File "c:\users\jovyan\python\myenv.venv\lib\site-packages\jupyter_core\paths.py", line 424, in secure_write
win32_restrict_file_to_user(fname)
File "c:\users\jovyan\python\myenv.venv\lib\site-packages\jupyter_core\paths.py", line 359, in win32_restrict_file_to_user
import win32api
ImportError: DLL load failed: The specified module could not be found.
```

- As noted earlier, the installation of `pywin32` can be problematic on Windows configurations. When such an issue occurs, you may need to revisit how the environment was setup. Pay careful attention to whether you're running the 32 or 64 bit versions of Windows and be sure to install appropriate packages for that environment.

Here's a portion of such a traceback:

```
File "C:\Users\jovyan\AppData\Roaming\Python\Python37\site-packages\jupyter_core\paths.py", line 435, in secure_write
win32_restrict_file_to_user(fname)
File "C:\Users\jovyan\AppData\Roaming\Python\Python37\site-packages\jupyter_core\paths.py", line 361, in win32_restrict_file_to_user
import win32api
ImportError: DLL load failed: %1 is not a valid Win32 application
```

Resolving pywin32 Issues

In this case, your `pywin32` module may not be installed correctly and the following should be attempted:

```
pip install --upgrade pywin32
```

or:

```
conda install --force-reinstall pywin32
```

followed by:

```
python.exe Scripts/pywin32_postinstall.py -install
```

where `Scripts` is located in the active Python's installation location.

- Another common failure specific to Windows environments is the location of various python commands. On `*nix` systems, these typically reside in the `bin` directory of the active Python environment. However, on Windows, these tend to reside in the `Scripts` folder - which is a sibling to `bin`. As a result, when encountering kernel startup issues, again, check the `argv` stanza and verify it's pointing to a valid file. You may find that it's pointing in `bin` when `Scripts` is correct, or the referenced file does not include its `.exe` extension - typically resulting in `FileNotFoundError` exceptions.

4.6 This Worked An Hour Ago

The Jupyter stack is very complex and rightfully so, there's a lot going on. On occasion you might find the system working perfectly well, then, suddenly, you can't get past a certain cell due to `import` failures. In these situations, it's best to ask yourself if any new python files were added to your notebook development area.

These issues are usually evident by carefully analyzing the traceback produced in the notebook error or the Notebook server's command window. In these cases, you'll typically find the Python kernel code (from `IPython` and `ipykernel`) performing *its* imports and notice a file from your Notebook development error included in that traceback followed by an `AttributeError`:

```
File "C:\Users\jovyan\anaconda3\lib\site-packages\ipykernel\connect.py", line 13, in
from IPython.core.profilendir import ProfileDir
File "C:\Users\jovyan\anaconda3\lib\site-packages\IPython_init.py", line 55, in
from .core.application import Application
...
File "C:\Users\jovyan\anaconda3\lib\site-packages\ipython_genutils\path.py", line 13,
↳ in
import random
File "C:\Users\jovyan\Desktop\Notebooks\random.py", line 4, in
rand_set = random.sample(english_words_lower_set, 12)
AttributeError: module 'random' has no attribute 'sample'
```

What has happened is that you have named a file that conflicts with an installed package that is used by the kernel software and now introduces a conflict preventing the kernel's startup.

Resolution: You'll need to rename your file. A best practice would be to prefix or *namespace* your files so as not to conflict with any python package.

4.7 Asking for help

As with any problem, try searching to see if someone has already found an answer. If you can't find an existing answer, you can ask questions at:

- The [Jupyter Discourse Forum](#)
- The [jupyter-notebook](#) tag on [Stackoverflow](#)
- Peruse the [jupyter/help](#) repository on [Github](#) (read-only)
- Or in an issue on another repository, if it's clear which component is responsible. Typical repositories include:
 - `jupyter_core` - `secure_write()` and file path issues
 - `jupyter_client` - kernel management issues found in Notebook server's command window.
 - `IPython` and `ipykernel` - kernel runtime issues typically found in Notebook server's command window and/or Notebook cell execution.

4.7.1 Gathering Information

Should you find that your problem warrants that an issue be opened in [notebook](#) please don't forget to provide details like the following:

- What error messages do you see (within your notebook and, more importantly, in the Notebook server's command window)?
- What platform are you on?
- How did you install Jupyter?
- What have you tried already?

The `jupyter troubleshoot` command collects a lot of information about your installation, which can also be useful.

When providing textual information, it's most helpful if you can *scrape* the contents into the issue rather than providing a screenshot. This enables others to select pieces of that content so they can search more efficiently and try to help.

Remember that it's not anyone's job to help you. We want Jupyter to work for you, but we can't always help everyone individually.

CHANGELOG

A summary of changes in the Jupyter notebook. For more detailed information, see [GitHub](#).

Use `pip install notebook --upgrade` or `conda upgrade notebook` to upgrade to the latest release.

We strongly recommend that you upgrade pip to version 9+ of pip before upgrading notebook.

Use `pip install pip --upgrade` to upgrade pip. Check pip version with `pip --version`.

5.1 6.3.0

5.1.1 Merged PRs

- Add square logo and desktop entry files #6010 (@befeleme)
- Modernize Changelog #6008 (@afshin)
- Add missing “import inspect” #5999 (@mgeier)
- Add Codecov badge to README #5989 (@thomasrockhu)
- Remove configuration for nosetests from setup.cfg #5986 (@frenzymadness)
- Update security.rst #5978 (@dlrice)
- Docs-Translations: Updated Hindi and Chinese Readme.md #5976 (@rjn01)
- Allow /metrics by default if auth is off #5974 (@blairdrummond)
- Skip terminal tests on Windows 3.9+ (temporary) #5968 (@kevin-bates)
- Update GatewayKernelManager to derive from AsyncMappingKernelManager #5966 (@kevin-bates)
- Drop use of deprecated pyzmq.ioloop #5965 (@kevin-bates)
- Drop support for Python 3.5 #5962 (@kevin-bates)
- Allow jupyter_server-based contents managers in notebook #5957 (@afshin)
- Russian translation fixes #5954 (@insolor)
- Increase culling test idle timeout #5952 (@kevin-bates)
- Re-enable support for answer_yes flag #5941 (@afshin)
- Replace Travis and Appveyor with Github Actions #5938 (@kevin-bates)
- DOC: Server extension, extra docs on configuration/authentication. #5937 (@Carreau)

5.1.2 Contributors to this release

(GitHub contributors page for this release)

@abielhammonds | @afshin | @ajharry | @Alokrar | @befeleme | @blairdrummond | @blink1073 | @bollwyvl | @Carreau | @ChenChenDS | @cosmoscalibur | @dlrice | @dwanneruchi | @ElisonSherton | @FazeelUsmani | @frenzymadness | @goerz | @insolor | @jasongrout | @JianghuiDu | @JuzerShakir | @kevin-bates | @Khalilsqu | @meeseeksdev | @mgeier | @michaelpedota | @mjbright | @MSeal | @ncoughlin | @NTimmons | @ProsperousHeart | @rjn01 | @slw07g | @stenivan | @takluyver | @thomasrockhu | @wgilpin | @wxtt522 | @yuvipanda | @Zsailer

5.2 6.2.0

5.3 Merged PRs

- Increase minimum tornado version (5933)
- Adjust skip decorators to avoid remaining dependency on nose (5932)
- Ensure that cell ids persist after save (5928)
- Add reconnection to Gateway (form nb2kg) (5924)
- Fix some typos (5917)
- Handle TrashPermissionError, now that it exist (5894)

Thank you to all the contributors:

- @kevin-bates
- @mishaschwartz
- @oyvsyo
- @user202729
- @stefanor

5.4 6.1.6

5.5 Merged PRs

- do not require nose for testing (5826)
- [docs] Update Chinese and Hindi readme.md (5823)
- Add support for creating terminals via GET (5813)
- Made doc translations in Hindi and Chinese (5787)

Thank you to all the contributors:

- @pgajdos
- @rjn01
- @kevin-bates
- @virejdasani

5.6 6.1.5

6.1.5 is a security release, fixing one vulnerability:

- Fix open redirect vulnerability GHSA-c7vm-f5p4-8fqh (CVE to be assigned)

5.7 6.1.4

- Fix broken links to jupyter documentation ([5686](#))
- Add additional entries to troubleshooting section ([5695](#))
- Revert change in page alignment ([5703](#))
- Bug fix: remove double encoding in download files ([5720](#))
- Fix typo for Check in zh_CN ([5730](#))
- Require a file name in the “Save As” dialog ([5733](#))

Thank you to all the contributors:

- bdbai
- Jaipreet Singh
- Kevin Bates
- Pavel Panchekha
- Zach Sailer

5.8 6.1.3

- Title new buttons with label if action undefined ([5676](#))

Thank you to all the contributors:

- Kyle Kelley

5.9 6.1.2

- Fix russian message format for delete/duplicate actions ([5662](#))
- Remove unnecessary import of bind_unix_socket ([5666](#))
- Tooltip style scope fix ([5672](#))

Thank you to all the contributors:

- Dmitry Akatov
- Kevin Bates
- Magda Stenius

5.10 6.1.1

- Prevent inclusion of `requests_unixsocket` on Windows (5650)

Thank you to all the contributors:

- Kevin Bates

5.11 6.1.0

Please note that this repository is currently maintained by a skeleton crew of maintainers from the Jupyter community. For our approach moving forward, please see this [notice](#) from the README. Thank you.

Here is an enumeration of changes made since the last release and included in 6.1.0.

- Remove deprecated encoding parameter for Python 3.9 compatibility. (5174)
- Add support for async kernel management (4479)
- Fix typo in `password_required` help message (5320)
- Gateway only: Ensure launch and request timeouts are in sync (5317)
- Update Markdown Cells example to HTML5 video tag (5411)
- Integrated `LoginWidget` into edit to enable users to logout from the t... (5406)
- Update message about minimum Tornado version (5222)
- Logged notebook type (5425)
- Added `nl` language (5354)
- Add UNIX socket support to notebook server. (4835)
- Update `CodeMirror` dependency (5198)
- Tree added download multiple files (5351)
- Toolbar buttons tooltip: show help instead of label (5107)
- Remove unnecessary import of `requests_unixsocket` (5451)
- Add ability to cull terminals and track last activity (5372)
- Code refactoring `notebook.js` (5352)
- Install `terminado` for docs build (5462)
- Convert notifications JS test to selenium (5455)
- Add cell attachments to markdown example (5412)
- Add Japanese document (5231)
- Migrate Move multiselection test to selenium (5158)
- Use `cmdtrl-enter` to run a cell (5120)
- Fix broken “Raw cell MIME type” dialog (5385)
- Make a notebook writable after successful save-as (5296)
- Add actual watch script (4738)
- Added `--autoreload` flag to `NotebookApp` (4795)

- Enable `check_origin` on gateway websocket communication (5471)
- Restore detection of missing terminado package (5465)
- Culling: ensure `last_activity` attr exists before use (5355)
- Added functionality to allow filter kernels by Jupyter Enterprise Gat... (5484)
- ‘Play’ icon for run-cell toolbar button (2922)
- Bump minimum version of jQuery to 3.5.0 (5491)
- Remove old JS markdown tests, add a new one in selenium (5497)
- Add support for more RTL languages (5036)
- Make markdown cells stay RTL in edit mode (5037)
- Unforce RTL output display (5039)
- Fixed multicursor backspacing (4880)
- Implemented Split Cell for multicursor (4824)
- Alignment issue [FIXED] (3173)
- MathJax: Support for `\gdef` (4407)
- Another (Minor) Duplicate Code Reduction (5316)
- Update readme regarding maintenance (5500)
- Document contents chunks (5508)
- Backspace deletes empty line (5516)
- The dropdown submenu at notebook page is not keyboard accessible (4732)
- Tooltips visible through keyboard navigation for specified buttons (4729)
- Fix for recursive symlink (4670)
- Fix for the terminal shutdown issue (4180)
- Add japanese translation files (4490)
- Workaround for socket permission errors on Cygwin (4584)
- Implement optional markdown header and footer files (4043)
- Remove double link when using `custom_display_url` (5544)
- Respect `cell.is_editable` during find-and-replace (5545)
- Fix exception causes all over the codebase (5556)
- Improve login shell heuristics (5588)
- Added support for `JUPYTER_TOKEN_FILE` (5587)
- Kill notebook itself when server cull idle kernel (5593)
- Implement password hashing with bcrypt (3793)
- Fix broken links (5600)
- Russian internationalization support (5571)
- Add a metadata tag to override notebook direction (ltr/rtl) (5052)
- Paste two images from clipboard in markdown cell (5598)

- Add keyboard shortcuts to menu dropdowns ([5525](#))
- Update codemirror to 5.56.0+components1 ([5637](#))

Thank you to all the contributors:

- Aaron Myatt
- Adam Blake
- Afshin Taylor Darian
- Aman Bansal
- Ben Thayer
- berendjan
- Bruno P. Kinoshita
- bzinberg
- Christophe Cadilhac
- Daiki Katsuragawa
- David Lukes
- Dmitriy Q
- dmpe
- dylanzjy
- dSchurch
- E. M. Bray
- ErwinRussel
- Felix Mönckemeyer
- Grant Nestor
- Jarrad Whitaker
- Jesus Panales Castillo
- Joshua Zeltser
- Karthikeyan Singaravelan
- Kenichi Ito
- Kevin Bates
- Koki Nishihara
- Kris Wilson
- Kyle Kelley
- Laura Merlo
- levinxo
- Luciano Resende
- Luis Cabezon Manchado
- Madhusudhan Srinivasa

- Matthias Geier
- mattn
- Max Klein
- Min RK
- Mingxuan Lin
- Mohammad Mostafa Farzan
- Niko Felger
- Norah Abanumay
- Onno Broekmans
- PierreMB
- pinarkavak
- Ram Rachum
- Reece Hart
- Remi Rampin
- Rohit Sanjay
- Shane Canon
- Simon Li
- Steinar Sturlaugsson
- Steven Silvester
- taohan16
- Thew Dhanat
- Thomas Kluyver
- Toon Baeyens
- Vidar Tonaas Fauske
- Zachary Sailer

5.12 6.0.3

- Dependency updates to fix startup issues on Windows platform
- Add support for nbconvert 6.x
- Creation of recent tab

Thanks for all the contributors:

- Luciano Resende
- Kevin Bates
- ahangleben
- Zachary Sailer

- Pallavi Bharadwaj
- Thomas Kluyver
- Min RK
- forest0
- Bibo Hao
- Michal Charemza
- Sergey Shevelev
- Shuichiro MAKIGAKI
- krinsman
- TPartida
- Landen McDonald
- Tres DuBiel

5.13 6.0.2

- Update JQuery dependency to version 3.4.1 to fix security vulnerability (CVE-2019-11358)
- Update CodeMirror to version 5.48.4 to fix Python formatting issues
- Continue removing obsolete Python 2.x code/dependencies
- Multiple documentation updates

Thanks for all the contributors:

- David Robles
- Jason Grout
- Kerwin Sun
- Kevin Bates
- Kyle Kelley
- Luciano Resende
- Marcus D Sherman
- Sasaki Takeru
- Tom Jarosz
- Vidar Tonaas Fauske
- Wes Turner
- Zachary Sailer

5.14 6.0.1

- Attempt to re-establish websocket connection to Gateway ([4777](#))
- Add missing react-dom js to package data ([4772](#))

Thanks for all the contributors:

- Eunsoo Park
- Min RK

5.15 6.0

This is the first major release of the Jupyter Notebook since version 5.0 (March 2017).

We encourage users to start trying JupyterLab, which has just announced it's 1.0 release in preparation for a future transition.

- Remove Python 2.x support in favor of Python 3.5 and higher.
- Multiple accessibility enhancements and bug-fixes.
- Multiple translation enhancements and bug-fixes.
- Remove deprecated ANSI CSS styles.
- Native support to forward requests to Jupyter Gateway(s) (Embedded NB2KG).
- Use JavaScript to redirect users to notebook homepage.
- Enhanced SSL/TLS security by using `PROTOCOL_TLS` which selects the highest ssl/tls protocol version available that both the client and server support. When `PROTOCOL_TLS` is not available use `PROTOCOL_SSLv23`.
- Add `?no_track_activity=1` argument to allow API requests. to not be registered as activity (e.g. API calls by external activity monitors).
- Kernels shutting down due to an idle timeout is no longer considered an activity-updating event.
- Further improve compatibility with tornado 6 with improved checks for when websockets are closed.
- Launch the browser with a local file which redirects to the server address including the authentication token. This prevents another logged-in user from stealing the token from command line arguments and authenticating to the server. The single-use token previously used to mitigate this has been removed. Thanks to Dr. Owain Kenway for suggesting the local file approach.
- Respect nbconvert entrypoints as sources for exporters
- Update to CodeMirror to 5.37, which includes f-string syntax for Python 3.6.
- Update jquery-ui to 1.12
- Execute cells by clicking icon in input prompt.
- New “Save as” menu option.
- When serving on a loopback interface, protect against DNS rebinding by checking the `Host` header from the browser. This check can be disabled if necessary by setting `NotebookApp.allow_remote_access`. (Disabled by default while we work out some Mac issues in [3754](#)).
- Add `kernel_info_timeout` traitlet to enable restarting slow kernels.
- Add `custom_display_host` config option to override displayed URL.

- Add /metrics endpoint for Prometheus Metrics.
- Optimize large file uploads.
- Allow access control headers to be overridden in `jupyter_notebook_config.py` to support greater CORS and proxy configuration flexibility.
- Add support for terminals on windows.
- Add a “restart and run all” button to the toolbar.
- Frontend/extension-config: allow default json files in a .d directory.
- Allow setting token via `jupyter_token` env.
- Cull idle kernels using `--MappingKernelManager.cull_idle_timeout`.
- Allow read-only notebooks to be trusted.
- Convert JS tests to Selenium.

Security Fixes included in previous minor releases of Jupyter Notebook and also included in version 6.0.

- Fix Open Redirect vulnerability (CVE-2019-10255) where certain malicious URLs could redirect from the Jupyter login page to a malicious site after a successful login.
- Contains a security fix for a cross-site inclusion (XSSI) vulnerability (CVE-2019-9644), where files at a known URL could be included in a page from an unauthorized website if the user is logged into a Jupyter server. The fix involves setting the `X-Content-Type-Options: nosniff` header, and applying CSRF checks previously on all non-GET API requests to GET requests to API endpoints and the `/files/` endpoint.
- Check Host header to more securely protect localhost deployments from DNS rebinding. This is a pre-emptive measure, not fixing a known vulnerability. Use `.NotebookApp.allow_remote_access` and `.NotebookApp.local_hostnames` to configure access.
- Upgrade bootstrap to 3.4, fixing an XSS vulnerability, which has been assigned [CVE-2018-14041](#).
- Contains a security fix preventing malicious directory names from being able to execute javascript.
- Contains a security fix preventing nbconvert endpoints from executing javascript with access to the server API. CVE request pending.

Thanks for all the contributors:

- AAYUSH SINHA
- Aaron Hall, MBA
- Abhinav Sagar
- Adam Rule
- Adeel Ahmad
- Alex Rothberg
- Amy Skerry-Ryan
- Anastasis Germanidis
- Andrés Sánchez
- Arjun Radhakrishna
- Arovit Narula
- Benda Xu
- Björn Grüning

- Brian E. Granger
- Carol Willing
- Celina Kilcrease
- Chris Holdgraf
- Chris Miller
- Ciaran Langton
- Damian Avila
- Dana Lee
- Daniel Farrell
- Daniel Nicolai
- Darío Hereñú
- Dave Aitken
- Dave Foster
- Dave Hirschfeld
- Denis Ledoux
- Dmitry Mikushin
- Dominic Kuang
- Douglas Hanley
- Elliott Sales de Andrade
- Emilio Talamante Lugo
- Eric Perry
- Ethan T. Hendrix
- Evan Van Dam
- Francesco Franchina
- Frédéric Chapoton
- Félix-Antoine Fortin
- Gabriel
- Gabriel Nützi
- Gabriel Ruiz
- Gestalt LUR
- Grant Nestor
- Gustavo Efeiche
- Harsh Vardhan
- Heng GAO
- Hisham Elsheshtawy
- Hong Xu

- Ian Rose
- Ivan Ogasawara
- J Forde
- Jason Grout
- Jessica B. Hamrick
- Jiaqi Liu
- John Emmons
- Josh Barnes
- Karthik Balakrishnan
- Kevin Bates
- Kirit Thadaka
- Kristian Gregorius Hustad
- Kyle Kelley
- Leo Gallucci
- Lilian Besson
- Lucas Seiki Oshiro
- Luciano Resende
- Luis Angel Rodriguez Guerrero
- M Pacer
- Maarten Breddels
- Mac Knight
- Madicken Munk
- Maitiú Ó Ciaráin
- Marc Udoff
- Mathis HAMMEL
- Mathis Rosenhauer
- Matthias Bussonnier
- Matthias Geier
- Max Vovshin
- Maxime Mouchet
- Michael Chirico
- Michael Droettboom
- Michael Heilman
- Michael Scott Cuthbert
- Michal Charemza
- Mike Boyle

- Milos Miljkovic
- Min RK
- Miro Hrončok
- Nicholas Bollweg
- Nitesh Sawant
- Ondrej Jariabka
- Park Hae Jin
- Paul Ivanov
- Paul Masson
- Peter Parente
- Pierre Tholoniati
- Remco Verhoef
- Roland Weber
- Roman Kornev
- Rosa Swaby
- Roy Hyunjin Han
- Sally
- Sam Lau
- Samar Sultan
- Shiti Saxena
- Simon Biggs
- Spencer Park
- Stephen Ward
- Steve (Gadget) Barnes
- Steven Silvester
- Surya Prakash Susarla
- Syed Shah
- Sylvain Corlay
- Thomas Aarholt
- Thomas Kluyver
- Tim
- Tim Head
- Tim Klever
- Tim Metzler
- Todd
- Tom Jorquera

- Tyler Makaro
- Vaibhav Sagar
- Victor
- Vidar Tonaas Fauske
- Vu Minh Tam
- Vít Tuček
- Will Costello
- Will Starms
- William Hosford
- Xiaohan Li
- Yuvi Panda
- ashley teoh
- nullptr

5.16 5.7.8

- Fix regression in restarting kernels in 5.7.5. The restart handler would return before restart was completed.
- Further improve compatibility with tornado 6 with improved checks for when websockets are closed.
- Fix regression in 5.7.6 on Windows where .js files could have the wrong mime-type.
- Fix Open Redirect vulnerability (CVE-2019-10255) where certain malicious URLs could redirect from the Jupyter login page to a malicious site after a successful login. 5.7.7 contained only a partial fix for this issue.

5.17 5.7.6

5.7.6 contains a security fix for a cross-site inclusion (XSSI) vulnerability (CVE-2019-9644), where files at a known URL could be included in a page from an unauthorized website if the user is logged into a Jupyter server. The fix involves setting the `X-Content-Type-Options: nosniff` header, and applying CSRF checks previously on all non-GET API requests to GET requests to API endpoints and the `/files/` endpoint.

The attacking page is able to access some contents of files when using Internet Explorer through script errors, but this has not been demonstrated with other browsers.

5.18 5.7.5

- Fix compatibility with tornado 6 ([4392](#), [4449](#)).
- Fix opening integer filedescriptor during startup on Python 2 ([4349](#))
- Fix compatibility with asynchronous `[KernelManager.restart_kernel][.title-ref]` methods ([4412](#))

5.19 5.7.4

5.7.4 fixes a bug introduced in 5.7.3, in which the `list_running_servers()` function attempts to parse HTML files as JSON, and consequently crashes (4284).

5.20 5.7.3

5.7.3 contains one security improvement and one security fix:

- Launch the browser with a local file which redirects to the server address including the authentication token (4260). This prevents another logged-in user from stealing the token from command line arguments and authenticating to the server. The single-use token previously used to mitigate this has been removed. Thanks to Dr. Owain Kenway for suggesting the local file approach.
- Upgrade bootstrap to 3.4, fixing an XSS vulnerability, which has been assigned CVE-2018-14041 (4271).

5.21 5.7.2

5.7.2 contains a security fix preventing malicious directory names from being able to execute javascript. CVE request pending.

5.22 5.7.1

5.7.1 contains a security fix preventing nbconvert endpoints from executing javascript with access to the server API. CVE request pending.

5.23 5.7.0

New features:

- Update to CodeMirror to 5.37, which includes f-string syntax for Python 3.6 (3816)
- Update jquery-ui to 1.12 (3836)
- Check Host header to more securely protect localhost deployments from DNS rebinding. This is a pre-emptive measure, not fixing a known vulnerability (3766). Use `.NotebookApp.allow_remote_access` and `.NotebookApp.local_hostnames` to configure access.
- Allow access-control-allow-headers to be overridden (3886)
- Allow configuring `max_body_size` and `max_buffer_size` (3829)
- Allow configuring `get_secure_cookie` keyword-args (3778)
- Respect nbconvert endpoints as sources for exporters (3879)
- Include translation sources in source distributions (3925, 3931)
- Various improvements to documentation (3799, 3800, 3806, 3883, 3908)

Fixing problems:

- Fix breadcrumb link when running with a base url (3905)

- Fix possible type error when closing activity stream (3907)
- Disable metadata editing for non-editable cells (3744)
- Fix some styling and alignment of prompts caused by regressions in 5.6.0.
- Enter causing page reload in shortcuts editor (3871)
- Fix uploading to the same file twice (3712)

See the 5.7 milestone on GitHub for a complete list of [pull requests](#) involved in this release.

Thanks to the following contributors:

- Aaron Hall
- Benjamin Ragan-Kelley
- Bill Major
- bxy007
- Dave Aitken
- Denis Ledoux
- Félix-Antoine Fortin
- Gabriel
- Grant Nestor
- Kevin Bates
- Kristian Gregorius Hustad
- M Pacer
- Madicken Munk
- Maitiu O Ciarain
- Matthias Bussonnier
- Michael Boyle
- Michael Chirico
- Mokkapati, Praneet(ES)
- Peter Parente
- Sally Wilsak
- Steven Silvester
- Thomas Kluyver
- Walter Martin

5.24 5.6.0

New features:

- Execute cells by clicking icon in input prompt (3535, 3687)
- New “Save as” menu option (3289)
- When serving on a loopback interface, protect against DNS rebinding by checking the `Host` header from the browser (3714). This check can be disabled if necessary by setting `NotebookApp.allow_remote_access`. (Disabled by default while we work out some Mac issues in 3754).
- Add `kernel_info_timeout` traitlet to enable restarting slow kernels (3665)
- Add `custom_display_host` config option to override displayed URL (3668)
- Add `/metrics` endpoint for Prometheus Metrics (3490)
- Update to MathJax 2.7.4 (3751)
- Update to jQuery 3.3 (3655)
- Update marked to 0.4 (3686)

Fixing problems:

- Don’t duplicate token in displayed URL (3656)
- Clarify displayed URL when listening on all interfaces (3703)
- Don’t trash non-empty directories on Windows (3673)
- Include LICENSE file in wheels (3671)
- Don’t show “0 active kernels” when starting the notebook (3696)

Testing:

- Add find replace test (3630)
- Selenium test for deleting all cells (3601)
- Make creating a new notebook more robust (3726)

Thanks to the following contributors:

- Arovit Narula (arovit)
- lucasoshiro (lucasoshiro)
- M Pacer (mpacer)
- Thomas Kluyver (takluyver)
- Todd (toddrme2178)
- Yuvi Panda (yuvipanda)

See the 5.6 milestone on GitHub for a complete list of [pull requests](#) involved in this release.

5.25 5.5.0

New features:

- The files list now shows file sizes (3539)
- Add a quit button in the dashboard (3004)
- Display hostname in the terminal when running remotely (3356, 3593)
- Add slides exportation/download to the menu (3287)
- Add any extra installed nbconvert exporters to the “Download as” menu (3323)
- Editor: warning when overwriting a file that is modified on disk (2783)
- Display a warning message if cookies are not enabled (3511)
- Basic `__version__` reporting for extensions (3541)
- Add `NotebookApp.terminals_enabled` config option (3478)
- Make buffer time between last modified on disk and last modified on last save configurable (3273)
- Allow binding custom shortcuts for ‘close and halt’ (3314)
- Add description for ‘Trusted’ notification (3386)
- Add `settings['activity_sources']` (3401)
- Add an `output_updated.OutputArea` event (3560)

Fixing problems:

- Fixes to improve web accessibility (3507)
- Fixed color contrast issue in `tree.less` (3336)
- Allow cancelling upload of large files (3373)
- Don’t clear login cookie on requests without cookie (3380)
- Don’t trash files on different device to home dir on Linux (3304)
- Clear waiting asterisks when restarting kernel (3494)
- Fix output prompt when `execution_count` missing (3236)
- Make the ‘changed on disk’ dialog work when displayed twice (3589)
- Fix going back to root directory with history in notebook list (3411)
- Allow defining keyboard shortcuts for missing actions (3561)
- Prevent default on pageup/pagedown when completer is active (3500)
- Prevent default event handling on new terminal (3497)
- `ConfigManager` should not write out default values found in the `.d` directory (3485)
- Fix leak of `iopub` object in activity monitoring (3424)
- Javascript lint in `notebooklist.js` (3409)
- Some Javascript syntax fixes (3294)
- Convert native for loop to `Array.forEach()` (3477)
- Disable cache when downloading nbconvert output (3484)

- Add missing digestmod arg to HMAC (3399)
- Log OSError failing to create less-critical files during startup (3384)
- Use powershell on Windows (3379)
- API spec improvements, API handler improvements (3368)
- Set notebook to dirty state after change to kernel metadata (3350)
- Use CSP header to treat served files as belonging to a separate origin (3341)
- Don't install gettext into builtins (3330)
- Add missing `import _` (3316, 3326)
- Write `notebook.json` file atomically (3305)
- Fix clicking with modifiers, page title updates (3282)
- Upgrade jQuery to version 2.2 (3428)
- Upgrade xterm.js to 3.1.0 (3189)
- Upgrade moment.js to 2.19.3 (3562)
- Upgrade CodeMirror to 5.35 (3372)
- "Require" `pymzmq>=17` (3586)

Documentation:

- Documentation updates and organisation (3584)
- Add section in docs about privacy (3571)
- Add explanation on how to change the type of a cell to Markdown (3377)
- Update docs with confd implementation details (3520)
- Add more information for where `jupyter_notebook_config.py` is located (3346)
- Document options to enable nbextensions in specific sections (3525)
- jQuery attribute selector value MUST be surrounded by quotes (3527)
- Do not execute special notebooks with nbsphinx (3360)
- Other minor fixes in 3288, 3528, 3293, 3367

Testing:

- Testing with Selenium & Sauce labs (3321)
- Selenium utils + markdown rendering tests (3458)
- Convert insert cell tests to Selenium (3508)
- Convert prompt numbers tests to Selenium (3554)
- Convert delete cells tests to Selenium (3465)
- Convert undelete cell tests to Selenium (3475)
- More selenium testing utilities (3412)
- Only check links when build is trigger by Travis Cron job (3493)
- Fix Appveyor build errors (3430)
- Undo patches in teardown before attempting to delete files (3459)

- Get tests running with tornado 5 (3398)
- Unpin ipykernel version on Travis (3223)

Thanks to the following contributors:

- Arovit Narula ([arovit](#))
- Ashley Teoh ([ashleytqy](#))
- Nicholas Bollweg ([bollwyvl](#))
- Alex Rothberg ([cancan101](#))
- Celina Kilcrease ([ckilcrease](#))
- dabuside ([dabuside](#))
- Damian Avila ([damianavila](#))
- Dana Lee ([danagilliann](#))
- Dave Hirschfeld ([dhirschfeld](#))
- Heng GAO ([ehengao](#))
- Leo Gallucci ([elgalu](#))
- Evan Van Dam ([evandam](#))
- forbxy ([forbxy](#))
- Grant Nestor ([gnestor](#))
- Ethan T. Hendrix ([hendrixet](#))
- Miro Hrončok ([hroncok](#))
- Paul Ivanov ([ivanov](#))
- Darío Hereñú ([kant](#))
- Kevin Bates ([kevin-bates](#))
- Maarten Breddels ([maartenbreddels](#))
- Michael Droettboom ([mdboom](#))
- Min RK ([minrk](#))
- M Pacer ([mpacer](#))
- Peter Parente ([parente](#))
- Paul Masson ([paulmasson](#))
- Philipp Rudiger ([philippjfr](#))
- Mac Knight ([Shels1909](#))
- Hisham Elsheshtawy ([Sheshtawy](#))
- Simon Biggs ([SimonBiggs](#))
- Sunil Hari ([@sunilhari](#))
- Thomas Kluyver ([takluyver](#))
- Tim Klever ([tklever](#))
- Gabriel Ruiz ([unnamedplay-r](#))

- Vaibhav Sagar ([vaibhavsagar](#))
- William Hosford ([whosford](#))
- Hong ([xuhdev](#))

See the 5.5 milestone on GitHub for a complete list of [pull requests](#) involved in this release.

5.26 5.4.1

A security release to fix [CVE-2018-8768](#).

Thanks to [Alex](#) for identifying this bug, and Jonathan Kamens and Scott Sanderson at Quantopian for verifying it and bringing it to our attention.

5.27 5.4.0

- Fix creating files and folders after navigating directories in the dashboard ([3264](#)).
- Enable printing notebooks in colour, removing the CSS that made everything black and white ([3212](#)).
- Limit the completion options displayed in the notebook to 1000, to avoid performance issues with very long lists ([3195](#)).
- Accessibility improvements in `tree.html` ([3271](#)).
- Added alt-text to the kernel logo image in the notebook UI ([3228](#)).
- Added a test on Travis CI to flag if symlinks are accidentally introduced in the future. This should prevent the issue that necessitated `release-5.3.1` `{.interpreted-text role="ref"}` ([3227](#)).
- Use lowercase letters for random IDs generated in our Javascript ([3264](#)).
- Removed duplicate code setting `TextCell.notebook` ([3256](#)).

Thanks to the following contributors:

- Alex Soderman ([asoderman](#))
- Matthias Bussonnier ([Carreau](#))
- Min RK ([minrk](#))
- Nitesh Sawant ([ns23](#))
- Thomas Kluyver ([takluyver](#))
- Yuvi Panda ([yuvipanda](#))

See the 5.4 milestone on GitHub for a complete list of [pull requests](#) involved in this release.

5.28 5.3.1

Replaced a symlink in the repository with a copy, to fix issues installing on Windows (3220).

5.29 5.3.0

This release introduces a couple notable improvements, such as terminal support for Windows and support for OS trash (files deleted from the notebook dashboard are moved to the OS trash vs. deleted permanently).

- Add support for terminals on windows (3087).
- Add a “restart and run all” button to the toolbar (2965).
- Send files to os trash mechanism on delete (1968).
- Allow programmatic copy to clipboard (3088).
- Use DOM History API for navigating between directories in the file browser (3115).
- Add translated files to folder(docs-translations) (3065).
- Allow non empty dirs to be deleted (3108).
- Set cookie on base_url (2959).
- Allow token-authenticated requests cross-origin by default (2920).
- Change cull_idle_timeout_minimum to 1 from 300 (2910).
- Config option to shut down server after n seconds with no kernels (2963).
- Display a “close” button on load notebook error (3176).
- Add action to command palette to run CodeMirror’s “indentAuto” on selection (3175).
- Add option to specify extra services (3158).
- Warn_bad_name should not use global name (3160).
- Avoid overflow of hidden form (3148).
- Fix shutdown trans loss (3147).
- Find available kernelspecs more efficiently (3136).
- Don’t try to translate missing help strings (3122).
- Frontend/extension-config: allow default json files in a .d directory (3116).
- Use [requirejs]{.title-ref} vs. [require]{.title-ref} (3097).
- Fixes some ui bugs in firefox #3044 (3058).
- Compare non-specific language code when choosing to use arabic numerals (3055).
- Fix save-script deprecation (3053).
- Include moment locales in package_data (3051).
- Fix moment locale loading in bidi support (3048).
- Tornado 5: periodiccallback loop arg will be removed (3034).
- Use [/files]{.title-ref} prefix for pdf-like files (3031).
- Add folder for document translation (3022).

- When login-in via token, let a chance for user to set the password (3008).
- Switch to jupyter_core implementation of ensure_dir_exists (3002).
- Send http shutdown request on 'stop' subcommand (3000).
- Work on loading ui translations (2969).
- Fix ansi inverse (2967).
- Add send2trash to requirements for building docs (2964).
- I18n readme.md improvement (2962).
- Add 'reason' field to json error responses (2958).
- Add some padding for stream outputs (3194).
- Always use setuptools in setup.py (3206).
- Fix clearing cookies on logout when base_url is configured (3207).

Thanks to the following contributors:

- bacboc (bacboc)
- Steven Silvester (blink1073)
- Matthias Bussonnier (Carreau)
- ChungJooHo (ChungJooHo)
- edida (edida)
- Francesco Franchina (ferdas)
- forbxy (forbxy)
- Grant Nestor (gnestor)
- Josh Barnes (jcb91)
- JocelynDelalande (JocelynDelalande)
- Karthik Balakrishnan (karthikb351)
- Kevin Bates (kevin-bates)
- Kirit Thadaka (kirit93)
- Lilian Besson (Naareen)
- Maarten Breddels (maartenbreddels)
- Madhu94 (Madhu94)
- Matthias Geier (mgeier)
- Michael Heilman (mheilman)
- Min RK (minrk)
- PHaeJin (PHaeJin)
- Sukneet (Sukneet)
- Thomas Kluyver (takluyver)

See the 5.3 milestone on GitHub for a complete list of [pull requests](#) involved in this release.

5.30 5.2.1

- Fix invisible CodeMirror cursor at specific browser zoom levels (2983).
- Fix nbconvert handler causing broken export to PDF (2981).
- Fix the prompt_area argument of the output area constructor. (2961).
- Handle a compound extension in newUntitled (2949).
- Allow disabling offline message buffering (2916).

Thanks to the following contributors:

- Steven Silvester ([blink1073](#))
- Grant Nestor ([gnestor](#))
- Jason Grout ([jasongrout](#))
- Min RK ([minrk](#))
- M Pacer ([mpacer](#))

See the 5.2.1 milestone on GitHub for a complete list of [pull requests](#) involved in this release.

5.31 5.2.0

- Allow setting token via jupyter_token env (2921).
- Fix some errors caused by raising 403 in get_current_user (2919).
- Register contents_manager.files_handler_class directly (2917).
- Update viewable_extensions (2913).
- Show edit shortcuts modal after shortcuts modal is hidden (2912).
- Improve edit/view behavior (2911).
- The root directory of the notebook server should never be hidden (2907).
- Fix notebook require config to match tools/build-main (2888).
- Give page constructor default arguments (2887).
- Fix codemirror.less to match codemirror's expected padding layout (2880).
- Add x-xsrftoken to access-control-allow-headers (2876).
- Buffer messages when websocket connection is interrupted (2871).
- Load locale dynamically only when not en-us (2866).
- Changed key strength to 2048 bits (2861).
- Resync jsversion with python version (2860).
- Allow copy operation on modified, read-only notebook (2854).
- Update error handling on apihandlers (2853).
- Test python 3.6 on travis, drop 3.3 (2852).
- Avoid base64-literals in image tests (2851).
- Upgrade xterm.js to 2.9.2 (2849).

- Changed all python variables named file to file_name to not override built_in file (2830).
- Add more doc tests (2823).
- Typos fix (2815).
- Rename and update license [ci skip] (2810).
- Travis builds doc (2808).
- Pull request i18n (2804).
- Factor out output_prompt_function, as is done with input prompt (2774).
- Use rfc5987 encoding for filenames (2767).
- Added path to the resources metadata, the same as in from_filename(...) in nbconvert.exporters.py (2753).
- Make “extrakeys” consistent for notebook and editor (2745).
- Bidi support (2357).

Special thanks to [samarsultan](#) and the Arabic Competence and Globalization Center Team at IBM Egypt for adding RTL (right-to-left) support to the notebook!

See the 5.2 milestone on GitHub for a complete list of [issues](#) and [pull requests](#) involved in this release.

5.32 5.1.0

- Preliminary i18n implementation (2140).
- Expose URL with auth token in notebook UI (2666).
- Fix search background style (2387).
- List running notebooks without requiring --allow-root (2421).
- Allow session of type other than notebook (2559).
- Fix search background style (2387).
- Fix some Markdown styling issues (2571), (2691) and (2534).
- Remove keymaps that conflict with non-English keyboards (2535).
- Add session-specific favicons (notebook, terminal, file) (2452).
- Add /api/shutdown handler (2507).
- Include metadata when copying a cell (2349).
- Stop notebook server from command line (2388).
- Improve “View” and “Edit” file handling in dashboard (2449) and (2402).
- Provide a promise to replace use of the app_initialized.NotebookApp event (2710).
- Fix disabled collapse/expand output button (2681).
- Cull idle kernels using --MappingKernelManager.cull_idle_timeout (2215).
- Allow read-only notebooks to be trusted (2718).

See the 5.1 milestone on GitHub for a complete list of [issues](#) and [pull requests](#) involved in this release.

5.33 5.0.0

This is the first major release of the Jupyter Notebook since version 4.0 was created by the “Big Split” of IPython and Jupyter.

We encourage users to start trying JupyterLab in preparation for a future transition.

We have merged more than 300 pull requests since 4.0. Some of the major user-facing changes are described here.

5.33.1 File sorting in the dashboard

Files in the dashboard may now be sorted by last modified date or name (943):

The screenshot shows the Jupyter dashboard interface. At the top left is the Jupyter logo and the word "jupyter". At the top right is a "Logout" button. Below the logo are tabs for "Files", "Running", "Clusters", and "Nbextensions". The "Files" tab is selected. Below the tabs is a message: "Select items to perform actions on them." To the right of this message are buttons for "Upload", "New" (with a dropdown arrow), and a refresh icon. Below this is a file browser view for the "/ Months" directory. It has a header with "Name" (sorted descending) and "Last Modified" (sorted ascending). The file list shows:

- .. (parent directory)
- March.ipynb (Running, 2 minutes ago)
- February.ipynb (Running, 2 minutes ago)
- January.ipynb (Running, 3 minutes ago)

 Each file has a checkbox to its left.

5.33.2 Cell tags

There is a new cell toolbar for adding *cell tags* (2048):

The screenshot shows three Jupyter Notebook cells. Each cell has a toolbar at the top with a "nbconvert-hide" button, a "nbval-ignore-output" button, a dropdown menu (indicated by "..."), and an "Add tag" button.

- Cell 1:** The code is `import numpy as np` and `print(np.__version__)`. The output is `1.12.0`.
- Cell 2:** The code is `a = np.arange(15).reshape(3, 5)` and `a`. The output is a 3x5 array: `array([[0, 1, 2, 3, 4], [5, 6, 7, 8, 9], [10, 11, 12, 13, 14]])`.
- Cell 3:** The code is `a.shape`. The output is `(3, 5)`.

 Each cell's toolbar is visible above its code input area.

Cell tags are a lightweight way to customise the behaviour of tools working with notebooks; we're working on building support for them into tools like `nbconvert` and `nbval`. To start using tags, select `Tags` in the `View > Cell Toolbar` menu in a notebook.

The UI for editing cell tags is basic for now; we hope to improve it in future releases.

5.33.3 Table style

The default styling for tables in the notebook has been updated (1776).

Before:

| | Acceleration | Cylinders | Displacement | Horsepower | Miles_per_Gallon | Name | Origin | Weight_in_lbs | Year |
|---|--------------|-----------|--------------|------------|------------------|---------------------------|--------|---------------|------------|
| 0 | 12.0 | 8 | 307.0 | 130.0 | 18.0 | chevrolet chevelle malibu | USA | 3504 | 1970-01-01 |
| 1 | 11.5 | 8 | 350.0 | 165.0 | 15.0 | buick skylark 320 | USA | 3693 | 1970-01-01 |
| 2 | 11.0 | 8 | 318.0 | 150.0 | 18.0 | plymouth satellite | USA | 3436 | 1970-01-01 |
| 3 | 12.0 | 8 | 304.0 | 150.0 | 16.0 | amc rebel sst | USA | 3433 | 1970-01-01 |
| 4 | 10.5 | 8 | 302.0 | 140.0 | 17.0 | ford torino | USA | 3449 | 1970-01-01 |
| 5 | 10.0 | 8 | 429.0 | 198.0 | 15.0 | ford galaxie 500 | USA | 4341 | 1970-01-01 |
| 6 | 9.0 | 8 | 454.0 | 220.0 | 14.0 | chevrolet impala | USA | 4354 | 1970-01-01 |
| 7 | 8.5 | 8 | 440.0 | 215.0 | 14.0 | plymouth fury iii | USA | 4312 | 1970-01-01 |
| 8 | 10.0 | 8 | 455.0 | 225.0 | 14.0 | pontiac catalina | USA | 4425 | 1970-01-01 |
| 9 | 8.5 | 8 | 390.0 | 190.0 | 15.0 | amc ambassador dpl | USA | 3850 | 1970-01-01 |

After:

| | Acceleration | Cylinders | Displacement | Horsepower | Miles_per_Gallon | Name | Origin | Weight_in_lbs | Year |
|---|--------------|-----------|--------------|------------|------------------|---------------------------|--------|---------------|------------|
| 0 | 12.0 | 8 | 307.0 | 130.0 | 18.0 | chevrolet chevelle malibu | USA | 3504 | 1970-01-01 |
| 1 | 11.5 | 8 | 350.0 | 165.0 | 15.0 | buick skylark 320 | USA | 3693 | 1970-01-01 |
| 2 | 11.0 | 8 | 318.0 | 150.0 | 18.0 | plymouth satellite | USA | 3436 | 1970-01-01 |
| 3 | 12.0 | 8 | 304.0 | 150.0 | 16.0 | amc rebel sst | USA | 3433 | 1970-01-01 |
| 4 | 10.5 | 8 | 302.0 | 140.0 | 17.0 | ford torino | USA | 3449 | 1970-01-01 |
| 5 | 10.0 | 8 | 429.0 | 198.0 | 15.0 | ford galaxie 500 | USA | 4341 | 1970-01-01 |
| 6 | 9.0 | 8 | 454.0 | 220.0 | 14.0 | chevrolet impala | USA | 4354 | 1970-01-01 |
| 7 | 8.5 | 8 | 440.0 | 215.0 | 14.0 | plymouth fury iii | USA | 4312 | 1970-01-01 |
| 8 | 10.0 | 8 | 455.0 | 225.0 | 14.0 | pontiac catalina | USA | 4425 | 1970-01-01 |
| 9 | 8.5 | 8 | 390.0 | 190.0 | 15.0 | amc ambassador dpl | USA | 3850 | 1970-01-01 |

5.33.4 Customise keyboard shortcuts

You can now edit keyboard shortcuts for *Command Mode* within the UI (1347):

Edit Command mode Shortcuts ×

Here you can modify the keyboard shortcuts available in command mode. Your changes will be stored for later sessions. See more [details of defining keyboard shortcuts](#) below.

| | | |
|--|---|---|
| toggle rtl layout | add shortcut | + |
| edit command mode keyboard shortcuts | add shortcut | + |
| shutdown kernel | add shortcut | + |
| confirm shutdown kernel | add shortcut | + |
| restart kernel | add shortcut | + |
| confirm restart kernel | <input type="text" value="0,0"/> add shortcut | + |
| restart kernel and run all cells | add shortcut | + |
| confirm restart kernel and run all cells | add shortcut | + |
| restart kernel and clear output | add shortcut | + |
| confirm restart kernel and clear output | add shortcut | + |
| interrupt kernel | <input type="text" value="I,I"/> add shortcut | + |

See the `Help > Edit Keyboard Shortcuts` menu item and follow the instructions.

5.33.5 Other additions

- You can copy and paste cells between notebooks, using `Ctrl-C` and `Ctrl-V` (Cmd-C and Cmd-V on Mac).
- It's easier to configure a password for the notebook with the new `jupyter notebook password command` (2007).
- The file list can now be ordered by *last modified* or by *name* (943).
- Markdown cells now support attachments. Simply drag and drop an image from your desktop to a markdown cell to add it. Unlike relative links that you enter manually, attachments are embedded in the notebook itself. An unreferenced attachment will be automatically scrubbed from the notebook on save (621).
- Undoing cell deletion now supports undeleting multiple cells. Cells may not be in the same order as before their deletion, depending on the actions you did on the meantime, but this should help reduce the impact of accidentally deleting code.
- The file browser now has *Edit* and *View* buttons.
- The file browser now supports moving multiple files at once (1088).
- The Notebook will refuse to run as root unless the `--allow-root` flag is given (1115).
- Keyboard shortcuts are now declarative (1234).
- Toggling line numbers can now affect all cells (1312).

- Add more visible *Trusted* and *Untrusted* notifications (1658).
- The favicon (browser shortcut icon) now changes to indicate when the kernel is busy (1837).
- Header and toolbar visibility is now persisted in nbconfig and across sessions (1769).
- Load server extensions with ConfigManager so that merge happens recursively, unlike normal config values, to make it load more consistently with frontend extensions(2108).
- The notebook server now supports the [bundler API](#) from the [jupyter_cms incubator project](#) (1579).
- The notebook server now provides information about kernel activity in its kernel resource API (1827).

Remember that upgrading `notebook` only affects the user interface. Upgrading kernels and libraries may also provide new features, better stability and integration with the notebook interface.

5.34 4.4.0

- Allow override of output callbacks to redirect output messages. This is used to implement the ipywidgets Output widget, for example.
- Fix an async bug in message handling by allowing comm message handlers to return a promise which halts message processing until the promise resolves.

See the 4.4 milestone on GitHub for a complete list of [issues](#) and [pull requests](#) involved in this release.

5.35 4.3.2

4.3.2 is a patch release with a bug fix for CodeMirror and improved handling of the “editable” cell metadata field.

- Monkey-patch for CodeMirror that resolves [#2037](#) without breaking [#1967](#)
- Read-only (`"editable": false`) cells can be executed but cannot be split, merged, or deleted

See the 4.3.2 milestone on GitHub for a complete list of [issues](#) and [pull requests](#) involved in this release.

5.36 4.3.1

4.3.1 is a patch release with a security patch, a couple bug fixes, and improvements to the newly-released token authentication.

Security fix:

- CVE-2016-9971. Fix CSRF vulnerability, where malicious forms could create untitled files and start kernels (no remote execution or modification of existing files) for users of certain browsers (Firefox, Internet Explorer / Edge). All previous notebook releases are affected.

Bug fixes:

- Fix carriage return handling
- Make the font size more robust against fickle browsers
- Ignore resize events that bubbled up and didn’t come from window
- Add Authorization to allowed CORS headers
- Downgrade CodeMirror to 5.16 while we figure out issues in Safari

Other improvements:

- Better docs for token-based authentication
- Further highlight token info in log output when autogenerated

See the 4.3.1 milestone on GitHub for a complete list of [issues](#) and [pull requests](#) involved in this release.

5.37 4.3.0

4.3 is a minor release with many bug fixes and improvements. The biggest user-facing change is the addition of token authentication, which is enabled by default. A token is generated and used when your browser is opened automatically, so you shouldn't have to enter anything in the default circumstances. If you see a login page (e.g. by switching browsers, or launching on a new port with `--no-browser`), you get a login URL with the token from the command `jupyter notebook list`, which you can paste into your browser.

Highlights:

- API for creating mime-type based renderer extensions using `OutputArea.register_mime_type` and `Notebook.render_cell_output` methods. See [mimerender-cookiecutter](#) for reference implementations and cookiecutter.
- Enable token authentication by default. See `server_security{.interpreted-text role="ref"}` for more details.
- Update security docs to reflect new signature system
- Switched from `term.js` to `xterm.js`

Bug fixes:

- Ensure variable is set if `exc_info` is falsey
- Catch and log handler exceptions in `events.trigger`
- Add debug log for static file paths
- Don't check origin on token-authenticated requests
- Remove leftover print statement
- Fix highlighting of Python code blocks
- `json_errors` should be outermost decorator on API handlers
- Fix remove old nbserver info files
- Fix notebook mime type on download links
- Fix carriage symbol behavior
- Fix terminal styles
- Update dead links in docs
- If kernel is broken, start a new session
- Include cross-origin check when allowing login URL redirects

Other improvements:

- Allow JSON output data with mime type `application/*+json`
- Allow kernelspecs to have spaces in them for backward compat
- Allow websocket connections from scripts

- Allow `None` for `post_save_hook`
- Upgrade CodeMirror to 5.21
- Upgrade xterm to 2.1.0
- Docs for using comms
- Set `dirty` flag when output arrives
- Set `ws-url` data attribute when accessing a notebook terminal
- Add base aliases for nbextensions
- Include `@` operator in CodeMirror IPython mode
- Extend `mathjax_url` docstring
- Load nbextension in predictable order
- Improve the error messages for nbextensions
- Include cross-origin check when allowing login URL redirects

See the 4.3 milestone on GitHub for a complete list of [issues](#) and [pull requests](#) involved in this release.

5.38 4.2.3

4.2.3 is a small bugfix release on 4.2.

Highlights:

- Fix regression in 4.2.2 that delayed loading `custom.js` until after `notebook_loaded` and `app_initialized` events have fired.
- Fix some outdated docs and links.

5.39 4.2.2

4.2.2 is a small bugfix release on 4.2, with an important security fix. All users are strongly encouraged to upgrade to 4.2.2.

Highlights:

- **Security fix:** CVE-2016-6524, where untrusted latex output could be added to the page in a way that could execute javascript.
- Fix missing POST in OPTIONS responses.
- Fix for downloading non-ascii filenames.
- Avoid clobbering `ssl_options`, so that users can specify more detailed SSL configuration.
- Fix inverted load order in `nbconfig`, so user config has highest priority.
- Improved error messages here and there.

5.40 4.2.1

4.2.1 is a small bugfix release on 4.2. Highlights:

- Compatibility fixes for some versions of ipywidgets
- Fix for ignored CSS on Windows
- Fix specifying destination when installing nbextensions

5.41 4.2.0

Release 4.2 adds a new API for enabling and installing extensions. Extensions can now be enabled at the system-level, rather than just per-user. An API is defined for installing directly from a Python package, as well.

Highlighted changes:

- Upgrade MathJax to 2.6 to fix vertical-bar appearing on some equations.
- Restore ability for notebook directory to be root (4.1 regression)
- Large outputs are now throttled, reducing the ability of output floods to kill the browser.
- Fix the notebook ignoring cell executions while a kernel is starting by queueing the messages.
- Fix handling of url prefixes (e.g. JupyterHub) in terminal and edit pages.
- Support nested SVGs in output.

And various other fixes and improvements.


5.42 4.1.0

Bug fixes:

- Properly reap zombie subprocesses
- Fix cross-origin problems
- Fix double-escaping of the base URL prefix
- Handle invalid unicode filenames more gracefully
- Fix ANSI color-processing
- Send keepalive messages for web terminals
- Fix bugs in the notebook tour

UI changes:

- Moved the cell toolbar selector into the *View* menu. Added a button that triggers a “hint” animation to the main toolbar so users can find the new location. (Click here to see a [screencast](#))



- Added *Restart & Run All* to the *Kernel* menu. Users can also bind it to a keyboard shortcut on action `restart-kernel-and-run-all-cells`.
- Added multiple-cell selection. Users press `Shift-Up/Down` or `Shift-K/J` to extend selection in command mode. Various actions such as cut/copy/paste, execute, and cell type conversions apply to all selected cells.

Code cells allow you to enter and run code

Run a code cell using `Shift-Enter` or pressing the button in the toolbar above:


```
In [ ]: a = 10
```

```
In [ ]: print(a)
```


There are two other keyboard shortcuts for running code:

- `Alt-Enter` runs the current cell and inserts a new one below.
- `Ctrl-Enter` run the current cell and enters command mode.

- Added a command palette for executing Jupyter actions by name. Users press `Cmd/Ctrl-Shift-P` or click the new command palette icon on the toolbar.



- Added a *Find and Replace* dialog to the *Edit* menu. Users can also press `F` in command mode to show the dialog.



Other improvements:

- Custom `KernelManager` methods can be Tornado coroutines, allowing async operations.
- Make clearing output optional when rewriting input with `set_next_input(replace=True)`.
- Added support for TLS client authentication via `--NotebookApp.client-ca`.
- Added tags to `jupyter/notebook` releases on [DockerHub](#). `latest` continues to track the master branch.

See the 4.1 milestone on [GitHub](#) for a complete list of [issues](#) and [pull requests](#) handled.

5.43 4.0.x

5.43.1 4.0.6

- fix installation of mathjax support files
- fix some double-escape regressions in 4.0.5
- fix a couple of cases where errors could prevent opening a notebook

5.43.2 4.0.5

Security fixes for maliciously crafted files.

- [CVE-2015-6938](#): malicious filenames
- [CVE-2015-7337](#): malicious binary files in text editor.

Thanks to Jonathan Kamens at Quantopian and Juan Broullón for the reports.

5.43.3 4.0.4

- Fix inclusion of mathjax-safe extension

5.43.4 4.0.2

- Fix launching the notebook on Windows
- Fix the path searched for frontend config

5.43.5 4.0.0

First release of the notebook as a standalone package.

COMMS

Comms allow custom messages between the frontend and the kernel. They are used, for instance, in [ipywidgets](#) to update widget state.

A comm consists of a pair of objects, in the kernel and the frontend, with an automatically assigned unique ID. When one side sends a message, a callback on the other side is triggered with that message data. Either side, the frontend or kernel, can open or close the comm.

See also:

Custom Messages The messaging specification section on comms

6.1 Opening a comm from the kernel

First, the function to accept the comm must be available on the frontend. This can either be specified in a *requires* module, or registered in a registry, for example when an *extension* is loaded. This example shows a frontend comm target registered in a registry:

```
Jupyter.notebook.kernel.comm_manager.register_target('my_comm_target',
    function(comm, msg) {
        // comm is the frontend comm instance
        // msg is the comm_open message, which can carry data

        // Register handlers for later messages:
        comm.on_msg(function(msg) {...});
        comm.on_close(function(msg) {...});
        comm.send({'foo': 0});
    });
```

Now that the frontend comm is registered, you can open the comm from the kernel:

```
from ipykernel.comm import Comm

# Use comm to send a message from the kernel
my_comm = Comm(target_name='my_comm_target', data={'foo': 1})
my_comm.send({'foo': 2})

# Add a callback for received messages.
@my_comm.on_msg
def _recv(msg):
    # Use msg['content']['data'] for the data in the message
```

This example uses the IPython kernel; it's up to each language kernel what API, if any, it offers for using comms.

6.2 Opening a comm from the frontend

This is very similar to above, but in reverse. First, a comm target must be registered in the kernel. For instance, this may be done by code displaying output: it will register a target in the kernel, and then display output containing Javascript to connect to it.

```
def target_func(comm, open_msg):
    # comm is the kernel Comm instance
    # msg is the comm_open message

    # Register handler for later messages
    @comm.on_msg
    def _recv(msg):
        # Use msg['content']['data'] for the data in the message
        comm.send({'echo': msg['content']['data']})

    # Send data to the frontend on creation
    comm.send({'foo': 5})

get_ipython().kernel.comm_manager.register_target('my_comm_target', target_func)
```

This example uses the IPython kernel again; this example will be different in other kernels that support comms. Refer to the specific language kernel’s documentation for comms support.

And then open the comm from the frontend:

```
const comm = Jupyter.notebook.kernel.comm_manager.new_comm('my_comm_target', {'foo': 6})
// Send data
comm.send({'foo': 7})

// Register a handler
comm.on_msg(function(msg) {
    console.log(msg.content.data.foo);
});
```


CONFIGURATION OVERVIEW

Beyond the default configuration settings, you can configure a rich array of options to suit your workflow. Here are areas that are commonly configured when using Jupyter Notebook:

- *Jupyter’s common configuration system*
- *Notebook server*
- *Notebook front-end client*
- *Notebook extensions*

Let’s look at highlights of each area.

7.1 Jupyter’s Common Configuration system

Jupyter applications, from the Notebook to JupyterHub to nbgrader, share a common configuration system. The process for creating a configuration file and editing settings is similar for all the Jupyter applications.

- [Jupyter’s Common Configuration Approach](#)
- [Common Directories and File Locations](#)
- [Language kernels](#)
- [traitlets](#) provide a low-level architecture for configuration.

7.2 Notebook server

The Notebook server runs the language kernel and communicates with the front-end Notebook client (i.e. the familiar notebook interface).

- **Configuring the Notebook server**

To create a `jupyter_notebook_config.py` file in the `.jupyter` directory, with all the defaults commented out, use the following command:

```
$ jupyter notebook --generate-config
```

:ref:`Command line arguments for configuration <config>` settings are documented in the configuration file and the user documentation.

- *Running a Notebook server*

- Related: [Configuring a language kernel](#) to run in the Notebook server enables your server to run other languages, like R or Julia.

7.3 Notebook front-end client

7.3.1 Configuring the notebook frontend

Note: The ability to configure the notebook frontend UI and preferences is still a work in progress.

This document is a rough explanation on how you can persist some configuration options for the notebook JavaScript. There is no exhaustive list of all the configuration options as most options are passed down to other libraries, which means that non valid configuration can be ignored without any error messages.

How front end configuration works

The frontend configuration system works as follows:

- get a handle of a configurable JavaScript object.
- access its configuration attribute.
- update its configuration attribute with a JSON patch.

Example - Changing the notebook's default indentation

This example explains how to change the default setting `indentUnit` for CodeMirror Code Cells:

```
var cell = Jupyter.notebook.get_selected_cell();
var config = cell.config;
var patch = {
  CodeCell: {
    cm_config: { indentUnit: 2 }
  }
}
config.update(patch)
```

You can enter the previous snippet in your browser's JavaScript console once. Then reload the notebook page in your browser. Now, the preferred indent unit should be equal to two spaces. The custom setting persists and you do not need to reissue the patch on new notebooks.

`indentUnit`, used in this example, is one of the many [CodeMirror options](#) which are available for configuration.

You can similarly change the options of the file editor by entering the following snippet in the browser's Javascript console once (from a file editing page).:

```
var config = Jupyter.editor.config
var patch = {
  Editor: {
    codemirror_options: {
      indentUnit: 2
    }
  }
}
```

(continues on next page)

(continued from previous page)

```
}  
config.update(patch)
```

Example - Restoring the notebook's default indentation

If you want to restore a notebook frontend preference to its default value, you will enter a JSON patch with a `null` value for the preference setting.

For example, let's restore the indent setting `indentUnit` to its default of four spaces. Enter the following code snippet in your JavaScript console:

```
var cell = Jupyter.notebook.get_selected_cell();  
var config = cell.config;  
var patch = {  
  CodeCell: {  
    cm_config: { indentUnit: null } // only change here.  
  }  
}  
config.update(patch)
```

Reload the notebook in your browser and the default indent should again be two spaces.

Persisting configuration settings

Under the hood, Jupyter will persist the preferred configuration settings in `~/.jupyter/nbconfig/<section>.json`, with `<section>` taking various value depending on the page where the configuration is issued. `<section>` can take various values like `notebook`, `tree`, and `editor`. A common section contains configuration settings shared by all pages.

7.4 Notebook extensions

- [Distributing Jupyter Extensions as Python Packages](#)
- [Extending the Notebook](#)

Security in Jupyter notebooks: Since security policies vary from organization to organization, we encourage you to consult with your security team on settings that would be best for your use cases. Our documentation offers some responsible security practices, and we recommend becoming familiar with the practices.

CONFIG FILE AND COMMAND LINE OPTIONS

The notebook server can be run with a variety of command line arguments. A list of available options can be found below in the *options section*.

Defaults for these options can also be set by creating a file named `jupyter_notebook_config.py` in your Jupyter folder. The Jupyter folder is in your home directory, `~/ .jupyter`.

To create a `jupyter_notebook_config.py` file, with all the defaults commented out, you can use the following command line:

```
$ jupyter notebook --generate-config
```

8.1 Options

This list of options can be generated by running the following and hitting enter:

```
$ jupyter notebook --help
```

Application.log_datefmt [Unicode] Default: '%Y-%m-%d %H:%M:%S'

The date format used by logging formatters for `%(asctime)s`

Application.log_format [Unicode] Default: '%(name)s %(levelname)s %(message)s'

The Logging format template

Application.log_level [any of 0, 10, 20, 30, 40, 50, 'DEBUG', 'INFO', 'WARN', 'ERROR', 'CRITICAL']
Default: 30

Set the log level by value or name.

Application.show_config [Bool] Default: False

Instead of starting the Application, dump configuration to stdout

Application.show_config_json [Bool] Default: False

Instead of starting the Application, dump configuration to stdout (as JSON)

JupyterApp.answer_yes [Bool] Default: False

Answer yes to any prompts.

JupyterApp.config_file [Unicode] Default: ''

Full path of a config file.

JupyterApp.config_file_name [Unicode] Default: ''

Specify a config file to load.

JupyterApp.generate_config [Bool] Default: False

Generate default config file.

JupyterApp.log_datefmt [Unicode] Default: '%Y-%m-%d %H:%M:%S'

The date format used by logging formatters for %(asctime)s

JupyterApp.log_format [Unicode] Default: '%(name)s %(levelname)s %(message)s'

The Logging format template

JupyterApp.log_level [any of 0 `` | ``10 `` | ``20 `` | ``30 `` | ``40 `` | ``50 `` | 'DEBUG' | 'INFO' | 'WARN' | 'ERROR' | 'CRITICAL']

Default: 30

Set the log level by value or name.

JupyterApp.show_config [Bool] Default: False

Instead of starting the Application, dump configuration to stdout

JupyterApp.show_config_json [Bool] Default: False

Instead of starting the Application, dump configuration to stdout (as JSON)

NotebookApp.allow_credentials [Bool] Default: False

Set the Access-Control-Allow-Credentials: true header

NotebookApp.allow_origin [Unicode] Default: ''

Set the Access-Control-Allow-Origin header

Use '*' to allow any origin to access your server.

Takes precedence over allow_origin_pat.

NotebookApp.allow_origin_pat [Unicode] Default: ''

Use a regular expression for the Access-Control-Allow-Origin header

Requests from an origin matching the expression will get replies with:

Access-Control-Allow-Origin: origin

where *origin* is the origin of the request.

Ignored if allow_origin is set.

NotebookApp.allow_password_change [Bool] Default: True

Allow password to be changed at login for the notebook server.

While logging in with a token, the notebook server UI will give the opportunity to the user to enter a new password at the same time that will replace the token login mechanism.

This can be set to false to prevent changing password from the UI/API.

NotebookApp.allow_remote_access [Bool] Default: False

Allow requests where the Host header doesn't point to a local server

By default, requests get a 403 forbidden response if the 'Host' header shows that the browser thinks it's on a non-local domain. Setting this option to True disables this check.

This protects against ‘DNS rebinding’ attacks, where a remote web server serves you a page and then changes its DNS to send later requests to a local IP, bypassing same-origin checks.

Local IP addresses (such as 127.0.0.1 and ::1) are allowed as local, along with hostnames configured in `local_hostnames`.

NotebookApp.allow_root [Bool] Default: `False`

Whether to allow the user to run the notebook as root.

NotebookApp.answer_yes [Bool] Default: `False`

Answer yes to any prompts.

NotebookApp.authenticate_prometheus [Bool] Default: `True`

” Require authentication to access prometheus metrics.

NotebookApp.autoreload [Bool] Default: `False`

Reload the webapp when changes are made to any Python src files.

NotebookApp.base_project_url [Unicode] Default: `'/'`

DEPRECATED use `base_url`

NotebookApp.base_url [Unicode] Default: `'/'`

The base URL for the notebook server.

Leading and trailing slashes can be omitted, and will automatically be added.

NotebookApp.browser [Unicode] Default: `''`

Specify what command to use to invoke a web browser when opening the notebook. If not specified, the default browser will be determined by the *webbrowser* standard library module, which allows setting of the `BROWSER` environment variable to override it.

NotebookApp.certfile [Unicode] Default: `''`

The full path to an SSL/TLS certificate file.

NotebookApp.client_ca [Unicode] Default: `''`

The full path to a certificate authority certificate for SSL/TLS client authentication.

NotebookApp.config_file [Unicode] Default: `''`

Full path of a config file.

NotebookApp.config_file_name [Unicode] Default: `''`

Specify a config file to load.

NotebookApp.config_manager_class [Type] Default: `'notebook.services.config.manager.ConfigManager'`

The config manager class to use

NotebookApp.contents_manager_class [TypeFromClasses] Default: `'notebook.services.contents.largefilemanager.LargeFileManager'`

The notebook manager class to use.

NotebookApp.cookie_options [Dict] Default: `{}`

Extra keyword arguments to pass to *set_secure_cookie*. See tornado’s *set_secure_cookie* docs for details.

NotebookApp.cookie_secret [Bytes] Default: `b''`

The random bytes used to secure cookies. By default this is a new random number every time you start the Notebook. Set it to a value in a config file to enable logins to persist across server sessions.

Note: Cookie secrets should be kept private, do not share config files with cookie_secret stored in plaintext (you can read the value from a file).

NotebookApp.cookie_secret_file [Unicode] Default: `''`

The file where the cookie secret is stored.

NotebookApp.custom_display_url [Unicode] Default: `''`

Override URL shown to users.

Replace actual URL, including protocol, address, port and base URL, with the given value when displaying URL to the users. Do not change the actual connection URL. If authentication token is enabled, the token is added to the custom URL automatically.

This option is intended to be used when the URL to display to the user cannot be determined reliably by the Jupyter notebook server (proxified or containerized setups for example).

NotebookApp.default_url [Unicode] Default: `' /tree '`

The default URL to redirect to from /

NotebookApp.disable_check_xsrf [Bool] Default: `False`

Disable cross-site-request-forgery protection

Jupyter notebook 4.3.1 introduces protection from cross-site request forgeries, requiring API requests to either:

- originate from pages served by this server (validated with XSRF cookie and token), or
- authenticate with a token

Some anonymous compute resources still desire the ability to run code, completely without authentication. These services can disable all authentication and security checks, with the full knowledge of what that implies.

NotebookApp.enable_mathjax [Bool] Default: `True`

Whether to enable MathJax for typesetting math/TeX

MathJax is the javascript library Jupyter uses to render math/LaTeX. It is very large, so you may want to disable it if you have a slow internet connection, or for offline use of the notebook.

When disabled, equations etc. will appear as their untransformed TeX source.

NotebookApp.extra_nbextensions_path [List] Default: `[]`

extra paths to look for Javascript notebook extensions

NotebookApp.extra_services [List] Default: `[]`

handlers that should be loaded at higher priority than the default services

NotebookApp.extra_static_paths [List] Default: `[]`

Extra paths to search for serving static files.

This allows adding javascript/css to be available from the notebook server machine, or overriding individual files in the IPython

NotebookApp.extra_template_paths [List] Default: `[]`

Extra paths to search for serving jinja templates.

Can be used to override templates from notebook.templates.

NotebookApp.file_to_run [Unicode] Default: ''

No description

NotebookApp.generate_config [Bool] Default: False

Generate default config file.

NotebookApp.get_secure_cookie_kwargs [Dict] Default: {}

Extra keyword arguments to pass to *get_secure_cookie*. See tornado's *get_secure_cookie* docs for details.

NotebookApp.ignore_minified_js [Bool] Default: False

Deprecated: Use minified JS file or not, mainly use during dev to avoid JS recompilation

NotebookApp.iopub_data_rate_limit [Float] Default: 1000000

(bytes/sec) Maximum rate at which stream output can be sent on iopub before they are limited.

NotebookApp.iopub_msg_rate_limit [Float] Default: 1000

(msgs/sec) Maximum rate at which messages can be sent on iopub before they are limited.

NotebookApp.ip [Unicode] Default: 'localhost'

The IP address the notebook server will listen on.

NotebookApp.jinja_environment_options [Dict] Default: {}

Supply extra arguments that will be passed to Jinja environment.

NotebookApp.jinja_template_vars [Dict] Default: {}

Extra variables to supply to jinja templates when rendering.

NotebookApp.kernel_manager_class [Type] Default: 'notebook.services.kernels.kernelmanager.MappingKernelManager'

The kernel manager class to use.

NotebookApp.kernel_spec_manager_class [Type] Default: 'jupyter_client.kernelspec.KernelSpecManager'

The kernel spec manager class to use. Should be a subclass of *jupyter_client.kernelspec.KernelSpecManager*.

The Api of *KernelSpecManager* is provisional and might change without warning between this version of Jupyter and the next stable one.

NotebookApp.keyfile [Unicode] Default: ''

The full path to a private key file for usage with SSL/TLS.

NotebookApp.local_hostnames [List] Default: ['localhost']

Hostnames to allow as local when *allow_remote_access* is False.

Local IP addresses (such as 127.0.0.1 and ::1) are automatically accepted as local as well.

NotebookApp.log_datefmt [Unicode] Default: '%Y-%m-%d %H:%M:%S'

The date format used by logging formatters for *%(asctime)s*

NotebookApp.log_format [Unicode] Default: '[(name)s]%(highlevel)s %(message)s'

The Logging format template

NotebookApp.log_json [Bool] Default: `False`

Set to `True` to enable JSON formatted logs. Run “`pip install notebook[json-logging]`” to install the required dependent packages. Can also be set using the environment variable `JUPYTER_ENABLE_JSON_LOGGING=true`.

NotebookApp.log_level [any of `0` | `10` | `20` | `30` | `40` | `50` | `'DEBUG'` | `'INFO'` | `'WARN'` | `'ERROR'` | `'CRITICAL'`] Default: `30`

Set the log level by value or name.

NotebookApp.login_handler_class [Type] Default: `'notebook.auth.login.LoginHandler'`

The login handler class to use.

NotebookApp.logout_handler_class [Type] Default: `'notebook.auth.logout.LogoutHandler'`

The logout handler class to use.

NotebookApp.mathjax_config [Unicode] Default: `'TeX-AMS-MML_HTMLorMML-full, Safe'`

The MathJax.js configuration file that is to be used.

NotebookApp.mathjax_url [Unicode] Default: `''`

A custom url for MathJax.js. Should be in the form of a case-sensitive url to MathJax, for example: `/static/components/MathJax/MathJax.js`

NotebookApp.max_body_size [Int] Default: `536870912`

Sets the maximum allowed size of the client request body, specified in the Content-Length request header field. If the size in a request exceeds the configured value, a malformed HTTP message is returned to the client.

Note: `max_body_size` is applied even in streaming mode.

NotebookApp.max_buffer_size [Int] Default: `536870912`

Gets or sets the maximum amount of memory, in bytes, that is allocated for use by the buffer manager.

NotebookApp.min_open_files_limit [Int] Default: `0`

Gets or sets a lower bound on the open file handles process resource limit. This may need to be increased if you run into an `OSError: [Errno 24] Too many open files`. This is not applicable when running on Windows.

NotebookApp.nbserver_extensions [Dict] Default: `{}`

Dict of Python modules to load as notebook server extensions. Entry values can be used to enable and disable the loading of the extensions. The extensions will be loaded in alphabetical order.

NotebookApp.notebook_dir [Unicode] Default: `''`

The directory to use for notebooks and kernels.

NotebookApp.open_browser [Bool] Default: `True`

Whether to open in a browser after starting. The specific browser used is platform dependent and determined by the python standard library `webbrowser` module, unless it is overridden using the `--browser` (`NotebookApp.browser`) configuration option.

NotebookApp.password [Unicode] Default: `''`

Hashed password to use for web authentication.

To generate, type in a python/IPython shell:

```
from notebook.auth import passwd; passwd()
```

The string should be of the form `type:salt:hashed-password`.

NotebookApp.password_required [Bool] Default: `False`

Forces users to use a password for the Notebook server. This is useful in a multi user environment, for instance when everybody in the LAN can access each other's machine through ssh.

In such a case, serving the notebook server on localhost is not secure since any user can connect to the notebook server via ssh.

NotebookApp.port [Int] Default: `8888`

The port the notebook server will listen on (env: `JUPYTER_PORT`).

NotebookApp.port_retries [Int] Default: `50`

The number of additional ports to try if the specified port is not available (env: `JUPYTER_PORT_RETRIES`).

NotebookApp.pylab [Unicode] Default: `'disabled'`

DISABLED: use `%pylab` or `%matplotlib` in the notebook to enable matplotlib.

NotebookApp.quit_button [Bool] Default: `True`

If `True`, display a button in the dashboard to quit (shutdown the notebook server).

NotebookApp.rate_limit_window [Float] Default: `3`

(sec) Time window used to check the message and data rate limits.

NotebookApp.reraise_server_extension_failures [Bool] Default: `False`

Reraise exceptions encountered loading server extensions?

NotebookApp.server_extensions [List] Default: `[]`

DEPRECATED use the `nbserver_extensions` dict instead

NotebookApp.session_manager_class [Type] Default: `'notebook.services.sessions.sessionmanager.SessionManager'`

The session manager class to use.

NotebookApp.show_config [Bool] Default: `False`

Instead of starting the Application, dump configuration to stdout

NotebookApp.show_config_json [Bool] Default: `False`

Instead of starting the Application, dump configuration to stdout (as JSON)

NotebookApp.shutdown_no_activity_timeout [Int] Default: `0`

Shut down the server after `N` seconds with no kernels or terminals running and no activity. This can be used together with culling idle kernels (`MappingKernelManager.cull_idle_timeout`) to shutdown the notebook server when it's not in use. This is not precisely timed: it may shut down up to a minute later. `0` (the default) disables this automatic shutdown.

NotebookApp.sock [Unicode] Default: `''`

The UNIX socket the notebook server will listen on.

NotebookApp.sock_mode [Unicode] Default: `'0600'`

The permissions mode for UNIX socket creation (default: `0600`).

NotebookApp.ssl_options [Dict] Default: `{}`

Supply SSL options for the tornado `HTTPServer`. See the tornado docs for details.

NotebookApp.terminado_settings [Dict] Default: {}

Supply overrides for terminado. Currently only supports “shell_command”. On Unix, if “shell_command” is not provided, a non-login shell is launched by default when the notebook server is connected to a terminal, a login shell otherwise.

NotebookApp.terminals_enabled [Bool] Default: True

Set to False to disable terminals.

This does *not* make the notebook server more secure by itself. Anything the user can in a terminal, they can also do in a notebook.

Terminals may also be automatically disabled if the terminado package is not available.

NotebookApp.token [Unicode] Default: '<generated>'

Token used for authenticating first-time connections to the server.

The token can be read from the file referenced by JUPYTER_TOKEN_FILE or set directly with the JUPYTER_TOKEN environment variable.

When no password is enabled, the default is to generate a new, random token.

Setting to an empty string disables authentication altogether, which is NOT RECOMMENDED.

NotebookApp.tornado_settings [Dict] Default: {}

Supply overrides for the tornado.web.Application that the Jupyter notebook uses.

NotebookApp.trust_xheaders [Bool] Default: False

Whether to trust or not X-Scheme/X-Forwarded-Proto and X-Real-Ip/X-Forwarded-For headers sent by the upstream reverse proxy. Necessary if the proxy handles SSL

NotebookApp.use_redirect_file [Bool] Default: True

Disable launching browser by redirect file

For versions of notebook > 5.7.2, a security feature measure was added that prevented the authentication token used to launch the browser from being visible. This feature makes it difficult for other users on a multi-user system from running code in your Jupyter session as you.

However, some environments (like Windows Subsystem for Linux (WSL) and Chromebooks), launching a browser using a redirect file can lead the browser failing to load. This is because of the difference in file structures/paths between the runtime and the browser.

Disabling this setting to False will disable this behavior, allowing the browser to launch by using a URL and visible token (as before).

NotebookApp.webapp_settings [Dict] Default: {}

DEPRECATED, use tornado_settings

NotebookApp.webbrowser_open_new [Int] Default: 2

Specify Where to open the notebook on startup. This is the *new* argument passed to the standard library method *webbrowser.open*. The behaviour is not guaranteed, but depends on browser support. Valid values are:

- 2 opens a new tab,
- 1 opens a new window,
- 0 opens in an existing window.

See the *webbrowser.open* documentation for details.

NotebookApp.websocket_compression_options [Any] Default: None

Set the tornado compression options for websocket connections.

This value will be returned from `WebSocketHandler.get_compression_options()`. None (default) will disable compression. A dict (even an empty one) will enable compression.

See the tornado docs for `WebSocketHandler.get_compression_options` for details.

NotebookApp.websocket_url [Unicode] Default: ''

The base URL for websockets, if it differs from the HTTP server (hint: it almost certainly doesn't).

Should be in the form of an HTTP origin: `ws[s]://hostname[:port]`

ConnectionFileMixin.connection_file [Unicode] Default: ''

JSON file in which to store connection info [default: `kernel-<pid>.json`]

This file will contain the IP, ports, and authentication key needed to connect clients to this kernel. By default, this file will be created in the security dir of the current profile, but can be specified by absolute path.

ConnectionFileMixin.control_port [Int] Default: 0

set the control (ROUTER) port [default: random]

ConnectionFileMixin.hb_port [Int] Default: 0

set the heartbeat port [default: random]

ConnectionFileMixin.iopub_port [Int] Default: 0

set the iopub (PUB) port [default: random]

ConnectionFileMixin.ip [Unicode] Default: ''

Set the kernel's IP address [default localhost]. If the IP address is something other than localhost, then Consoles on other machines will be able to connect to the Kernel, so be careful!

ConnectionFileMixin.shell_port [Int] Default: 0

set the shell (ROUTER) port [default: random]

ConnectionFileMixin.stdin_port [Int] Default: 0

set the stdin (ROUTER) port [default: random]

ConnectionFileMixin.transport [any of 'tcp' | 'ipc' (case-insensitive)] Default: 'tcp'

No description

KernelManager.autorestart [Bool] Default: True

Should we autorestart the kernel if it dies.

KernelManager.connection_file [Unicode] Default: ''

JSON file in which to store connection info [default: `kernel-<pid>.json`]

This file will contain the IP, ports, and authentication key needed to connect clients to this kernel. By default, this file will be created in the security dir of the current profile, but can be specified by absolute path.

KernelManager.control_port [Int] Default: 0

set the control (ROUTER) port [default: random]

KernelManager.hb_port [Int] Default: 0

set the heartbeat port [default: random]

KernelManager.iopub_port [Int] Default: 0

set the iopub (PUB) port [default: random]

KernelManager.ip [Unicode] Default: ''

Set the kernel's IP address [default localhost]. If the IP address is something other than localhost, then Consoles on other machines will be able to connect to the Kernel, so be careful!

KernelManager.kernel_cmd [List] Default: []

DEPRECATED: Use kernel_name instead.

The Popen Command to launch the kernel. Override this if you have a custom kernel. If kernel_cmd is specified in a configuration file, Jupyter does not pass any arguments to the kernel, because it cannot make any assumptions about the arguments that the kernel understands. In particular, this means that the kernel does not receive the option `--debug` if it given on the Jupyter command line.

KernelManager.shell_port [Int] Default: 0

set the shell (ROUTER) port [default: random]

KernelManager.shutdown_wait_time [Float] Default: 5.0

Time to wait for a kernel to terminate before killing it, in seconds. When a shutdown request is initiated, the kernel will be immediately send and interrupt (SIGINT), followed by a shutdown_request message, after 1/2 of *shutdown_wait_time* it will be sent a terminate (SIGTERM) request, and finally at the end of *shutdown_wait_time* will be killed (SIGKILL). terminate and kill may be equivalent on windows.

KernelManager.stdin_port [Int] Default: 0

set the stdin (ROUTER) port [default: random]

KernelManager.transport [any of 'tcp' or 'ipc' (case-insensitive)] Default: 'tcp'

No description

Session.buffer_threshold [Int] Default: 1024

Threshold (in bytes) beyond which an object's buffer should be extracted to avoid pickling.

Session.check_pid [Bool] Default: True

Whether to check PID to protect against calls after fork.

This check can be disabled if fork-safety is handled elsewhere.

Session.copy_threshold [Int] Default: 65536

Threshold (in bytes) beyond which a buffer should be sent without copying.

Session.debug [Bool] Default: False

Debug output in the Session

Session.digest_history_size [Int] Default: 65536

The maximum number of digests to remember.

The digest history will be culled when it exceeds this value.

Session.item_threshold [Int] Default: 64

The maximum number of items for a container to be introspected for custom serialization. Containers larger than this are pickled outright.

Session.key [CBytes] Default: b''

execution key, for signing messages.

Session.keyfile [Unicode] Default: ''

path to file containing execution key.

Session.metadata [Dict] Default: {}

Metadata dictionary, which serves as the default top-level metadata dict for each message.

Session.packer [DottedObjectName] Default: 'json'

The name of the packer for serializing messages. Should be one of 'json', 'pickle', or an import name for a custom callable serializer.

Session.session [Unicode] Default: ''

The UUID identifying this session.

Session.signature_scheme [Unicode] Default: 'hmac-sha256'

The digest scheme used to construct the message signatures. Must have the form 'hmac-HASH'.

Session.unpacker [DottedObjectName] Default: 'json'

The name of the unpacker for unserializing messages. Only used with custom functions for *packer*.

Session.username [Unicode] Default: 'username'

Username for the Session. Default is your system username.

MultiKernelManager.default_kernel_name [Unicode] Default: 'python3'

The name of the default kernel to start

MultiKernelManager.kernel_manager_class [DottedObjectName] Default: 'jupyter_client.ioloop.IOLoopKernelManager'

The kernel manager class. This is configurable to allow subclassing of the `KernelManager` for customized behavior.

MultiKernelManager.shared_context [Bool] Default: `True`

Share a single `zmq.Context` to talk to all my kernels

MappingKernelManager.allowed_message_types [List] Default: []

White list of allowed kernel message types. When the list is empty, all message types are allowed.

MappingKernelManager.buffer_offline_messages [Bool] Default: `True`

Whether messages from kernels whose frontends have disconnected should be buffered in-memory. When `True` (default), messages are buffered and replayed on reconnect, avoiding lost messages due to interrupted connectivity. Disable if long-running kernels will produce too much output while no frontends are connected.

MappingKernelManager.cull_busy [Bool] Default: `False`

Whether to consider culling kernels which are busy. Only effective if `cull_idle_timeout > 0`.

MappingKernelManager.cull_connected [Bool] Default: `False`

Whether to consider culling kernels which have one or more connections. Only effective if `cull_idle_timeout > 0`.

MappingKernelManager.cull_idle_timeout [Int] Default: 0

Timeout (in seconds) after which a kernel is considered idle and ready to be culled. Values of 0 or lower disable culling. Very short timeouts may result in kernels being culled for users with poor network connections.

MappingKernelManager.cull_interval [Int] Default: 300

The interval (in seconds) on which to check for idle kernels exceeding the cull timeout value.

MappingKernelManager.default_kernel_name [Unicode] Default: 'python3'

The name of the default kernel to start

MappingKernelManager.kernel_info_timeout [Float] Default: 60

Timeout for giving up on a kernel (in seconds). On starting and restarting kernels, we check whether the kernel is running and responsive by sending kernel_info_requests. This sets the timeout in seconds for how long the kernel can take before being presumed dead. This affects the MappingKernelManager (which handles kernel restarts) and the ZMQChannelsHandler (which handles the startup).

MappingKernelManager.kernel_manager_class [DottedObjectName] Default: 'jupyter_client.ioloop.IOLoopKernelManager'

The kernel manager class. This is configurable to allow subclassing of the KernelManager for customized behavior.

MappingKernelManager.root_dir [Unicode] Default: ''

No description

MappingKernelManager.shared_context [Bool] Default: True

Share a single zmq.Context to talk to all my kernels

KernelSpecManager.ensure_native_kernel [Bool] Default: True

If there is no Python kernelspec registered and the IPython kernel is available, ensure it is added to the spec list.

KernelSpecManager.kernel_spec_class [Type] Default: 'jupyter_client.kernelspec.KernelSpec'

The kernel spec class. This is configurable to allow subclassing of the KernelSpecManager for customized behavior.

KernelSpecManager.whitelist [Set] Default: set()

Whitelist of allowed kernel names.

By default, all installed kernels are allowed.

ContentsManager.allow_hidden [Bool] Default: False

Allow access to hidden files

ContentsManager.checkpoints [Instance] Default: None

No description

ContentsManager.checkpoints_class [Type] Default: 'notebook.services.contents.checkpoints.Checkpoints'

No description

ContentsManager.checkpoints_kwargs [Dict] Default: {}

No description

ContentsManager.files_handler_class [Type] Default: 'notebook.files.handlers.FilesHandler'

handler class to use when serving raw file requests.

Default is a fallback that talks to the ContentsManager API, which may be inefficient, especially for large files.

Local files-based ContentsManagers can use a StaticFileHandler subclass, which will be much more efficient.

Access to these files should be Authenticated.

ContentsManager.files_handler_params [Dict] Default: { }

Extra parameters to pass to files_handler_class.

For example, StaticFileHandlers generally expect a *path* argument specifying the root directory from which to serve files.

ContentsManager.hide_globs [List] Default: ['__pycache__', '*.pyc', '*.pyo', '.DS_Store', '*.so', '*.dyl...'

Glob patterns to hide in file and directory listings.

ContentsManager.pre_save_hook [Any] Default: None

Python callable or importstring thereof

To be called on a contents model prior to save.

This can be used to process the structure, such as removing notebook outputs or other side effects that should not be saved.

It will be called as (all arguments passed by keyword):

```
hook(path=path, model=model, contents_manager=self)
```

- model: the model to be saved. Includes file contents. Modifying this dict will affect the file that is stored.
- path: the API path of the save destination
- contents_manager: this ContentsManager instance

ContentsManager.root_dir [Unicode] Default: '/'

No description

ContentsManager.untitled_directory [Unicode] Default: 'Untitled Folder'

The base name used when creating untitled directories.

ContentsManager.untitled_file [Unicode] Default: 'untitled'

The base name used when creating untitled files.

ContentsManager.untitled_notebook [Unicode] Default: 'Untitled'

The base name used when creating untitled notebooks.

FileManagerMixin.use_atomic_writing [Bool] Default: True

By default notebooks are saved on disk on a temporary file and then if successfully written, it replaces the old ones. This procedure, namely 'atomic_writing', causes some bugs on file system without operation order enforcement (like some networked fs). If set to False, the new notebook is written directly on the old one which could fail (eg: full filesystem or quota)

FileContentsManager.allow_hidden [Bool] Default: False

Allow access to hidden files

FileContentsManager.checkpoints [Instance] Default: None

No description

FileContentsManager.checkpoints_class [Type] Default: 'notebook.services.contents.checkpoints.Checkpoints'

No description

FileContentsManager.checkpoints_kwargs [Dict] Default: {}

No description

FileContentsManager.delete_to_trash [Bool] Default: True

If True (default), deleting files will send them to the platform's trash/recycle bin, where they can be recovered. If False, deleting files really deletes them.

FileContentsManager.files_handler_class [Type] Default: `'notebook.files.handlers.FilesHandler'`

handler class to use when serving raw file requests.

Default is a fallback that talks to the ContentsManager API, which may be inefficient, especially for large files.

Local files-based ContentsManagers can use a StaticFileHandler subclass, which will be much more efficient.

Access to these files should be Authenticated.

FileContentsManager.files_handler_params [Dict] Default: {}

Extra parameters to pass to files_handler_class.

For example, StaticFileHandlers generally expect a *path* argument specifying the root directory from which to serve files.

FileContentsManager.hide_globs [List] Default: `['__pycache__', '*.pyc', '*.pyo', 'DS_Store', '*.so', '*.dyl...]`

Glob patterns to hide in file and directory listings.

FileContentsManager.post_save_hook [Any] Default: None

Python callable or importstring thereof

to be called on the path of a file just saved.

This can be used to process the file on disk, such as converting the notebook to a script or HTML via nbconvert.

It will be called as (all arguments passed by keyword):

```
hook(os_path=os_path, model=model, contents_manager=instance)
```

- path: the filesystem path to the file just written
- model: the model representing the file
- contents_manager: this ContentsManager instance

FileContentsManager.pre_save_hook [Any] Default: None

Python callable or importstring thereof

To be called on a contents model prior to save.

This can be used to process the structure, such as removing notebook outputs or other side effects that should not be saved.

It will be called as (all arguments passed by keyword):

```
hook(path=path, model=model, contents_manager=self)
```

- model: the model to be saved. Includes file contents. Modifying this dict will affect the file that is stored.
- path: the API path of the save destination
- contents_manager: this ContentsManager instance

FileContentsManager.root_dir [Unicode] Default: ''

No description

FileContentsManager.save_script [Bool] Default: False

DEPRECATED, use post_save_hook. Will be removed in Notebook 5.0

FileContentsManager.untitled_directory [Unicode] Default: 'Untitled Folder'

The base name used when creating untitled directories.

FileContentsManager.untitled_file [Unicode] Default: 'untitled'

The base name used when creating untitled files.

FileContentsManager.untitled_notebook [Unicode] Default: 'Untitled'

The base name used when creating untitled notebooks.

FileContentsManager.use_atomic_writing [Bool] Default: True

By default notebooks are saved on disk on a temporary file and then if successfully written, it replaces the old ones. This procedure, namely 'atomic_writing', causes some bugs on file system without operation order enforcement (like some networked fs). If set to False, the new notebook is written directly on the old one which could fail (eg: full filesystem or quota)

NotebookNotary.algorithm [any of 'sha3_384' | 'sha384' | 'blake2s' | 'md5' | 'sha224' | 'sha3_256' | 'sha256' | 'blake2b' | 'sha512'] Default: 'sha256'

The hashing algorithm used to sign notebooks.

NotebookNotary.data_dir [Unicode] Default: ''

The storage directory for notary secret and database.

NotebookNotary.db_file [Unicode] Default: ''

The sqlite file in which to store notebook signatures. By default, this will be in your Jupyter data directory. You can set it to 'memory:' to disable sqlite writing to the filesystem.

NotebookNotary.secret [Bytes] Default: b''

The secret key with which notebooks are signed.

NotebookNotary.secret_file [Unicode] Default: ''

The file where the secret key is stored.

NotebookNotary.store_factory [Callable] Default: traitlets.Undefined

A callable returning the storage backend for notebook signatures. The default uses an SQLite database.

AsyncMultiKernelManager.default_kernel_name [Unicode] Default: 'python3'

The name of the default kernel to start

AsyncMultiKernelManager.kernel_manager_class [DottedObjectName] Default: 'jupyter_client.ioloop.AsyncIOLoopKernelManager'

The kernel manager class. This is configurable to allow subclassing of the AsyncKernelManager for customized behavior.

AsyncMultiKernelManager.shared_context [Bool] Default: True

Share a single zmq.Context to talk to all my kernels

AsyncMappingKernelManager.allowed_message_types [List] Default: []

White list of allowed kernel message types. When the list is empty, all message types are allowed.

AsyncMappingKernelManager.buffer_offline_messages [Bool] Default: True

Whether messages from kernels whose frontends have disconnected should be buffered in-memory. When True (default), messages are buffered and replayed on reconnect, avoiding lost messages due to interrupted connectivity. Disable if long-running kernels will produce too much output while no frontends are connected.

AsyncMappingKernelManager.cull_busy [Bool] Default: False

Whether to consider culling kernels which are busy. Only effective if cull_idle_timeout > 0.

AsyncMappingKernelManager.cull_connected [Bool] Default: False

Whether to consider culling kernels which have one or more connections. Only effective if cull_idle_timeout > 0.

AsyncMappingKernelManager.cull_idle_timeout [Int] Default: 0

Timeout (in seconds) after which a kernel is considered idle and ready to be culled. Values of 0 or lower disable culling. Very short timeouts may result in kernels being culled for users with poor network connections.

AsyncMappingKernelManager.cull_interval [Int] Default: 300

The interval (in seconds) on which to check for idle kernels exceeding the cull timeout value.

AsyncMappingKernelManager.default_kernel_name [Unicode] Default: 'python3'

The name of the default kernel to start

AsyncMappingKernelManager.kernel_info_timeout [Float] Default: 60

Timeout for giving up on a kernel (in seconds). On starting and restarting kernels, we check whether the kernel is running and responsive by sending kernel_info_requests. This sets the timeout in seconds for how long the kernel can take before being presumed dead. This affects the MappingKernelManager (which handles kernel restarts) and the ZMQChannelsHandler (which handles the startup).

AsyncMappingKernelManager.kernel_manager_class [DottedObjectName] Default: 'jupyter_client.ioloop.AsyncIOLoopKernelManager'

The kernel manager class. This is configurable to allow subclassing of the AsyncKernelManager for customized behavior.

AsyncMappingKernelManager.root_dir [Unicode] Default: ''

No description

AsyncMappingKernelManager.shared_context [Bool] Default: True

Share a single zmq.Context to talk to all my kernels

GatewayKernelManager.allowed_message_types [List] Default: []

White list of allowed kernel message types. When the list is empty, all message types are allowed.

GatewayKernelManager.buffer_offline_messages [Bool] Default: True

Whether messages from kernels whose frontends have disconnected should be buffered in-memory. When True (default), messages are buffered and replayed on reconnect, avoiding lost messages due to interrupted connectivity. Disable if long-running kernels will produce too much output while no frontends are connected.

GatewayKernelManager.cull_busy [Bool] Default: False

Whether to consider culling kernels which are busy. Only effective if cull_idle_timeout > 0.

GatewayKernelManager.cull_connected [Bool] Default: `False`

Whether to consider culling kernels which have one or more connections. Only effective if `cull_idle_timeout > 0`.

GatewayKernelManager.cull_idle_timeout [Int] Default: `0`

Timeout (in seconds) after which a kernel is considered idle and ready to be culled. Values of 0 or lower disable culling. Very short timeouts may result in kernels being culled for users with poor network connections.

GatewayKernelManager.cull_interval [Int] Default: `300`

The interval (in seconds) on which to check for idle kernels exceeding the cull timeout value.

GatewayKernelManager.default_kernel_name [Unicode] Default: `'python3'`

The name of the default kernel to start

GatewayKernelManager.kernel_info_timeout [Float] Default: `60`

Timeout for giving up on a kernel (in seconds). On starting and restarting kernels, we check whether the kernel is running and responsive by sending `kernel_info_requests`. This sets the timeout in seconds for how long the kernel can take before being presumed dead. This affects the `MappingKernelManager` (which handles kernel restarts) and the `ZMQChannelsHandler` (which handles the startup).

GatewayKernelManager.kernel_manager_class [DottedObjectName] Default: `'jupyter_client.ioloop.AsyncIOLoopKernelManager'`

The kernel manager class. This is configurable to allow subclassing of the `AsyncKernelManager` for customized behavior.

GatewayKernelManager.root_dir [Unicode] Default: `''`

No description

GatewayKernelManager.shared_context [Bool] Default: `True`

Share a single `zmq.Context` to talk to all my kernels

GatewayKernelSpecManager.ensure_native_kernel [Bool] Default: `True`

If there is no Python kernelspec registered and the IPython kernel is available, ensure it is added to the spec list.

GatewayKernelSpecManager.kernel_spec_class [Type] Default: `'jupyter_client.kernelspec.KernelSpec'`

The kernel spec class. This is configurable to allow subclassing of the `KernelSpecManager` for customized behavior.

GatewayKernelSpecManager.whitelist [Set] Default: `set()`

Whitelist of allowed kernel names.

By default, all installed kernels are allowed.

GatewayClient.auth_token [Unicode] Default: `None`

The authorization token used in the HTTP headers. (`JUPYTER_GATEWAY_AUTH_TOKEN` env var)

GatewayClient.ca_certs [Unicode] Default: `None`

The filename of CA certificates or `None` to use defaults. (`JUPYTER_GATEWAY_CA_CERTS` env var)

GatewayClient.client_cert [Unicode] Default: `None`

The filename for client SSL certificate, if any. (`JUPYTER_GATEWAY_CLIENT_CERT` env var)

GatewayClient.client_key [Unicode] Default: None

The filename for client SSL key, if any. (JUPYTER_GATEWAY_CLIENT_KEY env var)

GatewayClient.connect_timeout [Float] Default: 40.0

The time allowed for HTTP connection establishment with the Gateway server. (JUPYTER_GATEWAY_CONNECT_TIMEOUT env var)

GatewayClient.env_whitelist [Unicode] Default: ''

A comma-separated list of environment variable names that will be included, along with their values, in the kernel startup request. The corresponding *env_whitelist* configuration value must also be set on the Gateway server - since that configuration value indicates which environmental values to make available to the kernel. (JUPYTER_GATEWAY_ENV_WHITELIST env var)

GatewayClient.gateway_retry_interval [Float] Default: 1.0

The time allowed for HTTP reconnection with the Gateway server for the first time. Next will be JUPYTER_GATEWAY_RETRY_INTERVAL multiplied by two in factor of numbers of retries but less than JUPYTER_GATEWAY_RETRY_INTERVAL_MAX. (JUPYTER_GATEWAY_RETRY_INTERVAL env var)

GatewayClient.gateway_retry_interval_max [Float] Default: 30.0

The maximum time allowed for HTTP reconnection retry with the Gateway server. (JUPYTER_GATEWAY_RETRY_INTERVAL_MAX env var)

GatewayClient.gateway_retry_max [Int] Default: 5

The maximum retries allowed for HTTP reconnection with the Gateway server. (JUPYTER_GATEWAY_RETRY_MAX env var)

GatewayClient.headers [Unicode] Default: '{}'

Additional HTTP headers to pass on the request. This value will be converted to a dict. (JUPYTER_GATEWAY_HEADERS env var)

GatewayClient.http_pwd [Unicode] Default: None

The password for HTTP authentication. (JUPYTER_GATEWAY_HTTP_PWD env var)

GatewayClient.http_user [Unicode] Default: None

The username for HTTP authentication. (JUPYTER_GATEWAY_HTTP_USER env var)

GatewayClient.kernels_endpoint [Unicode] Default: '/api/kernels'

The gateway API endpoint for accessing kernel resources (JUPYTER_GATEWAY_KERNELS_ENDPOINT env var)

GatewayClient.kernelspecs_endpoint [Unicode] Default: '/api/kernelspecs'

The gateway API endpoint for accessing kernelspecs (JUPYTER_GATEWAY_KERNELSPECS_ENDPOINT env var)

GatewayClient.kernelspecs_resource_endpoint [Unicode] Default: '/kernelspecs'

The gateway endpoint for accessing kernelspecs resources (JUPYTER_GATEWAY_KERNELSPECS_RESOURCE_ENDPOINT env var)

GatewayClient.request_timeout [Float] Default: 40.0

The time allowed for HTTP request completion. (JUPYTER_GATEWAY_REQUEST_TIMEOUT env var)

GatewayClient.url [Unicode] Default: None

The url of the Kernel or Enterprise Gateway server where kernel specifications are defined and kernel management takes place. If defined, this Notebook server acts as a proxy for all kernel management and kernel specification retrieval. (JUPYTER_GATEWAY_URL env var)

GatewayClient.validate_cert [Bool] Default: `True`

For HTTPS requests, determines if server's certificate should be validated or not. (JUPYTER_GATEWAY_VALIDATE_CERT env var)

GatewayClient.ws_url [Unicode] Default: `None`

The websocket url of the Kernel or Enterprise Gateway server. If not provided, this value will correspond to the value of the Gateway url with 'ws' in place of 'http'. (JUPYTER_GATEWAY_WS_URL env var)

TerminalManager.cull_inactive_timeout [Int] Default: `0`

Timeout (in seconds) in which a terminal has been inactive and ready to be culled. Values of 0 or lower disable culling.

TerminalManager.cull_interval [Int] Default: `300`

The interval (in seconds) on which to check for terminals exceeding the inactive timeout value.

RUNNING A NOTEBOOK SERVER

The *Jupyter notebook* web application is based on a server-client structure. The notebook server uses a **two-process kernel architecture** based on **ZeroMQ**, as well as **Tornado** for serving HTTP requests.

Note: By default, a notebook server runs locally at 127.0.0.1:8888 and is accessible only from *localhost*. You may access the notebook server from the browser using *http://127.0.0.1:8888*.

This document describes how you can *secure a notebook server* and how to *run it on a public interface*.

Important: **This is not the multi-user server you are looking for.** This document describes how you can run a public server with a single user. This should only be done by someone who wants remote access to their personal machine. Even so, doing this requires a thorough understanding of the set-ups limitations and security implications. If you allow multiple users to access a notebook server as it is described in this document, their commands may collide, clobber and overwrite each other.

If you want a multi-user server, the official solution is **JupyterHub**. To use JupyterHub, you need a Unix server (typically Linux) running somewhere that is accessible to your users on a network. This may run over the public internet, but doing so introduces additional **security concerns**.

9.1 Securing a notebook server

You can protect your notebook server with a simple single password. As of notebook 5.0 this can be done automatically. To set up a password manually you can configure the `NotebookApp.password` setting in `jupyter_notebook_config.py`.

9.1.1 Prerequisite: A notebook configuration file

Check to see if you have a notebook configuration file, `jupyter_notebook_config.py`. The default location for this file is your Jupyter folder located in your home directory:

- Windows: `C:\Users\USERNAME\.jupyter\jupyter_notebook_config.py`
- OS X: `/Users/USERNAME/.jupyter/jupyter_notebook_config.py`
- Linux: `/home/USERNAME/.jupyter/jupyter_notebook_config.py`

If you don't already have a Jupyter folder, or if your Jupyter folder doesn't contain a notebook configuration file, run the following command:

```
$ jupyter notebook --generate-config
```

This command will create the Jupyter folder if necessary, and create notebook configuration file, `jupyter_notebook_config.py`, in this folder.

9.1.2 Automatic Password setup

As of notebook 5.3, the first time you log-in using a token, the notebook server should give you the opportunity to setup a password from the user interface.

You will be presented with a form asking for the current `_token_`, as well as your `_new_` `_password_`; enter both and click on Login and setup new password.

Next time you need to log in you'll be able to use the new password instead of the login token, otherwise follow the procedure to set a password from the command line.

The ability to change the password at first login time may be disabled by integrations by setting the `--NotebookApp.allow_password_change=False`

Starting at notebook version 5.0, you can enter and store a password for your notebook server with a single command. **jupyter notebook password** will prompt you for your password and record the hashed password in your `jupyter_notebook_config.json`.

```
$ jupyter notebook password
Enter password: ****
Verify password: ****
[NotebookPasswordApp] Wrote hashed password to /Users/you/.jupyter/jupyter_notebook_
↪config.json
```

This can be used to reset a lost password; or if you believe your credentials have been leaked and desire to change your password. Changing your password will invalidate all logged-in sessions after a server restart.

9.1.3 Preparing a hashed password

You can prepare a hashed password manually, using the function `notebook.auth.security.passwd()`:

```
In [1]: from notebook.auth import passwd
In [2]: passwd()
Enter password:
Verify password:
Out[2]: 'sha1:67c9e60bb8b6:9ffede0825894254b2e042ea597d771089e11aed'
```

Caution: `passwd()` when called with no arguments will prompt you to enter and verify your password such as in the above code snippet. Although the function can also be passed a string as an argument such as `passwd('mypassword')`, please **do not** pass a string as an argument inside an IPython session, as it will be saved in your input history.

9.1.4 Adding hashed password to your notebook configuration file

You can then add the hashed password to your `jupyter_notebook_config.py`. The default location for this file `jupyter_notebook_config.py` is in your Jupyter folder in your home directory, `~/ .jupyter`, e.g.:

```
c.NotebookApp.password = u'sha1:67c9e60bb8b6:9ffede0825894254b2e042ea597d771089e11aed'
```

Automatic password setup will store the hash in `jupyter_notebook_config.json` while this method stores the hash in `jupyter_notebook_config.py`. The `.json` configuration options take precedence over the `.py` one, thus the manual password may not take effect if the `Json` file has a password set.

9.1.5 Using SSL for encrypted communication

When using a password, it is a good idea to also use SSL with a web certificate, so that your hashed password is not sent unencrypted by your browser.

Important: Web security is rapidly changing and evolving. We provide this document as a convenience to the user, and recommend that the user keep current on changes that may impact security, such as new releases of OpenSSL. The Open Web Application Security Project ([OWASP](#)) website is a good resource on general security issues and web practices.

You can start the notebook to communicate via a secure protocol mode by setting the `certfile` option to your self-signed certificate, i.e. `mycert.pem`, with the command:

```
$ jupyter notebook --certfile=mycert.pem --keyfile mykey.key
```

Tip: A self-signed certificate can be generated with `openssl`. For example, the following command will create a certificate valid for 365 days with both the key and certificate data written to the same file:

```
$ openssl req -x509 -nodes -days 365 -newkey rsa:2048 -keyout mykey.key -out mycert.  
→pem
```

When starting the notebook server, your browser may warn that your self-signed certificate is insecure or unrecognized. If you wish to have a fully compliant self-signed certificate that will not raise warnings, it is possible (but rather involved) to create one, as explained in detail in this [tutorial](#). Alternatively, you may use [Let's Encrypt](#) to acquire a free SSL certificate and follow the steps in [Using Let's Encrypt](#) to set up a public server.

9.2 Running a public notebook server

If you want to access your notebook server remotely via a web browser, you can do so by running a public notebook server. For optimal security when running a public notebook server, you should first secure the server with a password and SSL/HTTPS as described in [Securing a notebook server](#).

Start by creating a certificate file and a hashed password, as explained in [Securing a notebook server](#).

If you don't already have one, create a config file for the notebook using the following command line:

```
$ jupyter notebook --generate-config
```

In the `~/.` `jupyter` directory, edit the notebook config file, `jupyter_notebook_config.py`. By default, the notebook config file has all fields commented out. The minimum set of configuration options that you should uncomment and edit in `jupyter_notebook_config.py` is the following:

```
# Set options for certfile, ip, password, and toggle off
# browser auto-opening
c.NotebookApp.certfile = u'/absolute/path/to/your/certificate/mycert.pem'
c.NotebookApp.keyfile = u'/absolute/path/to/your/certificate/mykey.key'
# Set ip to '*' to bind on all interfaces (ips) for the public server
c.NotebookApp.ip = '*'
c.NotebookApp.password = u'sha1:bcd259ccf...<your hashed password here>'
c.NotebookApp.open_browser = False

# It is a good idea to set a known, fixed port for server access
c.NotebookApp.port = 9999
```

You can then start the notebook using the `jupyter notebook` command.

9.2.1 Using Let's Encrypt

Let's Encrypt provides free SSL/TLS certificates. You can also set up a public server using a Let's Encrypt certificate.

Running a public notebook server will be similar when using a Let's Encrypt certificate with a few configuration changes. Here are the steps:

1. Create a Let's Encrypt certificate.
2. Use *Preparing a hashed password* to create one.
3. If you don't already have config file for the notebook, create one using the following command:

```
$ jupyter notebook --generate-config
```

4. In the `~/.` `jupyter` directory, edit the notebook config file, `jupyter_notebook_config.py`. By default, the notebook config file has all fields commented out. The minimum set of configuration options that you should uncomment and edit in `jupyter_notebook_config.py` is the following:

```
# Set options for certfile, ip, password, and toggle off
# browser auto-opening
c.NotebookApp.certfile = u'/absolute/path/to/your/certificate/fullchain.pem'
c.NotebookApp.keyfile = u'/absolute/path/to/your/certificate/privkey.pem'
# Set ip to '*' to bind on all interfaces (ips) for the public server
c.NotebookApp.ip = '*'
c.NotebookApp.password = u'sha1:bcd259ccf...<your hashed password here>'
c.NotebookApp.open_browser = False

# It is a good idea to set a known, fixed port for server access
c.NotebookApp.port = 9999
```

You can then start the notebook using the `jupyter notebook` command.

Important: Use **'https'**. Keep in mind that when you enable SSL support, you must access the notebook server over `https://`, not over plain `http://`. The startup message from the server prints a reminder in the console, but *it is easy to overlook this detail and think the server is for some reason non-responsive*.

When using SSL, always access the notebook server with **'https://'.**

You may now access the public server by pointing your browser to `https://your.host.com:9999` where `your.host.com` is your public server's domain.

9.2.2 Firewall Setup

To function correctly, the firewall on the computer running the jupyter notebook server must be configured to allow connections from client machines on the access port `c.NotebookApp.port` set in `jupyter_notebook_config.py` to allow connections to the web interface. The firewall must also allow connections from 127.0.0.1 (localhost) on ports from 49152 to 65535. These ports are used by the server to communicate with the notebook kernels. The kernel communication ports are chosen randomly by ZeroMQ, and may require multiple connections per kernel, so a large range of ports must be accessible.

9.3 Running the notebook with a customized URL prefix

The notebook dashboard, which is the landing page with an overview of the notebooks in your working directory, is typically found and accessed at the default URL `http://localhost:8888/`.

If you prefer to customize the URL prefix for the notebook dashboard, you can do so through modifying `jupyter_notebook_config.py`. For example, if you prefer that the notebook dashboard be located with a sub-directory that contains other ipython files, e.g. `http://localhost:8888/ipython/`, you can do so with configuration options like the following (see above for instructions about modifying `jupyter_notebook_config.py`):

```
c.NotebookApp.base_url = '/ipython/'
```

9.4 Embedding the notebook in another website

Sometimes you may want to embed the notebook somewhere on your website, e.g. in an `IFrame`. To do this, you may need to override the Content-Security-Policy to allow embedding. Assuming your website is at `https://mywebsite.example.com`, you can embed the notebook on your website with the following configuration setting in `jupyter_notebook_config.py`:

```
c.NotebookApp.tornado_settings = {
    'headers': {
        'Content-Security-Policy': "frame-ancestors https://mywebsite.example.com
↪ 'self' "
    }
}
```

When embedding the notebook in a website using an `iframe`, consider putting the notebook in single-tab mode. Since the notebook opens some links in new tabs by default, single-tab mode keeps the notebook from opening additional tabs. Adding the following to `~/ .jupyter/custom/custom.js` will enable single-tab mode:

```
define(['base/js/namespace'], function(Jupyter) {
    Jupyter._target = '_self';
});
```

9.5 Using a gateway server for kernel management

You are now able to redirect the management of your kernels to a Gateway Server (i.e., [Jupyter Kernel Gateway](#) or [Jupyter Enterprise Gateway](#)) simply by specifying a Gateway url via the following command-line option:

```
$ jupyter notebook --gateway-url=http://my-gateway-server:8888
```

the environment:

```
JUPYTER_GATEWAY_URL=http://my-gateway-server:8888
```

or in `jupyter_notebook_config.py`:

```
c.GatewayClient.url = http://my-gateway-server:8888
```

When provided, all kernel specifications will be retrieved from the specified Gateway server and all kernels will be managed by that server. This option enables the ability to target kernel processes against managed clusters while allowing for the notebook’s management to remain local to the Notebook server.

9.6 Known issues

9.6.1 Proxies

When behind a proxy, especially if your system or browser is set to autodetect the proxy, the notebook web application might fail to connect to the server’s websockets, and present you with a warning at startup. In this case, you need to configure your system not to use the proxy for the server’s address.

For example, in Firefox, go to the Preferences panel, Advanced section, Network tab, click ‘Settings...’, and add the address of the notebook server to the ‘No proxy for’ field.

9.6.2 Content-Security-Policy (CSP)

Certain [security guidelines](#) recommend that servers use a Content-Security-Policy (CSP) header to prevent cross-site scripting vulnerabilities, specifically limiting to `default-src: https:` when possible. This directive causes two problems with Jupyter. First, it disables execution of inline javascript code, which is used extensively by Jupyter. Second, it limits communication to the https scheme, and prevents WebSockets from working because they communicate via the wss scheme (or ws for insecure communication). Jupyter uses WebSockets for interacting with kernels, so when you visit a server with such a CSP, your browser will block attempts to use wss, which will cause you to see “Connection failed” messages from jupyter notebooks, or simply no response from jupyter terminals. By looking in your browser’s javascript console, you can see any error messages that will explain what is failing.

To avoid these problem, you need to add `'unsafe-inline'` and `connect-src https: wss:` to your CSP header, at least for pages served by jupyter. (That is, you can leave your CSP unchanged for other parts of your website.) Note that multiple CSP headers are allowed, but successive CSP headers can only restrict the policy; they cannot loosen it. For example, if your server sends both of these headers

```
Content-Security-Policy “default-src https: ‘unsafe-inline’” Content-Security-Policy “connect-src https: wss:”
```

the first policy will already eliminate wss connections, so the second has no effect. Therefore, you can’t simply add the second header; you have to actually modify your CSP header to look more like this:

```
Content-Security-Policy “default-src https: ‘unsafe-inline’; connect-src https: wss:”
```

9.6.3 Docker CMD

Using `jupyter notebook` as a **Docker CMD** results in kernels repeatedly crashing, likely due to a lack of **PID reaping**. To avoid this, use the `tini` init as your Dockerfile *ENTRYPOINT*:

```
# Add Tini. Tini operates as a process subreaper for jupyter. This prevents
# kernel crashes.
ENV TINI_VERSION v0.6.0
ADD https://github.com/krallin/tini/releases/download/${TINI_VERSION}/tini /usr/bin/
↳ tini
RUN chmod +x /usr/bin/tini
ENTRYPOINT ["/usr/bin/tini", "--"]

EXPOSE 8888
CMD ["jupyter", "notebook", "--port=8888", "--no-browser", "--ip=0.0.0.0"]
```


SECURITY IN THE JUPYTER NOTEBOOK SERVER

Since access to the Jupyter notebook server means access to running arbitrary code, it is important to restrict access to the notebook server. For this reason, notebook 4.3 introduces token-based authentication that is **on by default**.

Note: If you enable a password for your notebook server, token authentication is not enabled by default, and the behavior of the notebook server is unchanged from versions earlier than 4.3.

When token authentication is enabled, the notebook uses a token to authenticate requests. This token can be provided to login to the notebook server in three ways:

- in the `Authorization` header, e.g.:

```
Authorization: token abcdef...
```

- In a URL parameter, e.g.:

```
https://my-notebook/tree/?token=abcdef...
```

- In the password field of the login form that will be shown to you if you are not logged in.

When you start a notebook server with token authentication enabled (default), a token is generated to use for authentication. This token is logged to the terminal, so that you can copy/paste the URL into your browser:

```
[I 11:59:16.597 NotebookApp] The Jupyter Notebook is running at:
http://localhost:8888/?token=c8de56fa4deed24899803e93c227592aef6538f93025fe01
```

If the notebook server is going to open your browser automatically (the default, unless `--no-browser` has been passed), an *additional* token is generated for launching the browser. This additional token can be used only once, and is used to set a cookie for your browser once it connects. After your browser has made its first request with this one-time-token, the token is discarded and a cookie is set in your browser.

At any later time, you can see the tokens and URLs for all of your running servers with **jupyter notebook list**:

```
$ jupyter notebook list
Currently running servers:
http://localhost:8888/?token=abc... :: /home/you/notebooks
https://0.0.0.0:9999/?token=123... :: /tmp/public
http://localhost:8889/ :: /tmp/has-password
```

For servers with token-authentication enabled, the URL in the above listing will include the token, so you can copy and paste that URL into your browser to login. If a server has no token (e.g. it has a password or has authentication disabled), the URL will not include the token argument. Once you have visited this URL, a cookie will be set in your

browser and you won't need to use the token again, unless you switch browsers, clear your cookies, or start a notebook server on a new port.

10.1 Alternatives to token authentication

If a generated token doesn't work well for you, you can set a password for your notebook. **jupyter notebook password** will prompt you for a password, and store the hashed password in your `jupyter_notebook_config.json`.

New in version 5.0: **jupyter notebook password** command is added.

It is possible to disable authentication altogether by setting the token and password to empty strings, but this is **NOT RECOMMENDED**, unless authentication or access restrictions are handled at a different layer in your web application:

```
c.NotebookApp.token = ''
c.NotebookApp.password = ''
```

SECURITY IN NOTEBOOK DOCUMENTS

As Jupyter notebooks become more popular for sharing and collaboration, the potential for malicious people to attempt to exploit the notebook for their nefarious purposes increases. IPython 2.0 introduced a security model to prevent execution of untrusted code without explicit user input.

11.1 The problem

The whole point of Jupyter is arbitrary code execution. We have no desire to limit what can be done with a notebook, which would negatively impact its utility.

Unlike other programs, a Jupyter notebook document includes output. Unlike other documents, that output exists in a context that can execute code (via Javascript).

The security problem we need to solve is that no code should execute just because a user has **opened** a notebook that **they did not write**. Like any other program, once a user decides to execute code in a notebook, it is considered trusted, and should be allowed to do anything.

11.2 Our security model

- Untrusted HTML is always sanitized
- Untrusted Javascript is never executed
- HTML and Javascript in Markdown cells are never trusted
- **Outputs** generated by the user are trusted
- Any other HTML or Javascript (in Markdown cells, output generated by others) is never trusted
- The central question of trust is “Did the current user do this?”

11.3 The details of trust

When a notebook is executed and saved, a signature is computed from a digest of the notebook’s contents plus a secret key. This is stored in a database, writable only by the current user. By default, this is located at:

```
~/local/share/jupyter/nbsignatures.db # Linux
~/Library/Jupyter/nbsignatures.db    # OS X
%APPDATA%/jupyter/nbsignatures.db    # Windows
```

Each signature represents a series of outputs which were produced by code the current user executed, and are therefore trusted.

When you open a notebook, the server computes its signature, and checks if it's in the database. If a match is found, HTML and Javascript output in the notebook will be trusted at load, otherwise it will be untrusted.

Any output generated during an interactive session is trusted.

11.3.1 Updating trust

A notebook's trust is updated when the notebook is saved. If there are any untrusted outputs still in the notebook, the notebook will not be trusted, and no signature will be stored. If all untrusted outputs have been removed (either via `Clear Output` or re-execution), then the notebook will become trusted.

While trust is updated per output, this is only for the duration of a single session. A newly loaded notebook file is either trusted or not in its entirety.

11.3.2 Explicit trust

Sometimes re-executing a notebook to generate trusted output is not an option, either because dependencies are unavailable, or it would take a long time. Users can explicitly trust a notebook in two ways:

- At the command-line, with:

```
jupyter trust /path/to/notebook.ipynb
```

- After loading the untrusted notebook, with `File / Trust Notebook`

These two methods simply load the notebook, compute a new signature, and add that signature to the user's database.

11.4 Reporting security issues

If you find a security vulnerability in Jupyter, either a failure of the code to properly implement the model described here, or a failure of the model itself, please report it to security@ipython.org.

If you prefer to encrypt your security reports, you can use [this PGP public key](#).

11.5 Affected use cases

Some use cases that work in Jupyter 1.0 became less convenient in 2.0 as a result of the security changes. We do our best to minimize these annoyances, but security is always at odds with convenience.

11.5.1 Javascript and CSS in Markdown cells

While never officially supported, it had become common practice to put hidden Javascript or CSS styling in Markdown cells, so that they would not be visible on the page. Since Markdown cells are now sanitized (by [Google Caja](#)), all Javascript (including click event handlers, etc.) and CSS will be stripped.

We plan to provide a mechanism for notebook themes, but in the meantime styling the notebook can only be done via either `custom.css` or CSS in HTML output. The latter only have an effect if the notebook is trusted, because otherwise the output will be sanitized just like Markdown.

11.5.2 Collaboration

When collaborating on a notebook, people probably want to see the outputs produced by their colleagues' most recent executions. Since each collaborator's key will differ, this will result in each share starting in an untrusted state. There are three basic approaches to this:

- re-run notebooks when you get them (not always viable)
- explicitly trust notebooks via `jupyter trust` or the notebook menu (annoying, but easy)
- share a notebook signatures database, and use configuration dedicated to the collaboration while working on the project.

To share a signatures database among users, you can configure:

```
c.NotebookNotary.data_dir = "/path/to/signature_dir"
```

to specify a non-default path to the SQLite database (of notebook hashes, essentially). We are aware that SQLite doesn't work well on NFS and we are [working out better ways to do this](#).

DISTRIBUTING JUPYTER EXTENSIONS AS PYTHON PACKAGES

12.1 Overview

12.1.1 How can the notebook be extended?

The Jupyter Notebook client and server application are both deeply customizable. Their behavior can be extended by creating, respectively:

- nbextension: a notebook extension
 - a single JS file, or directory of JavaScript, Cascading StyleSheets, etc. that contain at minimum a JavaScript module packaged as an `AMD modules` that exports a function `load_ipython_extension`
- server extension: an importable Python module
 - that implements `load_jupyter_server_extension`
- bundler extension: an importable Python module with generated File -> Download as / Deploy as menu item trigger
 - that implements `bundle`

12.1.2 Why create a Python package for Jupyter extensions?

Since it is rare to have a server extension that does not have any frontend components (an nbextension), for convenience and consistency, all these client and server extensions with their assets can be packaged and versioned together as a Python package with a few simple commands, or as of Notebook 5.3, handled automatically by your package manager of choice. This makes installing the package of extensions easier and less error-prone for the user.

12.2 Installation of Jupyter Extensions

12.2.1 Install a Python package containing Jupyter Extensions

There are several ways that you may get a Python package containing Jupyter Extensions. Commonly, you will use a package manager for your system:

```
pip install helpful_package
# or
conda install helpful_package
# or
apt-get install helpful_package
```

(continues on next page)

(continued from previous page)

```
# where 'helpful_package' is a Python package containing one or more Jupyter_
↳Extensions
```

12.2.2 Automatic installation and Enabling

New in Notebook 5.3

The absolute simplest case requires no user interaction at all! Configured correctly, after installing with their package manager of choice, both server and frontend extensions can be enabled by default in the environment where they were installed, i.e. `--sys-prefix`. See the `setup.py` in the example below.

12.2.3 Enable a Server Extension

The simplest case would be to enable a server extension which has no frontend components.

A pip user that wants their configuration stored in their home directory would type the following command:

```
jupyter serverextension enable --py helpful_package
```

Alternatively, a virtualenv or conda user can pass `--sys-prefix` which keeps their environment isolated and reproducible. For example:

```
# Make sure that your virtualenv or conda environment is activated
[source] activate my-environment

jupyter serverextension enable --py helpful_package --sys-prefix
```

12.2.4 Install the nbextension assets

If a package also has an nbextension with frontend assets that must be available (but not necessarily enabled by default), install these assets with the following command:

```
jupyter nbextension install --py helpful_package # or --sys-prefix if using_
↳virtualenv or conda
```

12.2.5 Enable nbextension assets

If a package has assets that should be loaded every time a Jupyter app (e.g. lab, notebook, dashboard, terminal) is loaded in the browser, the following command can be used to enable the nbextension:

```
jupyter nbextension enable --py helpful_package # or --sys-prefix if using virtualenv_
↳or conda
```


12.3 Did it work? Check by listing Jupyter Extensions.

After running one or more extension installation steps, you can list what is presently known about nbextensions, server extensions, or bundler extensions. The following commands will list which extensions are available, whether they are enabled, and other extension details:

```
jupyter nbextension list
jupyter serverextension list
jupyter bundlerextension list
```

12.4 Additional resources on creating and distributing packages

Of course, in addition to the files listed, there are number of other files one needs to build a proper package. Here are some good resources: - [The Hitchhiker's Guide to Packaging - Repository Structure and Python](#) by Kenneth Reitz

How you distribute them, too, is important: - [Packaging and Distributing Projects](#) - [conda: Building packages](#)

12.5 Example - Server extension

12.5.1 Creating a Python package with a server extension

Here is an example of a python module which contains a server extension directly on itself. It has this directory structure:

```
- setup.py
- MANIFEST.in
- my_module/
  - __init__.py
```

12.5.2 Defining the server extension

This example shows that the server extension and its `load_jupyter_server_extension` function are defined in the `__init__.py` file.

`my_module/__init__.py`

```
def _jupyter_server_extension_paths():
    return [{
        "module": "my_module"
    }]

def load_jupyter_server_extension(nbapp):
    nbapp.log.info("my module enabled!")
```

12.5.3 Install and enable the server extension

Which a user can install with:

```
jupyter serverextension enable --py my_module [--sys-prefix]
```

12.6 Example - Server extension and nbextension

12.6.1 Creating a Python package with a server extension and nbextension

Here is another server extension, with a front-end module. It assumes this directory structure:

```
- setup.py
- MANIFEST.in
- my_fancy_module/
  - __init__.py
  - static/
    index.js
```

12.6.2 Defining the server extension and nbextension

This example again shows that the server extension and its `load_jupyter_server_extension` function are defined in the `__init__.py` file. This time, there is also a function `_jupyter_nbextension_paths` for the nbextension.

`my_fancy_module/__init__.py`

```
def _jupyter_server_extension_paths():
    return [{
        "module": "my_fancy_module"
    }]

# Jupyter Extension points
def _jupyter_nbextension_paths():
    return [dict(
        section="notebook",
        # the path is relative to the `my_fancy_module` directory
        src="static",
        # directory in the `nbextension/` namespace
        dest="my_fancy_module",
        # _also_ in the `nbextension/` namespace
        require="my_fancy_module/index")]

def load_jupyter_server_extension(nbapp):
    nbapp.log.info("my module enabled!")
```

12.6.3 Install and enable the server extension and nbextension

The user can install and enable the extensions with the following set of commands:

```
jupyter nbextension install --py my_fancy_module [--sys-prefix|--user]
jupyter nbextension enable --py my_fancy_module [--sys-prefix|--system]
jupyter serverextension enable --py my_fancy_module [--sys-prefix|--system]
```

12.6.4 Automatically enabling a server extension and nbextension

New in Notebook 5.3

Server extensions and nbextensions can be installed and enabled without any user intervention or post-install scripts beyond `<package manager> install <extension package name>`

In addition to the `my_fancy_module` file tree, assume:

```
jupyter-config/
├── jupyter_notebook_config.d/
│   └── my_fancy_module.json
└── nbconfig/
    ├── notebook.d/
    │   └── my_fancy_module.json
```

`jupyter-config/jupyter_notebook_config.d/my_fancy_module.json`

```
{
  "NotebookApp": {
    "nbserver_extensions": {
      "my_fancy_module": true
    }
  }
}
```

`jupyter-config/nbconfig/notebook.d/my_fancy_module.json`

```
{
  "load_extensions": {
    "my_fancy_module/index": true
  }
}
```

Put all of them in place via:

setup.py

```
import setuptools

setuptools.setup(
    name="MyFancyModule",
    ...
    include_package_data=True,
    data_files=[
        # like `jupyter nbextension install --sys-prefix`
        ("share/jupyter/nbextensions/my_fancy_module", [
            "my_fancy_module/static/index.js",
        ]),
        # like `jupyter nbextension enable --sys-prefix`
        ("etc/jupyter/nbconfig/notebook.d", [
            "jupyter-config/nbconfig/notebook.d/my_fancy_module.json"
        ]),
        # like `jupyter serverextension enable --sys-prefix`
        ("etc/jupyter/jupyter_notebook_config.d", [
            "jupyter-config/jupyter_notebook_config.d/my_fancy_module.json"
        ])
    ],
    ...
    zip_safe=False
)
```

and last, but not least:

MANIFEST.in

```
recursive-include jupyter-config *.json
recursive-include my_fancy_module/static *.js
```

As most package managers will only modify their environment, the eventual configuration will be as if the user had typed:

```
jupyter nbextension install --py my_fancy_module --sys-prefix
jupyter nbextension enable --py my_fancy_module --sys-prefix
jupyter serverextension enable --py my_fancy_module --sys-prefix
```

If a user manually disables an extension, that configuration will override the bundled package configuration.

When automagical install fails

Note this can still fail in certain situations with `pip`, requiring manual use of `install` and `enable` commands.

Non-python-specific package managers (e.g. `conda`, `apt`) may choose not to implement the above behavior at the `setup.py` level, having more ways to put data files in various places at build time.

12.7 Example - Bundler extension

12.7.1 Creating a Python package with a bundlerextension

Here is a bundler extension that adds a *Download as -> Notebook Tarball (tar.gz)* option to the notebook *File* menu. It assumes this directory structure:

```
- setup.py
- MANIFEST.in
- my_tarball_bundler/
  - __init__.py
```

12.7.2 Defining the bundler extension

This example shows that the bundler extension and its `bundle` function are defined in the `__init__.py` file.

`my_tarball_bundler/__init__.py`

```
import tarfile
import io
import os
import nbformat

def _jupyter_bundlerextension_paths():
    """Declare bundler extensions provided by this package."""
    return [{
        # unique bundler name
        "name": "tarball_bundler",
        # module containing bundle function
        "module_name": "my_tarball_bundler",
        # human-readable menu item label
        "label": "Notebook Tarball (tar.gz)",
        # group under 'deploy' or 'download' menu
        "group": "download",
    }]

def bundle(handler, model):
    """Create a compressed tarball containing the notebook document.

    Parameters
    -----
    handler : tornado.web.RequestHandler
        Handler that serviced the bundle request
    model : dict
        Notebook model from the configured ContentManager
    """
    notebook_filename = model['name']
    notebook_content = nbformat.writes(model['content']).encode('utf-8')
    notebook_name = os.path.splitext(notebook_filename)[0]
    tar_filename = '{}.tar.gz'.format(notebook_name)

    info = tarfile.TarInfo(notebook_filename)
```

(continues on next page)

(continued from previous page)

```
info.size = len(notebook_content)

with io.BytesIO() as tar_buffer:
    with tarfile.open(tar_filename, "w:gz", fileobj=tar_buffer) as tar:
        tar.addfile(info, io.BytesIO(notebook_content))

    # Set headers to trigger browser download
    handler.set_header('Content-Disposition',
                       'attachment; filename="{}".format(tar_filename))
    handler.set_header('Content-Type', 'application/gzip')

    # Return the buffer value as the response
    handler.finish(tar_buffer.getvalue())
```

See [Extending the Notebook](#) for more documentation about writing nbextensions, server extensions, and bundler extensions.

EXTENDING THE NOTEBOOK

Certain subsystems of the notebook server are designed to be extended or overridden by users. These documents explain these systems, and show how to override the notebook's defaults with your own custom behavior.

13.1 Contents API

The Jupyter Notebook web application provides a graphical interface for creating, opening, renaming, and deleting files in a virtual filesystem.

The `ContentsManager` class defines an abstract API for translating these interactions into operations on a particular storage medium. The default implementation, `FileContentsManager`, uses the local filesystem of the server for storage and straightforwardly serializes notebooks into JSON. Users can override these behaviors by supplying custom subclasses of `ContentsManager`.

This section describes the interface implemented by `ContentsManager` subclasses. We refer to this interface as the **Contents API**.

13.1.1 Data Model

Filesystem Entities

`ContentsManager` methods represent virtual filesystem entities as dictionaries, which we refer to as **models**.

Models may contain the following entries:

| Key | Type | Info |
|----------------------|-----------------|--|
| name | unicode | Basename of the entity. |
| path | unicode | Full (<i>API-style</i>) path to the entity. |
| type | unicode | The entity type. One of "notebook", "file" or "directory". |
| created | datetime | Creation date of the entity. |
| last_modified | datetime | Last modified date of the entity. |
| content | variable | The "content" of the entity. (<i>See Below</i>) |
| mimetype | unicode or None | The mimetype of content, if any. (<i>See Below</i>) |
| format | unicode or None | The format of content, if any. (<i>See Below</i>) |

Certain model fields vary in structure depending on the `type` field of the model. There are three model types: **notebook**, **file**, and **directory**.

- **notebook models**
 - The `format` field is always "json".

- The `mimetype` field is always `None`.
- The `content` field contains a `nbformat.notebooknode.NotebookNode` representing the `.ipynb` file represented by the model. See the [NBFormat](#) documentation for a full description.

- **file models**

- The `format` field is either `"text"` or `"base64"`.
- The `mimetype` field can be any `mimetype` string, but defaults to `text/plain` for text-format models and `application/octet-stream` for base64-format models. For files with unknown mime types (e.g. unknown file extensions), this field may be `None`.
- The `content` field is always of type `unicode`. For text-format file models, `content` simply contains the file's bytes after decoding as UTF-8. Non-text (base64) files are read as bytes, base64 encoded, and then decoded as UTF-8.

- **directory models**

- The `format` field is always `"json"`.
- The `mimetype` field is always `None`.
- The `content` field contains a list of *content-free* models representing the entities in the directory.

Note: In certain circumstances, we don't need the full content of an entity to complete a Contents API request. In such cases, we omit the `content`, and `format` keys from the model. The default values for the `mimetype` field will might also not be evaluated, in which case it will be set as `None`. This reduced reply most commonly occurs when listing a directory, in which circumstance we represent files within the directory as content-less models to avoid having to recursively traverse and serialize the entire filesystem.

Sample Models

```
# Notebook Model with Content
{
  'content': {
    'metadata': {},
    'nbformat': 4,
    'nbformat_minor': 0,
    'cells': [
      {
        'cell_type': 'markdown',
        'metadata': {},
        'source': 'Some **Markdown**',
      },
    ],
  },
  'created': datetime(2015, 7, 25, 19, 50, 19, 19865),
  'format': 'json',
  'last_modified': datetime(2015, 7, 25, 19, 50, 19, 19865),
  'mimetype': None,
  'name': 'a.ipynb',
  'path': 'foo/a.ipynb',
  'type': 'notebook',
  'writable': True,
}

# Notebook Model without Content
{
```

(continues on next page)

(continued from previous page)

```

'content': None,
'created': datetime.datetime(2015, 7, 25, 20, 17, 33, 271931),
'format': None,
'last_modified': datetime.datetime(2015, 7, 25, 20, 17, 33, 271931),
'mimetype': None,
'name': 'a.ipynb',
'path': 'foo/a.ipynb',
'type': 'notebook',
'writable': True
}

```

API Paths

ContentsManager methods represent the locations of filesystem resources as **API-style paths**. Such paths are interpreted as relative to the root directory of the notebook server. For compatibility across systems, the following guarantees are made:

- Paths are always unicode, not bytes.
- Paths are not URL-escaped.
- Paths are always forward-slash (/) delimited, even on Windows.
- Leading and trailing slashes are stripped. For example, /foo/bar/buzz/ becomes foo/bar/buzz.
- The empty string ("") represents the root directory.

13.1.2 Writing a Custom ContentsManager

The default ContentsManager is designed for users running the notebook as an application on a personal computer. It stores notebooks as .ipynb files on the local filesystem, and it maps files and directories in the Notebook UI to files and directories on disk. It is possible to override how notebooks are stored by implementing your own custom subclass of ContentsManager. For example, if you deploy the notebook in a context where you don't trust or don't have access to the filesystem of the notebook server, it's possible to write your own ContentsManager that stores notebooks and files in a database.

Required Methods

A minimal complete implementation of a custom ContentsManager must implement the following methods:

| | |
|---|---|
| ContentsManager.get(path[, content, type, ...]) | Get a file or directory model. |
| ContentsManager.save(model, path) | Save a file or directory model to path. |
| ContentsManager.delete_file(path) | Delete the file or directory at path. |
| ContentsManager.rename_file(old_path, new_path) | Rename a file or directory. |
| ContentsManager.file_exists([path]) | Does a file exist at the given path? |
| ContentsManager.dir_exists(path) | Does a directory exist at the given path? |
| ContentsManager.is_hidden(path) | Is path a hidden directory or file? |

You may be required to specify a Checkpoints object, as the default one, FileCheckpoints, could be incompatible with your custom ContentsManager.

Chunked Saving

The contents API allows for “chunked” saving of files, i.e. saving/transmitting in partial pieces:

- This can only be used when the `type` of the model is `file`.
- The model should be as otherwise expected for `save()`, with an added field `chunk`.
- The value of `chunk` should be an integer starting at 1, and incrementing for each subsequent chunk, except for the final chunk, which should be indicated with a value of `-1`.
- The model returned from using `save()` with `chunk` should be treated as unreliable for all chunks except the final one.
- Any interaction with a file being saved in a chunked manner is unreliable until the final chunk has been saved. This includes directory listings.

13.1.3 Customizing Checkpoints

Customized Checkpoint definitions allows behavior to be altered and extended.

The `Checkpoints` and `GenericCheckpointsMixin` classes (from `notebook.services.contents.checkpoints`) have reusable code and are intended to be used together, but require the following methods to be implemented.

| | |
|---|---|
| <code>Checkpoints.rename_checkpoint(checkpoint_id)</code> | Rename a single checkpoint from <code>old_path</code> to <code>new_path</code> .
... |
| <code>Checkpoints.list_checkpoints(path)</code> | Return a list of checkpoints for a given file |
| <code>Checkpoints.delete_checkpoint(checkpoint_id)</code> | delete a checkpoint for a file
... |
| <code>GenericCheckpointsMixin.create_file_checkpoint(...)</code> | Create a checkpoint of the current state of a file |
| <code>GenericCheckpointsMixin.create_notebook_checkpoint(nb,...)</code> | Create a checkpoint of the current state of a file |
| <code>GenericCheckpointsMixin.get_file_checkpoint(...)</code> | Get the content of a checkpoint for a non-notebook file. |
| <code>GenericCheckpointsMixin.get_notebook_checkpoint(...)</code> | Get the content of a checkpoint for a notebook. |

No-op example

Here is an example of a no-op checkpoints object - note the mixin comes first. The docstrings indicate what each method should do or return for a more complete implementation.

```
class NoOpCheckpoints(GenericCheckpointsMixin, Checkpoints):
    """requires the following methods:"""
    def create_file_checkpoint(self, content, format, path):
        """ -> checkpoint model"""
    def create_notebook_checkpoint(self, nb, path):
        """ -> checkpoint model"""
    def get_file_checkpoint(self, checkpoint_id, path):
        """ -> {'type': 'file', 'content': <str>, 'format': {'text', 'base64'}}"""
    def get_notebook_checkpoint(self, checkpoint_id, path):
        """ -> {'type': 'notebook', 'content': <output of nbformat.read>}"""
    def delete_checkpoint(self, checkpoint_id, path):
```

(continues on next page)

(continued from previous page)

```

"""deletes a checkpoint for a file"""
def list_checkpoints(self, path):
    """returns a list of checkpoint models for a given file,
    default just does one per file
    """
    return []
def rename_checkpoint(self, checkpoint_id, old_path, new_path):
    """renames checkpoint from old path to new path"""

```

See `GenericFileCheckpoints` in `notebook.services.contents.filecheckpoints` for a more complete example.

13.1.4 Testing

`notebook.services.contents.tests` includes several test suites written against the abstract `Contents` API. This means that an excellent way to test a new `ContentsManager` subclass is to subclass our tests to make them use your `ContentsManager`.

Note: `PGContents` is an example of a complete implementation of a custom `ContentsManager`. It stores notebooks and files in `PostgreSQL` and encodes directories as SQL relations. `PGContents` also provides an example of how to re-use the notebook's tests.

13.2 File save hooks

You can configure functions that are run whenever a file is saved. There are two hooks available:

- `ContentsManager.pre_save_hook` runs on the API path and model with content. This can be used for things like stripping output that people don't like adding to VCS noise.
- `FileContentsManager.post_save_hook` runs on the filesystem path and model without content. This could be used to commit changes after every save, for instance.

They are both called with keyword arguments:

```

pre_save_hook(model=model, path=path, contents_manager=cm)
post_save_hook(model=model, os_path=os_path, contents_manager=cm)

```

13.2.1 Examples

These can both be added to `jupyter_notebook_config.py`.

A pre-save hook for stripping output:

```

def scrub_output_pre_save(model, **kwargs):
    """scrub output before saving notebooks"""
    # only run on notebooks
    if model['type'] != 'notebook':
        return
    # only run on nbformat v4
    if model['content']['nbformat'] != 4:

```

(continues on next page)

(continued from previous page)

```

        return

    for cell in model['content']['cells']:
        if cell['cell_type'] != 'code':
            continue
        cell['outputs'] = []
        cell['execution_count'] = None

c.FileContentsManager.pre_save_hook = scrub_output_pre_save

```

A post-save hook to make a script equivalent whenever the notebook is saved (replacing the `--script` option in older versions of the notebook):

```

import io
import os
from notebook.utils import to_api_path

_script_exporter = None

def script_post_save(model, os_path, contents_manager, **kwargs):
    """convert notebooks to Python script after save with nbconvert

    replaces `jupyter notebook --script`
    """
    from nbconvert.exporters.script import ScriptExporter

    if model['type'] != 'notebook':
        return

    global _script_exporter

    if _script_exporter is None:
        _script_exporter = ScriptExporter(parent=contents_manager)

    log = contents_manager.log

    base, ext = os.path.splitext(os_path)
    script, resources = _script_exporter.from_filename(os_path)
    script_fname = base + resources.get('output_extension', '.txt')
    log.info("Saving script %s", to_api_path(script_fname, contents_manager.root_
↪dir))

    with io.open(script_fname, 'w', encoding='utf-8') as f:
        f.write(script)

c.FileContentsManager.post_save_hook = script_post_save

```

This could be a simple call to `jupyter nbconvert --to script`, but spawning the subprocess every time is quite slow.

13.3 Custom request handlers

The notebook webserver can be interacted with using a well defined [RESTful API](#). You can define custom RESTful API handlers in addition to the ones provided by the notebook. As described below, to define a custom handler you need to first write a notebook server extension. Then, in the extension, you can register the custom handler.

13.3.1 Writing a notebook server extension

The notebook webserver is written in Python, hence your server extension should be written in Python too. Server extensions, like IPython extensions, are Python modules that define a specially named load function, `load_jupyter_server_extension`. This function is called when the extension is loaded.

```
def load_jupyter_server_extension(nb_server_app):
    """
    Called when the extension is loaded.

    Args:
        nb_server_app (NotebookWebApplication): handle to the Notebook webserver_
    instance.
    """
    pass
```

To get the notebook server to load your custom extension, you'll need to add it to the list of extensions to be loaded. You can do this using the config system. `NotebookApp.nbserver_extensions` is a config variable which is a dictionary of strings, each a Python module to be imported, mapping to `True` to enable or `False` to disable each extension. Because this variable is notebook config, you can set it two different ways, using config files or via the command line.

For example, to get your extension to load via the command line add a double dash before the variable name, and put the Python dictionary in double quotes. If your package is “mypackage” and module is “mymodule”, this would look like `jupyter notebook --NotebookApp.nbserver_extensions='{ "mypackage.mymodule": True }'`. Basically the string should be Python importable.

Alternatively, you can have your extension loaded regardless of the command line args by setting the variable in the Jupyter config file. The default location of the Jupyter config file is `~/.jupyter/jupyter_notebook_config.py` (see [Configuration Overview](#)). Inside the config file, you can use Python to set the variable. For example, the following config does the same as the previous command line example.

```
c = get_config()
c.NotebookApp.nbserver_extensions = {
    'mypackage.mymodule': True,
}
```

Before continuing, it's a good idea to verify that your extension is being loaded. Use a print statement to print something unique. Launch the notebook server and you should see your statement printed to the console.

13.3.2 Registering custom handlers

Once you’ve defined a server extension, you can register custom handlers because you have a handle to the Notebook server app instance (`nb_server_app` above). However, you first need to define your custom handler. To declare a custom handler, inherit from `notebook.base.handlers.IPythonHandler`. The example below[1] is a Hello World handler:

```
from notebook.base.handlers import IPythonHandler

class HelloWorldHandler(IPythonHandler):
    def get(self):
        self.finish('Hello, world!')
```

The Jupyter Notebook server use [Tornado](#) as its web framework. For more information on how to implement request handlers, refer to the [Tornado documentation on the matter](#).

After defining the handler, you need to register the handler with the Notebook server. See the following example:

```
web_app = nb_server_app.web_app
host_pattern = '.*$'
route_pattern = url_path_join(web_app.settings['base_url'], '/hello')
web_app.add_handlers(host_pattern, [(route_pattern, HelloWorldHandler)])
```

Putting this together with the extension code, the example looks like the following:

```
from notebook.utils import url_path_join
from notebook.base.handlers import IPythonHandler

class HelloWorldHandler(IPythonHandler):
    def get(self):
        self.finish('Hello, world!')

def load_jupyter_server_extension(nb_server_app):
    """
    Called when the extension is loaded.

    Args:
        nb_server_app (NotebookWebApplication): handle to the Notebook webserver_
    instance.
    """
    web_app = nb_server_app.web_app
    host_pattern = '.*$'
    route_pattern = url_path_join(web_app.settings['base_url'], '/hello')
    web_app.add_handlers(host_pattern, [(route_pattern, HelloWorldHandler)])
```

13.4 Extra Parameters and authentication

Here is a quick rundown of what you need to know to pass extra parameters to the handler and enable authentication:

- extra arguments to the `__init__` constructor are given in a dictionary after the handler class in `add_handlers`:

```
class HelloWorldHandler(IPythonHandler):

    def __init__(self, *args, **kwargs):
```

(continues on next page)

(continued from previous page)

```

        self.extra = kwargs.pop('extra')
        ...

def load_jupyter_server_extension(nb_server_app):

    ...

    web_app.add_handlers(host_pattern,
        [
            (route_pattern, HelloWorldHandler, {"extra": nb_server_app.extra})
        ])

```

All handler methods that require authentication ***MUST*** be decorated with `@tornado.web.authenticated`:

```

from tornado import web

class HelloWorldHandler(IPythonHandler):

    ...

    @web.authenticated
    def get(self, *args, **kwargs):
        ...

    @web.authenticated
    def post(self, *args, **kwargs):
        ...

```

References:

1. [Peter Parente's Mindtrove](#)

13.5 Custom front-end extensions

This describes the basic steps to write a JavaScript extension for the Jupyter notebook front-end. This allows you to customize the behaviour of the various pages like the dashboard, the notebook, or the text editor.

13.5.1 The structure of a front-end extension

Note: The notebook front-end and Javascript API are not stable, and are subject to a lot of changes. Any extension written for the current notebook is almost guaranteed to break in the next release.

A front-end extension is a JavaScript file that defines an [AMD module](#) which exposes at least a function called `load_ipython_extension`, which takes no arguments. We will not get into the details of what each of these terms consists of yet, but here is the minimal code needed for a working extension:

```

// file my_extension/main.js

define(function() {

    function load_ipython_extension() {

```

(continues on next page)

(continued from previous page)

```

        console.info('this is my first extension');
    }

    return {
        load_ipython_extension: load_ipython_extension
    };
});

```

Note: Although for historical reasons the function is called `load_ipython_extension`, it does apply to the Jupyter notebook in general, and will work regardless of the kernel in use.

If you are familiar with JavaScript, you can use this template to require any Jupyter module and modify its configuration, or do anything else in client-side Javascript. Your extension will be loaded at the right time during the notebook page initialisation for you to set up a listener for the various events that the page can trigger.

You might want access to the current instances of the various Jupyter notebook components on the page, as opposed to the classes defined in the modules. The current instances are exposed by a module named `base/js/namespace`. If you plan on accessing instances on the page, you should `require` this module rather than accessing the global variable `Jupyter`, which will be removed in future. The following example demonstrates how to access the current notebook instance:

```

// file my_extension/main.js

define([
    'base/js/namespace'
], function(
    Jupyter
) {
    function load_ipython_extension() {
        console.log(
            'This is the current notebook application instance:',
            Jupyter.notebook
        );
    }

    return {
        load_ipython_extension: load_ipython_extension
    };
});

```

13.5.2 Modifying key bindings

One of the abilities of extensions is to modify key bindings, although once again this is an API which is not guaranteed to be stable. However, custom key bindings are frequently requested, and are helpful to increase accessibility, so in the following we show how to access them.

Here is an example of an extension that will unbind the shortcut `0, 0` in command mode, which normally restarts the kernel, and bind `0, 0, 0` in its place:

```

// file my_extension/main.js

define([
    'base/js/namespace'

```

(continues on next page)

(continued from previous page)

```

], function(
  Jupyter
) {

  function load_ipython_extension() {
    Jupyter.keyboard_manager.command_shortcuts.remove_shortcut('0,0');
    Jupyter.keyboard_manager.command_shortcuts.add_shortcut('0,0,0', 'jupyter-
↪notebook:restart-kernel');
  }

  return {
    load_ipython_extension: load_ipython_extension
  };
});

```

Note: The standard keybindings might not work correctly on non-US keyboards. Unfortunately, this is a limitation of browser implementations and the status of keyboard event handling on the web in general. We appreciate your feedback if you have issues binding keys, or have any ideas to help improve the situation.

You can see that I have used the **action name** `jupyter-notebook:restart-kernel` to bind the new shortcut. There is no API yet to access the list of all available *actions*, though the following in the JavaScript console of your browser on a notebook page should give you an idea of what is available:

```
Object.keys(require('base/js/namespace').actions._actions);
```

In this example, we changed a keyboard shortcut in **command mode**; you can also customize keyboard shortcuts in **edit mode**. However, most of the keyboard shortcuts in edit mode are handled by CodeMirror, which supports custom key bindings via a completely different API.

13.5.3 Defining and registering your own actions

As part of your front-end extension, you may wish to define actions, which can be attached to toolbar buttons, or called from the command palette. Here is an example of an extension that defines an (not very useful!) action to show an alert, and adds a toolbar button using the full action name:

```

// file my_extension/main.js

define([
  'base/js/namespace'
], function(
  Jupyter
) {
  function load_ipython_extension() {

    var handler = function () {
      alert('this is an alert from my_extension!');
    };

    var action = {
      icon: 'fa-comment-o', // a font-awesome class used on buttons, etc
      help : 'Show an alert',
      help_index : 'zz',
      handler : handler
    };
  }
});

```

(continues on next page)

(continued from previous page)

```

    };
    var prefix = 'my_extension';
    var action_name = 'show-alert';

    var full_action_name = Jupyter.actions.register(action, action_name, prefix);
    ↪ // returns 'my_extension:show-alert'
    Jupyter.toolbar.add_buttons_group([full_action_name]);
  }

  return {
    load_ipython_extension: load_ipython_extension
  };
});

```

Every action needs a name, which, when joined with its prefix to make the full action name, should be unique. Built-in actions, like the `jupyter-notebook:restart-kernel` we bound in the earlier [Modifying key bindings](#) example, use the prefix `jupyter-notebook`. For actions defined in an extension, it makes sense to use the extension name as the prefix. For the action name, the following guidelines should be considered:

- First pick a noun and a verb for the action. For example, if the action is “restart kernel,” the verb is “restart” and the noun is “kernel”.
- Omit terms like “selected” and “active” by default, so “delete-cell”, rather than “delete-selected-cell”. Only provide a scope like “-all-” if it is other than the default “selected” or “active” scope.
- If an action has a secondary action, separate the secondary action with “-and-”, so “restart-kernel-and-clear-output”.
- Use above/below or previous/next to indicate spatial and sequential relationships.
- Don’t ever use before/after as they have a temporal connotation that is confusing when used in a spatial context.
- For dialogs, use a verb that indicates what the dialog will accomplish, such as “confirm-restart-kernel”.

13.5.4 Installing and enabling extensions

You can install your nbextension with the command:

```
jupyter nbextension install path/to/my_extension/ [--user|--sys-prefix]
```

The default installation is system-wide. You can use `--user` to do a per-user installation, or `--sys-prefix` to install to Python’s prefix (e.g. in a virtual or conda environment). Where `my_extension` is the directory containing the Javascript files. This will copy it to a Jupyter data directory (the exact location is platform dependent - see [jupyter_path](#)).

For development, you can use the `--symlink` flag to symlink your extension rather than copying it, so there’s no need to reinstall after changes.

To use your extension, you’ll also need to **enable** it, which tells the notebook interface to load it. You can do that with another command:

```
jupyter nbextension enable my_extension/main [--sys-prefix][--section='common']
```

The argument refers to the Javascript module containing your `load_ipython_extension` function, which is `my_extension/main.js` in this example. The `--section='common'` argument will affect all pages, by default it will be loaded on the notebook view only. There is a corresponding `disable` command to stop using an extension without uninstalling it.

Changed in version 4.2: Added `--sys-prefix` argument

13.5.5 Kernel Specific extensions

Warning: This feature serves as a stopgap for kernel developers who need specific JavaScript injected onto the page. The availability and API are subject to change at anytime.

It is possible to load some JavaScript on the page on a per kernel basis. Be aware that doing so will make the browser page reload without warning as soon as the user switches the kernel without notice.

If you, a kernel developer, need a particular piece of JavaScript to be loaded on a “per kernel” basis, such as:

- if you are developing a CodeMirror mode for your language
- if you need to enable some specific debugging options

your `kernelspecs` are allowed to contain a `kernel.js` file that defines an AMD module. The AMD module should define an *onload* function that will be called when the kernelspec loads, such as:

- when you load a notebook that uses your kernelspec
- change the active kernelspec of a notebook to your kernelspec.

Note that adding a *kernel.js* to your kernelspec will add an unexpected side effect to changing a kernel in the notebook. As it is impossible to “unload” JavaScript, any attempt to change the kernelspec again will save the current notebook and reload the page without confirmations.

Here is an example of `kernel.js`:

```
define(function() {
  return {onload: function() {
    console.info('Kernel specific javascript loaded');

    // do more things here, like define a codemirror mode

  }}
});
```

13.6 Customize keymaps

Note: Declarative Custom Keymaps is a provisional feature with unstable API which is not guaranteed to be kept in future versions of the notebook, and can be removed or changed without warnings.

The notebook shortcuts that are defined by jupyter both in edit mode and command mode are configurable in the frontend configuration file `~/.jupyter/nbconfig/notebook.json`. The modification of keyboard shortcuts suffers from several limitations, mainly that your Browser and OS might prevent certain shortcuts from working correctly. If this is the case, there is unfortunately not much that can be done. The second issue can arise with keyboards that have a layout different than US English. Again, even if we are aware of the issue, there is not much that can be done.

Shortcuts are also limited by the underlying library that handles code and text editing: CodeMirror. If some keyboard shortcuts are conflicting, the method described below might not work to create new keyboard shortcuts, especially in the `edit` mode of the notebook.

The 4 sections of interest in `~/.jupyter/nbconfig/notebook.json` are the following:

- `keys.command.unbind`
- `keys.edit.unbind`
- `keys.command.bind`
- `keys.edit.bind`

The first two sections describe which default keyboard shortcuts not to register at notebook startup time. These are mostly useful if you need to `unbind` a default keyboard shortcut before binding it to a new `command`.

The first two sections apply respectively to the `command` and `edit` mode of the notebook. They take a list of shortcuts to `unbind`.

For example, to unbind the shortcut to split a cell at the position of the cursor (`Ctrl-Shift-Minus`) use the following:

```
// file ~/.jupyter/nbconfig/notebook.json

{
  "keys": {
    "edit": {
      "unbind": [
        "Ctrl-Shift-Minus"
      ]
    },
  },
}
```

The last two sections describe which new keyboard shortcuts to register at notebook startup time and which actions they trigger.

The last two sections apply respectively to the `command` and `edit` mode of the notebook. They take a dictionary with shortcuts as `keys` and `commands` name as value.

For example, to bind the shortcut `G, G, G` (Press `G` three time in a row) in `command` mode to the command that restarts the kernel and runs all cells, use the following:

```
// file ~/.jupyter/nbconfig/notebook.json

{
  "keys": {
    "command": {
      "bind": {
        "G,G,G": "jupyter-notebook:restart-kernel-and-run-all-cells"
      }
    },
  },
}
```

The name of the available `commands` can be find by hovering over the right end of a row in the command palette.

13.7 Custom bundler extensions

The notebook server supports the writing of *bundler extensions* that transform, package, and download/deploy notebook files. As a developer, you need only write a single Python function to implement a bundler. The notebook server automatically generates a *File -> Download as* or *File -> Deploy as* menu item in the notebook front-end to trigger your bundler.

Here are some examples of what you can implement using bundler extensions:

- Convert a notebook file to a HTML document and publish it as a post on a blog site
- Create a snapshot of the current notebook environment and bundle that definition plus notebook into a zip download
- Deploy a notebook as a standalone, interactive [dashboard](#)

To implement a bundler extension, you must do all of the following:

- Declare bundler extension metadata in your Python package
- Write a *bundle* function that responds to bundle requests
- Instruct your users on how to enable/disable your bundler extension

The following sections describe these steps in detail.

13.7.1 Declaring bundler metadata

You must provide information about the bundler extension(s) your package provides by implementing a `_jupyter_bundlerextensions_paths` function. This function can reside anywhere in your package so long as it can be imported when enabling the bundler extension. (See [Enabling/disabling bundler extensions](#).)

```
# in mypackage.hello_bundler

def _jupyter_bundlerextension_paths():
    """Example "hello world" bundler extension"""
    return [{
        'name': 'hello_bundler',          # unique bundler name
        'label': 'Hello Bundler',        # human-readable menu item label
        'module_name': 'mypackage.hello_bundler', # module containing bundle()
        'group': 'deploy'                # group under 'deploy' or
        ↪ 'download' menu
    ]
```

Note that the return value is a list. By returning multiple dictionaries in the list, you allow users to enable/disable sets of bundlers all at once.

13.7.2 Writing the *bundle* function

At runtime, a menu item with the given label appears either in the *File -> Deploy as* or *File -> Download as* menu depending on the *group* value in your metadata. When a user clicks the menu item, a new browser tab opens and notebook server invokes a *bundle* function in the *module_name* specified in the metadata.

You must implement a *bundle* function that matches the signature of the following example:

```
# in mypackage.hello_bundler

def bundle(handler, model):
    """Transform, convert, bundle, etc. the notebook referenced by the given
    model.

    Then issue a Tornado web response using the `handler` to redirect
    the user's browser, download a file, show a HTML page, etc. This function
    must finish the handler response before returning either explicitly or by
    raising an exception.

    Parameters
    -----
    handler : tornado.web.RequestHandler
        Handler that serviced the bundle request
    model : dict
        Notebook model from the configured ContentManager
    """
    handler.finish('I bundled {}'.format(model['path']))
```

Your *bundle* function is free to do whatever it wants with the request and respond in any manner. For example, it may read additional query parameters from the request, issue a redirect to another site, run a local process (e.g., *nbconvert*), make a HTTP request to another service, etc.

The caller of the *bundle* function is `@tornado.gen.coroutine` decorated and wraps its call with `torando.gen.maybe_future`. This behavior means you may handle the web request synchronously, as in the example above, or asynchronously using `@tornado.gen.coroutine` and `yield`, as in the example below.

```
from tornado import gen

@gen.coroutine
def bundle(handler, model):
    # simulate a long running IO op (e.g., deploying to a remote host)
    yield gen.sleep(10)

    # now respond
    handler.finish('I spent 10 seconds bundling {}'.format(model['path']))
```

You should prefer the second, asynchronous approach when your bundle operation is long-running and would otherwise block the notebook server main loop if handled synchronously.

For more details about the data flow from menu item click to bundle function invocation, see [Bundler invocation details](#).

13.7.3 Enabling/disabling bundler extensions

The notebook server includes a command line interface (CLI) for enabling and disabling bundler extensions.

You should document the basic commands for enabling and disabling your bundler. One possible command for enabling the *hello_bundler* example is the following:

```
jupyter bundlerextension enable --py mypackage.hello_bundler --sys-prefix
```

The above updates the notebook configuration file in the current conda/virtualenv environment (`--sys-prefix`) with the metadata returned by the `mypackage.hellow_bundler._jupyter_bundlerextension_paths` function.

The corresponding command to later disable the bundler extension is the following:

```
jupyter bundlerextension disable --py mypackage.hello_bundler --sys-prefix
```

For more help using the *bundlerextension* subcommand, run the following.

```
jupyter bundlerextension --help
```

The output describes options for listing enabled bundlers, configuring bundlers for single users, configuring bundlers system-wide, etc.

13.7.4 Example: IPython Notebook bundle (.zip)

The *hello_bundler* example in this documentation is simplistic in the name of brevity. For more meaningful examples, see *notebook/bundler/zip_bundler.py* and *notebook/bundler/tarball_bundler.py*. You can enable them to try them like so:

```
jupyter bundlerextension enable --py notebook.bundler.zip_bundler --sys-prefix  
jupyter bundlerextension enable --py notebook.bundler.tarball_bundler --sys-prefix
```

13.7.5 Bundler invocation details

Support for bundler extensions comes from Python modules in *notebook/bundler* and JavaScript in *notebook/static/notebook/js/menubar.js*. The flow of data between the various components proceeds roughly as follows:

1. User opens a notebook document
2. Notebook front-end JavaScript loads notebook configuration
3. Bundler front-end JS creates menu items for all bundler extensions in the config
4. User clicks a bundler menu item
5. JS click handler opens a new browser window/tab to `<notebook base_url>/bundle/<path/to/notebook>?bundler=<name>` (i.e., a HTTP GET request)
6. Bundle handler validates the notebook path and bundler *name*
7. Bundle handler delegates the request to the *bundle* function in the bundler's *module_name*
8. *bundle* function finishes the HTTP request

CONTRIBUTING TO THE JUPYTER NOTEBOOK

If you're reading this section, you're probably interested in contributing to Jupyter. Welcome and thanks for your interest in contributing!

Please take a look at the Contributor documentation, familiarize yourself with using the Jupyter Notebook, and introduce yourself on the mailing list and share what area of the project you are interested in working on.

14.1 General Guidelines

For general documentation about contributing to Jupyter projects, see the [Project Jupyter Contributor Documentation](#).

14.2 Setting Up a Development Environment

14.2.1 Installing Node.js and npm

Building the Notebook from its GitHub source code requires some tools to create and minify JavaScript components and the CSS, specifically Node.js and Node's package manager, npm. It should be node version 6.0.

If you use `conda`, you can get them with:

```
conda install -c conda-forge nodejs
```

If you use [Homebrew](#) on Mac OS X:

```
brew install node
```

Installation on Linux may vary, but be aware that the *nodejs* or *npm* packages included in the system package repository may be too old to work properly.

You can also use the installer from the [Node.js website](#).

14.2.2 Installing the Jupyter Notebook

Once you have installed the dependencies mentioned above, use the following steps:

```
pip install --upgrade setuptools pip
git clone https://github.com/jupyter/notebook
cd notebook
pip install -e .
```


If you are using a system-wide Python installation and you only want to install the notebook for you, you can add `--user` to the install commands.

Once you have done this, you can launch the master branch of Jupyter notebook from any directory in your system with:

```
jupyter notebook
```

14.2.3 Verification

While running the notebook, select one of your notebook files (the file will have the extension `.ipynb`). In the top tab you will click on “Help” and then click on “About”. In the pop window you will see information about the version of Jupyter that you are running. You will see “The version of the notebook server is:”. If you are working in development mode, you will see that your version of Jupyter notebook will include the word “dev”.



If it does not include the word “dev”, you are currently not working in development mode and should follow the steps below to uninstall and reinstall Jupyter.

14.2.4 Troubleshooting the Installation

If you do not see that your Jupyter Notebook is not running on dev mode, it’s possible that you are running other instances of Jupyter Notebook. You can try the following steps:

1. Uninstall all instances of the notebook package. These include any installations you made using pip or conda.
2. Run `python3 -m pip install -e .` in the notebook repository to install the notebook from there.
3. Run `npm run build` to make sure the Javascript and CSS are updated and compiled.
4. Launch with `python3 -m notebook --port 8989`, and check that the browser is pointing to `localhost:8989` (rather than the default 8888). You don’t necessarily have to launch with port 8989, as long as you use a port that is neither the default nor in use, then it should be fine.
5. Verify the installation with the steps in the previous section.

14.2.5 Rebuilding JavaScript and CSS

There is a build step for the JavaScript and CSS in the notebook. To make sure that you are working with up-to-date code, you will need to run this command whenever there are changes to JavaScript or LESS sources:

```
npm run build
```

IMPORTANT: Don't forget to run `npm run build` after switching branches. When switching between branches of different versions (e.g. `4.x` and `master`), run `pip install -e ..`. If you have tried the above and still find that the notebook is not reflecting the current source code, try cleaning the repo with `git clean -xfd` and reinstalling with `pip install -e ..`.

Development Tip

When doing development, you can use this command to automatically rebuild JavaScript and LESS sources as they are modified:

```
npm run build:watch
```

Git Hooks

If you want to automatically update dependencies and recompile JavaScript and CSS after checking out a new commit, you can install post-checkout and post-merge hooks which will do it for you:

```
git-hooks/install-hooks.sh
```

See `git-hooks/README.md` for more details.

14.3 Running Tests

14.3.1 Python Tests

Install dependencies:

```
pip install -e '.[test]'
```

To run the Python tests, use:

```
pytest
```

If you want coverage statistics as well, you can run:

```
py.test --cov notebook -v --pyargs notebook
```

14.3.2 JavaScript Tests

To run the JavaScript tests, you will need to have PhantomJS and CasperJS installed:

```
npm install -g casperjs phantomjs-prebuilt
```

Then, to run the JavaScript tests:

```
python -m notebook.jstest [group]
```

where `[group]` is an optional argument that is a path relative to `notebook/tests/`. For example, to run all tests in `notebook/tests/notebook`:

```
python -m notebook.jstest notebook
```

or to run just `notebook/tests/notebook/deletecell.js`:

```
python -m notebook.jstest notebook/deletecell.js
```

14.4 Building the Documentation

To build the documentation you'll need [Sphinx](#), [pandoc](#) and a few other packages.

To install (and activate) a conda environment named `notebook_docs` containing all the necessary packages (except `pandoc`), use:

```
conda create -n notebook_docs pip
conda activate notebook_docs # Linux and OS X
activate notebook_docs      # Windows
pip install .[docs]
```

If you want to install the necessary packages with `pip`, use the following instead:

```
pip install .[docs]
```

Once you have installed the required packages, you can build the docs with:

```
cd docs
make html
```

After that, the generated HTML files will be available at `build/html/index.html`. You may view the docs in your browser.

You can automatically check if all hyperlinks are still valid:

```
make linkcheck
```

Windows users can find `make.bat` in the docs folder.

You should also have a look at the [Project Jupyter Documentation Guide](#).

DEVELOPER FAQ

1. How do I install a prerelease version such as a beta or release candidate?

```
python -m pip install notebook --pre --upgrade
```


MY NOTEBOOK

```
[1]: def foo():  
     return "foo"
```

```
[2]: def has_ip_syntax():  
     listing = !ls  
     return listing
```

```
[4]: def whatsmyname():  
     return __name__
```


OTHER NOTEBOOK

This notebook just defines `bar`

```
[2]: def bar(x):  
      return "bar" * x
```